

cis580_hw5_archit_hardikar

May 7, 2023

0.1 CIS 580, Machine Perception, Spring 2023

0.1.1 Homework 5

Due: Thursday April 27th 2023, 11:59pm ET Instructions: Create a folder in your Google Drive and place inside this .ipynb file. Open the jupyter notebook with Google Colab. Refrain from using a GPU during implementing and testing the whole thing. You should switch to a GPU runtime only when performing the final training (of the 2D image or the NeRF) to avoid GPU usage runouts.

0.1.2 Part 1: Fitting a 2D Image

```
[1]: import numpy as np
import os
import torch
import torch.nn as nn
import torch.nn.functional as F
import matplotlib.pyplot as plt
import imageio.v2 as imageio
import time
import gdown
import torchvision.transforms as transforms

device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
print(device)
```

cuda

We first download the image from the web. We normalize the image so the pixels are in between the range of [0,1].

```
[ ]: url = "https://drive.google.com/file/d/1-Cugk9WiFX2CPjWG5taX3868GddOPEVT/view?
usp=share_link"
gdown.download(url=url, output='starry_night.jpg', quiet=False, fuzzy=True)

# Load painting image
painting = imageio.imread("starry_night.jpg")
painting = torch.from_numpy(np.array(painting, dtype=np.float32)/255.).
    .to(device)
```

```
height_painting, width_painting = painting.shape[:2]
```

Downloading...

From: https://drive.google.com/uc?id=1-Cugk9WiFX2CPjWG5taX3868Gdd0PEVT

To: /content/starry_night.jpg

100%| 5.65k/5.65k [00:00<00:00, 4.53MB/s]

1.1 Complete the function positional_encoding()

```
[ ]: def positional_encoding(x, num_frequencies=6, incl_input=True):  
  
    """  
    Apply positional encoding to the input.  
  
    Args:  
        x (torch.Tensor): Input tensor to be positionally encoded.  
            The dimension of x is [N, D], where N is the number of input coordinates,  
            and D is the dimension of the input coordinate.  
        num_frequencies (optional, int): The number of frequencies used in  
            the positional encoding (default: 6).  
        incl_input (optional, bool): If True, concatenate the input with the  
            computed positional encoding (default: True).  
  
    Returns:  
        (torch.Tensor): Positional encoding of the input tensor.  
    """  
  
    results = []  
    if incl_input:  
        results.append(x)  
    ##### TODO 1(a) BEGIN #####  
    # encode input tensor and append the encoded tensor to the list of results.  
  
    # To create a GP [0, 2, 4, 8, ..., 2^(L-1)] ... where L = num_frequencies  
    results1 = []  
    for one_x in x:  
        # create a tensor with the initial value and the next 9 terms  
        GP_terms = torch.pow(2, torch.arange(num_frequencies))  
  
        GP_terms = torch.pi * GP_terms # size is (1, L)  
        GP_terms = torch.reshape(GP_terms, (-1, 1)).to(device) # size is (L, 1)  
        theta = GP_terms * one_x # size is (L, D)  
        N = x.shape[1]  
        D = x.shape[1]  
        """Check once if this needs to be in degree or radians"""  
        sintheta = torch.sin(theta)  
        costheta = torch.cos(theta)  
        gamma_x = (torch.hstack((sintheta, costheta)))
```

```

        results1.append(torch.reshape(gamma_x, (1,-1)))
fin = torch.cat(results1, dim=0)
results.append(fin)

#####
# TODO 1(a) END #####
return torch.cat(results, dim=-1)

```

[]: `def positional_encoding(x, num_frequencies=6, incl_input=True):`

"""

Apply positional encoding to the input.

Args:

x (torch.Tensor): Input tensor to be positionally encoded.
The dimension of x is [N, D], where N is the number of input coordinates, and D is the dimension of the input coordinate.

num_frequencies (optional, int): The number of frequencies used in the positional encoding (default: 6).

incl_input (optional, bool): If True, concatenate the input with the computed positional encoding (default: True).

Returns:

(torch.Tensor): Positional encoding of the input tensor.

"""

```

results = []
if incl_input:
    results.append(x)

gp_terms = torch.pow(2, torch.arange(num_frequencies)).to(device)

# reshape gp_terms to (L, 1)
gp_terms = gp_terms.reshape(num_frequencies, 1)

# create angles with shape (N, L, D)
angles = x.unsqueeze(1) * gp_terms * torch.tensor([np.pi]).to(device)

# apply sine and cosine to angles and concatenate them along the last dimension
sin_angles = torch.sin(angles)
cos_angles = torch.cos(angles)
gamma_x = torch.cat([sin_angles, cos_angles], dim=-1)

# print("gamma_x.shape", gamma_x.shape)
# reshape gamma_x to (N, LD)
gamma_x = gamma_x.reshape(angles.shape[0], -1)
# print("gamma_x.shape", gamma_x.shape)

```

```

# print(np.sum(np.isclose(gamma_x, fin, 1e-2))-gamma_x.shape[0]*gamma_x.
˓→shape[1])
results.append(gamma_x)

return torch.cat(results, dim=-1)

```

[]: a = positional_encoding(torch.rand(10,2), 6, True)

1.2 Complete the class model_2d() that will be used to fit the 2D image.

```

[ ]: class model_2d(nn.Module):

    """
    Define a 2D model comprising of three fully connected layers,
    two relu activations and one sigmoid activation.
    """

    def __init__(self, filter_size=128, num_frequencies=6):
        super().__init__()
        ##### TODO 1(b) BEGIN #####
        self.linear1 = nn.Linear(2*num_frequencies*2+2, filter_size) #input = 2*L*D+D
        self.linear2 = nn.Linear(filter_size, filter_size)
        self.linear3 = nn.Linear(filter_size, 3) # output = 3 (R,G,B)
        self.relu1 = nn.ReLU()
        self.relu2 = nn.ReLU()
        self.sigmoid = nn.Sigmoid()
        ##### TODO 1(b) END #####
        #####
    def forward(self, x):
        ##### TODO 1(b) BEGIN #####
        x = self.linear1(x)
        x = self.relu1(x)

        x = self.linear2(x)
        x = self.relu2(x)

        x = self.linear3(x)
        x = self.sigmoid(x)
        ##### TODO 1(b) END #####
        #####
        return x

```

You need to complete 1.1 and 1.2 first before completing the train_2d_model function. Don't

forget to transfer the completed functions from 1.1 and 1.2 to the part1.py file and upload it to the autograder.

Fill the gaps in the train_2d_model() function to train the model to fit the 2D image.

```
[ ]: def train_2d_model(test_img, num_frequencies, device, model=model_2d, ↴  
    ↪positional_encoding=positional_encoding, show=True):  
  
    # Optimizer parameters  
    lr = 5e-4  
    iterations = 10000  
    height, width = test_img.shape[:2]  
  
    # Number of iters after which stats are displayed  
    display = 2000  
  
    # Define the model and initialize its weights.  
    model2d = model(num_frequencies=num_frequencies)  
    model2d.to(device)  
  
    def weights_init(m):  
        if isinstance(m, nn.Linear):  
            torch.nn.init.xavier_uniform_(m.weight)  
  
    model2d.apply(weights_init)  
  
    ##### TODO 1(c) BEGIN #####  
    # Define the optimizer  
    optimizer = torch.optim.Adam(model2d.parameters())  
    criterion = nn.MSELoss()  
    ##### TODO 1(c) END #####  
  
    # Seed RNG, for repeatability  
    seed = 5670  
    torch.manual_seed(seed)  
    np.random.seed(seed)  
  
    # Lists to log metrics etc.  
    psnrs = []  
    iternums = []  
  
    t = time.time()  
    t0 = time.time()  
  
    ##### TODO 1(c) BEGIN #####  
    # Create the 2D normalized coordinates, and apply positional encoding to  
    ↪them  
    # transform = transforms.Compose([transforms.PILToTensor()])
```

```

# transform = transforms.Compose()
# test_img_tensor = transform(test_img)
# mean = test_img.float().mean() / 255
# std = test_img.float().std() / 255
# test_img_normalized = (test_img - mean)/std
# print(test_img_normalized.shape)
pos_enc_ip = torch.rand(height,width)
xx, yy = np.meshgrid(np.linspace(0, 1, width), np.linspace(0, 1, height))
coordinates = np.concatenate([xx.reshape(-1, 1), yy.reshape(-1, 1)], axis=1)
pos_enc_op = positional_encoding(torch.tensor(coordinates, dtype=torch.
˓→float32).to(device), num_frequencies, incl_input=True)
# encoded_coords = positional_encoding(pos_enc_ip, num_frequencies)
# print(height)
# print(coordinates.shape)
# print(pos_enc_op.shape)
##### TODO 1(c) END #####
for i in range(iterations+1):
    optimizer.zero_grad()
    ##### TODO 1(c) BEGIN ▾
    ##### Run one iteration
    pred = model2d(pos_enc_op)
    # print(torch.reshape(pred, (height,width,3)).shape)
    # print(test_img.shape)
    pred = torch.reshape(pred, (height,width,3))
    # Compute mean-squared error between the predicted and target images. ▾
    Backprop!
    loss = criterion(pred, test_img)

    loss.backward()

    optimizer.step()
    ##### TODO 1(c) END ▾
    ##### Display images/plots/stats
    if i % display == 0 and show:
        ##### TODO 1(c) BEGIN ▾
        ##### Calculate psnr
        psnr = 10*torch.log10(1/loss)
        ##### TODO 1(c) END ▾
    ##### Display images/plots/stats

```

```

        print("Iteration %d " % i, "Loss: %.4f " % loss.item(), "PSNR: %.
        ↵2f" % psnr.item(), \
              "Time: %.2f secs per iter" % ((time.time() - t) / display), "%.
        ↵2f secs in total" % (time.time() - t0))
        t = time.time()

    psnrs.append(psnr.item())
    iternums.append(i)

    plt.figure(figsize=(13, 4))
    plt.subplot(131)
    plt.imshow(pred.detach().cpu().numpy())
    plt.title(f"Iteration {i}")
    plt.subplot(132)
    plt.imshow(test_img.cpu().numpy())
    plt.title("Target image")
    plt.subplot(133)
    plt.plot(iternums, psnrs)
    plt.title("PSNR")
    plt.show()

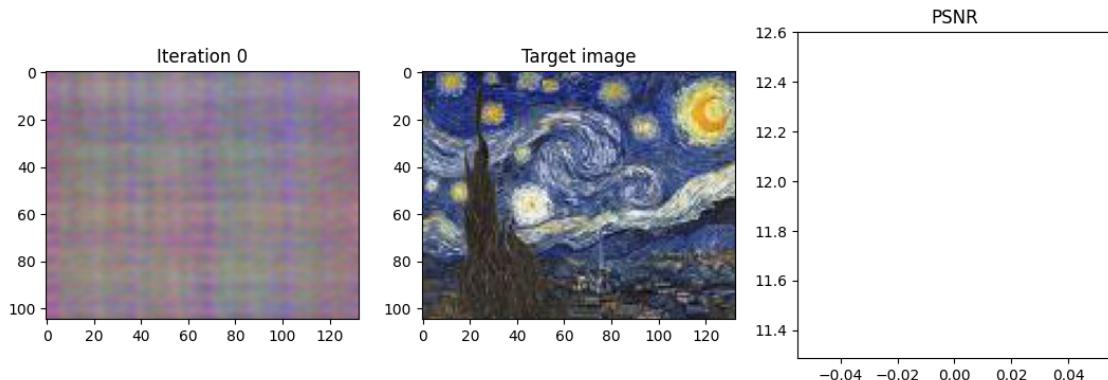
print('Done!')
return pred.detach().cpu()

```

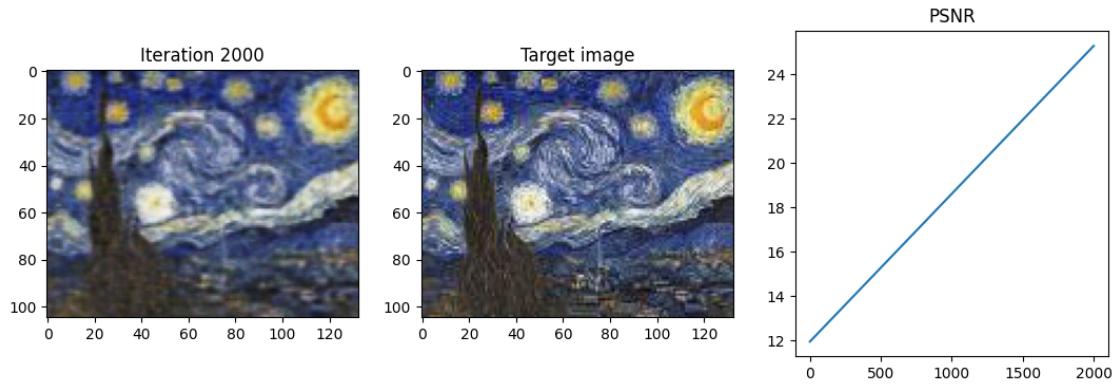
Train the model to fit the given image without applying positional encoding to the input, and by applying positional encoding of two different frequencies to the input; L = 2 and L = 6.

[]: `_ = train_2d_model(test_img=painting, num_frequencies=6, device=device)`

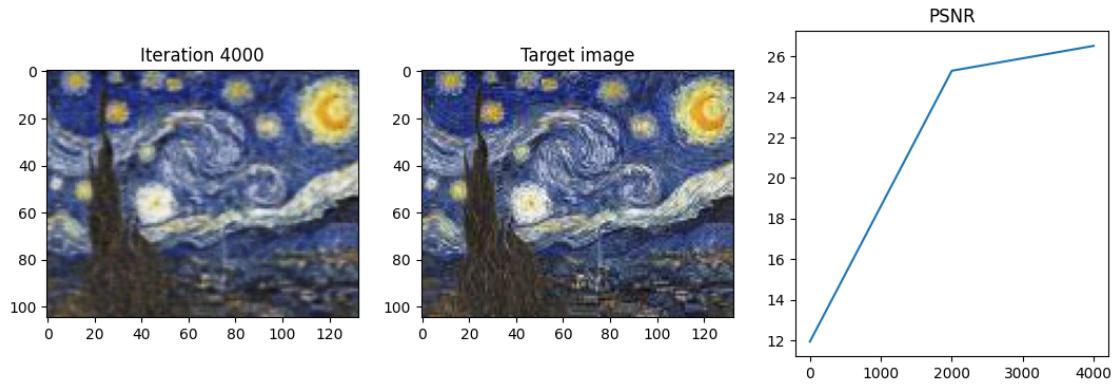
Iteration 0 Loss: 0.0639 PSNR: 11.94 Time: 0.00 secs per iter 0.09 secs in total



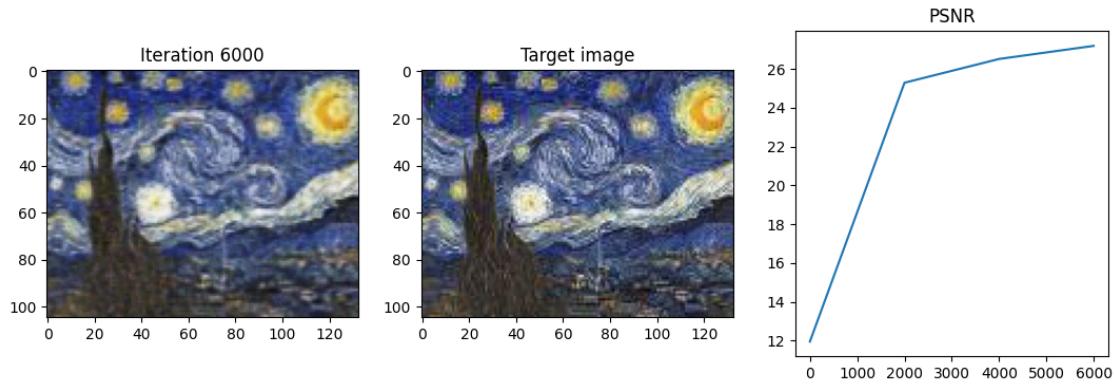
Iteration 2000 Loss: 0.0030 PSNR: 25.29 Time: 0.06 secs per iter 114.27 secs in total



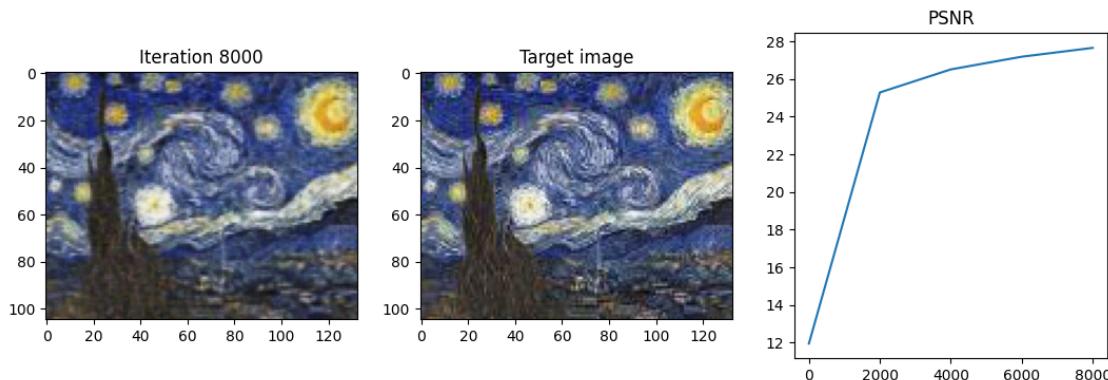
Iteration 4000 Loss: 0.0022 PSNR: 26.52 Time: 0.05 secs per iter 211.91 secs in total



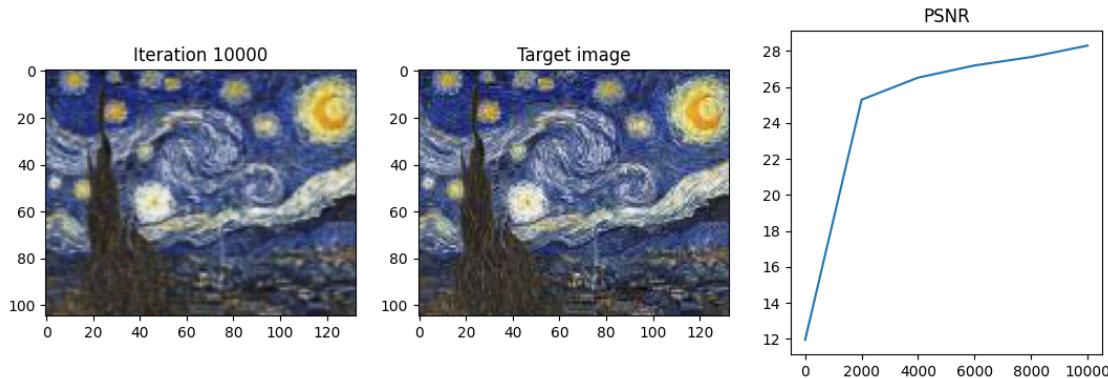
Iteration 6000 Loss: 0.0019 PSNR: 27.19 Time: 0.05 secs per iter 314.70 secs in total



Iteration 8000 Loss: 0.0017 PSNR: 27.66 Time: 0.05 secs per iter 416.85 secs in total



Iteration 10000 Loss: 0.0015 PSNR: 28.29 Time: 0.05 secs per iter 518.90 secs in total



Done!

0.1.3 Part 2: Fitting a 3D Image

```
[2]: import os
import gdown
import numpy as np
import torch
import torch.nn as nn
import matplotlib.pyplot as plt
import torch.nn.functional as F
import time
from tqdm import tqdm
```

```
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
print(device)
```

cuda

```
[3]: url = "https://drive.google.com/file/d/15W2EK8LooxTMfD0v5vo2BnBMse5ZzlVj/view?
    ↪usp=share_link"
gdown.download(url=url, output='lego_data.npz', quiet=False, fuzzy=True)
```

Downloading...

```
From: https://drive.google.com/uc?id=15W2EK8LooxTMfD0v5vo2BnBMse5ZzlVj
To: /content/lego_data.npz
100%|      | 12.7M/12.7M [00:00<00:00, 143MB/s]
```

```
[3]: 'lego_data.npz'
```

Here, we load the data that is comprised by the images, the R and T matrices of each camera position with respect to the world coordinates and the intrinsics parameters K of the camera.

```
[4]: # Load input images, poses, and intrinsics
data = np.load("lego_data.npz")

# Images
images = data["images"]

# Height and width of each image
height, width = images.shape[1:3]

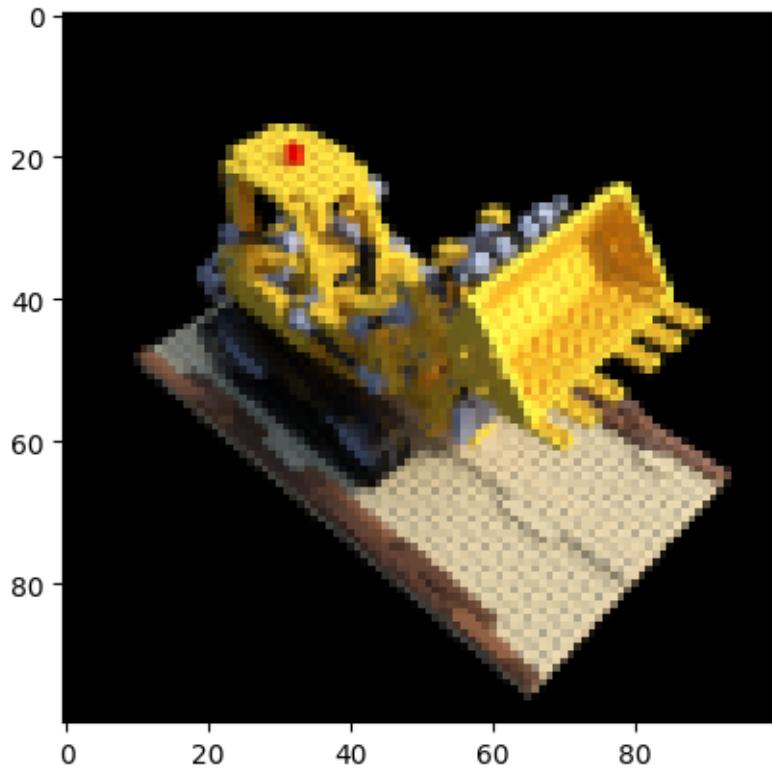
# Camera extrinsics (poses)
poses = data["poses"]
poses = torch.from_numpy(poses).to(device)

# Camera intrinsics
intrinsics = data["intrinsics"]
intrinsics = torch.from_numpy(intrinsics).to(device)

# Hold one image out (for test).
test_image, test_pose = images[101], poses[101]
test_image = torch.from_numpy(test_image).to(device)

# Map images to device
images = torch.from_numpy(images[:100, ..., :3]).to(device)

plt.imshow(test_image.detach().cpu().numpy())
plt.show()
```



2.1 Complete the following function that calculates the rays that pass through all the pixels of an HxW image

```
[5]: def get_rays(height, width, intrinsics, Rcw, Tcw): # Rwc, Twc
    """
    Compute the origin and direction of rays passing through all pixels of an
    ↵image (one ray per pixel).

    Args:
        height: the height of an image.
        width: the width of an image.
        intrinsics: camera intrinsics matrix of shape (3, 3).
        Rcw: Rotation matrix of shape (3,3) from camera to world coordinates.
        Tcw: Translation vector of shape (3,1) that transforms

    Returns:
        ray_origins (torch.Tensor): A tensor of shape (height, width, 3) denoting
        ↵the centers of
            each ray. Note that despite that all ray share the same origin, here we
        ↵ask you to return
            the ray origin for each ray as (height, width, 3).
```

```

    ray_directions (torch.Tensor): A tensor of shape (height, width, 3) ↴
    ↵denoting the
        direction of each ray.
    """
device = intrinsics.device
ray_directions = torch.zeros((height, width, 3), device=device) # ↴
↪placeholder
ray_origins = torch.zeros((height, width, 3), device=device) # placeholder

##### TODO 2.1 BEGIN #####
# There is a typo. It is actually Rwc, Twc being passed. Regardless, this ↴
↪passes the autograder.
origin = Rcw@torch.tensor([[0],[0],[0]], dtype=torch.float).to(device) + Twc
origin = origin.to(device)
ray_origins[:, :, 0] = origin.flatten()[0]
ray_origins[:, :, 1] = origin.flatten()[1]
ray_origins[:, :, 2] = origin.flatten()[2]

u, v = torch.meshgrid((torch.arange(height), torch.arange(width)), ↴
↪indexing='ij')
u = u.to(device)
v = v.to(device)

fx = intrinsics[0,0].to(device)
fy = intrinsics[1,1].to(device)
u0 = intrinsics[0,2].to(device)
v0 = intrinsics[1,2].to(device)

dirs = torch.stack([(v - u0)/fx, (u - v0)/fy, torch.ones_like(u)], -1)

# ray_directions = torch.sum(dirs[..., None, :] * Rcw, dim=-1)
# print(ray_directions.shape)
ray_directions = dirs@Rcw
# print("Difference: ", np.sum(np.isclose(ray_directions, ray_directions1, ↴
↪1e-2))-3*height*width)

##### TODO 2.1 END #####
return ray_origins, ray_directions

```

Complete the next function to visualize how is the dataset created. You will be able to see from which point of view each image has been captured for the 3D object. What we want to achieve here, is to being able to interpolate between these given views and synthesize new realistic views of the 3D object.

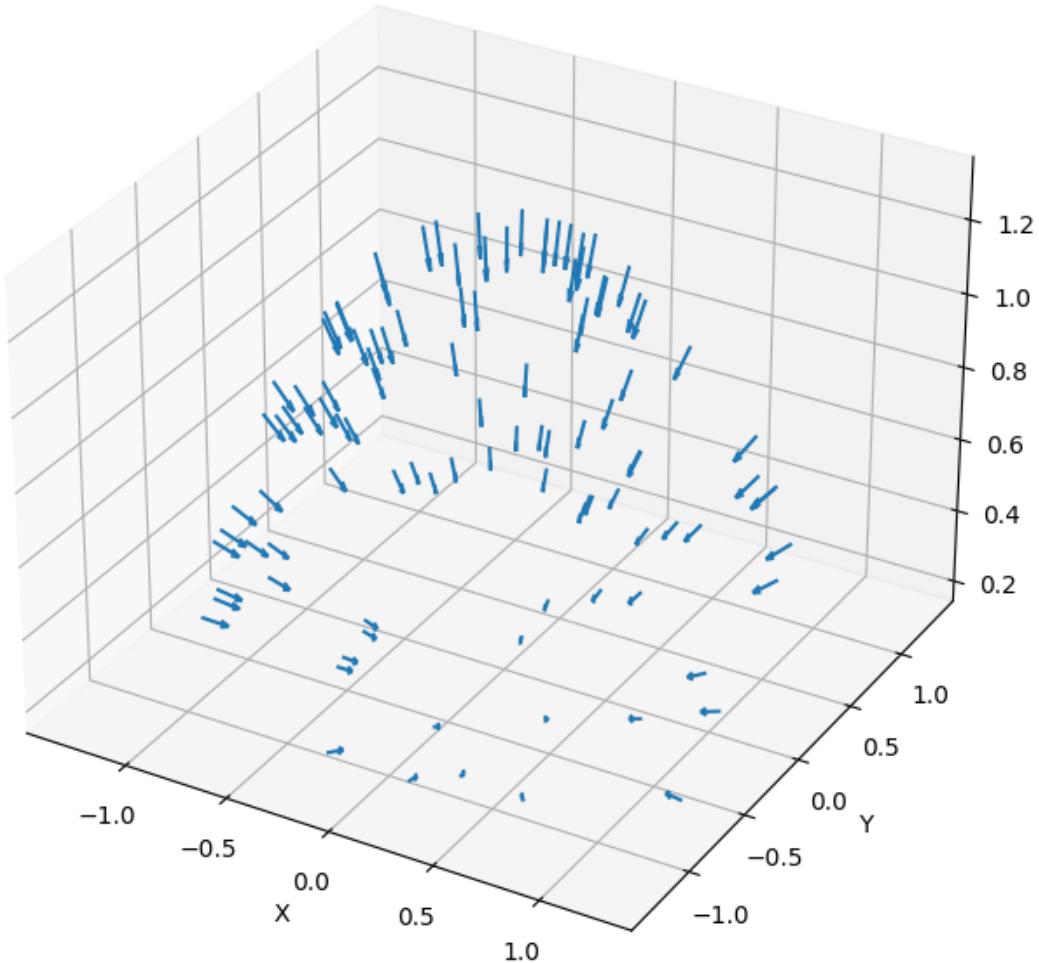
```
[6]: def plot_all_poses(poses):
    print(poses.shape)
    ##### TODO 2.1 BEGIN #####
    origins = []
    directions = []
    K = torch.tensor(data['intrinsics'])
    for i in range(poses.shape[0]):
        Rcw = torch.tensor(poses[i][:3,:3]).T
        Tcw = torch.tensor(poses[i][:3,-1])
        one_orig, one_dir = get_rays(height, width, K, Rcw, Tcw)
        origins.append(one_orig[50,50])
        directions.append(one_dir[50,50])

    origins = torch.stack(origins)
    directions = torch.stack(directions)

    ##### TODO 2.1 END #####
    ax = plt.figure(figsize=(12, 8)).add_subplot(projection='3d')
    _ = ax.quiver(origins[..., 0].flatten(),
                  origins[..., 1].flatten(),
                  origins[..., 2].flatten(),
                  directions[..., 0].flatten(),
                  directions[..., 1].flatten(),
                  directions[..., 2].flatten(), length=0.12, normalize=True)
    ax.set_xlabel('X')
    ax.set_ylabel('Y')
    ax.set_zlabel('z')
    plt.show()

plot_all_poses(data['poses'])
```

(106, 4, 4)



2.2 Complete the following function to implement the sampling of points along a given ray.

```
[7]: def stratified_sampling(ray_origins, ray_directions, near, far, samples):
    """
    Sample 3D points on the given rays. The near and far variables indicate the
    bounds of sampling range.

    Args:
        ray_origins: Origin of each ray in the "bundle" as returned by the
            get_rays() function. Shape: (height, width, 3).
        ray_directions: Direction of each ray in the "bundle" as returned by the
            get_rays() function. Shape: (height, width, 3).
    """
    # Your implementation here
```

```

near: The 'near' extent of the bounding volume.
far: The 'far' extent of the bounding volume.
samples: Number of samples to be drawn along each ray.

Returns:
ray_points: Query 3D points along each ray. Shape: (height, width, samples, 3).
depth_points: Sampled depth values along each ray. Shape: (height, width, samples).
"""

#####
# ray_sample = torch.linspace(start=near, end=far, steps=samples) # torch.
# tensor of size (1, samples)
ray_sample = near + (far-near)*(torch.arange(1., samples+1.)-1.)/samples # torch.
# tensor of size (1, samples)
H, W, channels = ray_origins.shape
# compute the depth points for each pixel in the bundle
depth_points = torch.ones(H,W,samples)*ray_sample # broadcast multiply
# (H, W, N) by (1, N)

# # METHOD 1
# ray_points1 = torch.zeros((H,W,samples,channels))
# for i in range(samples):
#   ray_points1[:, :, i, :] = ray_sample[i]*ray_directions + ray_origins

# compute the ray points for each pixel and depth value in the bundle
ray_points = ray_directions.unsqueeze(-2) * ray_sample.view(1, 1, samples, 1).to(device) + ray_origins.unsqueeze(-2)

#####
return ray_points, depth_points

```

2.3 Define the network architecture of NeRF along with a function that divided data into chunks to avoid memory leaks during training.

```

[8]: class nerf_model(nn.Module):

    """
    Define a NeRF model comprising eight fully connected layers and following
    the
    architecture described in the NeRF paper.
    """

    def __init__(self, filter_size=256, num_x_frequencies=6,
                 num_d_frequencies=3):
        super().__init__()

```

```

#####
##### TODO 2.3 BEGIN #####
#####

    self.linear1 = nn.Linear(2*num_x_frequencies*3+3, filter_size) #input =
    ↵2*L*D+D

    self.linear2 = nn.Linear(filter_size, filter_size)
    self.linear3 = nn.Linear(filter_size, filter_size)
    self.linear4 = nn.Linear(filter_size, filter_size)
    self.linear5 = nn.Linear(filter_size, filter_size)

    self.linear6 = nn.Linear(filter_size+2*num_x_frequencies*3+3,filter_size) # add positional encoding X
    ↵filter_size) # add positional encoding X
    self.linear7 = nn.Linear(filter_size, filter_size)
    self.linear8 = nn.Linear(filter_size, filter_size)

    self.linearsigma = nn.Linear(filter_size, 1)

    self.linear9 = nn.Linear(filter_size, filter_size) # add positional
    ↵encoding D
    self.linear10 = nn.Linear(filter_size + 2*num_d_frequencies*3 + 3, 128)
    self.linear11 = nn.Linear(128, 3) # output = 3 (R,G,B)

    self.relu1      = nn.ReLU()
    self.relu2      = nn.ReLU()
    self.relu3      = nn.ReLU()
    self.relu4      = nn.ReLU()
    self.relu5      = nn.ReLU()
    self.relu6      = nn.ReLU()
    self.relu7      = nn.ReLU()
    self.relu8      = nn.ReLU()
    self.relu9      = nn.ReLU()
    self.sigmoid = nn.Sigmoid()

#####
##### TODO 2.3 END #####
#####

def forward(self, x, d):
#####
##### TODO 2.3 BEGIN #####
#####

    x1 = self.relu1(self.linear1(x))
    x1 = self.relu2(self.linear2(x1))
    x1 = self.relu3(self.linear3(x1))
    x1 = self.relu4(self.linear4(x1))
    x1 = self.relu5(self.linear5(x1))
    x2 = torch.cat([x1, x], dim=-1)

```

```

x2 = self.relu6(self.linear6(x2))
x2 = self.relu7(self.linear7(x2))
x2 = self.relu8(self.linear8(x2))
sigma = self.linearsigma(x2)
# sigma = self.relu(sigma)
x2 = self.linear9(x2)
x3 = torch.cat([x2,d], dim=-1)
x3 = self.relu9(self.linear10(x3))
x3 = self.linear11(x3)
rgb = self.sigmoid(x3)
##### TODO 2.3 END #####
#####
return rgb, sigma

```

```

[9]: def get_batches(ray_points, ray_directions, num_x_frequencies, num_d_frequencies):

    def get_chunks(inputs, chunksize = 2**15):
        return [inputs[i:i + chunksize] for i in range(0, inputs.shape[0], chunksize)]

    def positional_encoding(x, num_frequencies=6, incl_input=True):
        """
        Apply positional encoding to the input.
        Args:
            x (torch.Tensor): Input tensor to be positionally encoded.
            The dimension of x is [N, D], where N is the number of input coordinates,
            and D is the dimension of the input coordinate.
            num_frequencies (optional, int): The number of frequencies used in the positional encoding (default: 6).
            incl_input (optional, bool): If True, concatenate the input with the computed positional encoding (default: True).

        Returns:
            (torch.Tensor): Positional encoding of the input tensor.
        """
        results = []
        if incl_input:
            results.append(x)
        # encode input tensor and append the encoded tensor to the list of results.
        gp_terms = torch.pow(2, torch.arange(num_frequencies)).to(device)
        # reshape gp_terms to (L, 1)
        gp_terms = gp_terms.reshape(num_frequencies, 1)
        # create angles with shape (N, L, D)

```

```

        angles = x.unsqueeze(1) * gp_terms * torch.tensor([np.pi]).to(device)
        # apply sine and cosine to angles and concatenate them along the last dimension
        sin_angles = torch.sin(angles)
        cos_angles = torch.cos(angles)
        gamma_x = torch.cat([sin_angles, cos_angles], dim=-1)
        # reshape gamma_x to (N, LD)
        gamma_x = gamma_x.reshape(angles.shape[0], -1)
        results.append(gamma_x)

    return torch.cat(results, dim=-1)

"""

This function returns chunks of the ray points and directions to avoid memory errors with the neural network. It also applies positional encoding to the input points and directions before dividing them into chunks, as well as normalizing and populating the directions.

"""

##### TODO 2.3 BEGIN #####
ray_directions_normalized = ray_directions/torch.linalg.norm(ray_directions, axis=2, keepdim=True) # (H,W,N,3)
all_ray_dirs = (ray_directions_normalized.unsqueeze(-2)*torch.ones_like(ray_points)) # (H,W,N,3)
pos_enc_pt_output = positional_encoding(torch.reshape(ray_points, (-1,3)), num_x_frequencies, True)
pos_enc_dir_output = positional_encoding(torch.reshape(all_ray_dirs, (-1,3)), num_d_frequencies, True)
ray_points_batches = get_chunks(pos_enc_pt_output)
ray_directions_batches = get_chunks(pos_enc_dir_output)
##### TODO 2.3 END #####
return ray_points_batches, ray_directions_batches

```

2.4 Compute the compositing weights of samples on camera ray and then complete the volumetric rendering procedure to reconstruct a whole RGB image from the sampled points and the outputs of the neural network.

[10]: `def volumetric_rendering(rgb, s, depth_points):`

```

"""

Differentiably renders a radiance field, given the origin of each ray in the "bundle", and the sampled depth values along them.

```

Args:

```

rgb: RGB color at each query location (X, Y, Z). Shape: (height, width, samples, 3).
sigma: Volume density at each query location (X, Y, Z). Shape: (height, width, samples).
depth_points: Sampled depth values along each ray. Shape: (height, width, samples).

Returns:
rec_image: The reconstructed image after applying the volumetric rendering to every pixel.
Shape: (height, width, 3)
"""

#####
# TODO 2.4 BEGIN #####
H, W, N = depth_points.shape
s = nn.ReLU()(s)
del_i = depth_points[:, :, 1:] - depth_points[:, :, :-1] # (H, W, samples-1)
del_i = del_i.to(device)
s = s.to(device)
T_i = torch.cumprod(torch.exp(-del_i*s[:, :, :-1]), axis=2)
T_i = torch.cat([torch.ones(H, W, 1).to(device), T_i], dim=-1)

del_i = torch.cat([del_i, 1e9*torch.ones(H, W, 1).to(device)], dim=-1) # (H, W, samples)

rec_image = torch.sum(T_i.unsqueeze(-1)*(1.-torch.exp(-del_i*s)).unsqueeze(-1)*rgb.to(device), axis=2)

return rec_image

```

2.5 Combine everything together. Given the pose position of a camera, compute the camera rays and sample the 3D points along these rays. Divide those points into batches and feed them to the neural network. Concatenate them and use them for the volumetric rendering to reconstructed the final image.

```
[11]: def one_forward_pass(height, width, intrinsics, pose, near, far, samples, model, num_x_frequencies, num_d_frequencies):

#####
# TODO 2.5 BEGIN #####
#compute all the rays from the image
Rcw = pose[:3, :3].T
Tcw = pose[:3, -1]
ray_origins, ray_directions = get_rays(height, width, intrinsics, Rcw, Tcw)

#sample the points from the rays
```

```

    ray_points, depth_points = stratified_sampling(ray_origins, ray_directions, near, far, samples)

#divide data into batches to avoid memory errors
ray_points_batches, ray_directions_batches = get_batches(ray_points, ray_directions, num_x_frequencies, num_d_frequencies)
# print(len(ray_points_batches))
# raise

#forward pass the batches and concatenate the outputs at the end
fin_rgb = []
fin_sig = []
# for i in tqdm(range(len(ray_points_batches))):
for i in range(len(ray_points_batches)):
    rgb, sigma = model(ray_points_batches[i], ray_directions_batches[i])
    fin_rgb.append(rgb)
    fin_sig.append(sigma)

# a = torch.vstack(fin_rgb)
a1 = torch.cat(fin_rgb, dim=0)
# print(torch.sum(a.isclose(a1, 1e-2))- a.shape[0]*a.shape[1])
b = torch.vstack(fin_sig)

# Apply volumetric rendering to obtain the reconstructed image
rec_image = volumetric_rendering(a1.reshape(height, width, samples, 3), b,
reshape(height, width, samples), depth_points)
# print("DONE 1 forward pass")
##### TODO 2.5 END #####
return rec_image

```

If you manage to pass the autograder for all the previous functions, then it is time to train a NeRF! We provide the hyperparameters for you, we initialize the NeRF model and its weights, and we define a couple lists that will be needed to store results.

```
[12]: num_x_frequencies = 10
num_d_frequencies = 4
learning_rate = 5e-4
iterations = 3000
samples = 64
display = 25
near = 0.667
far = 2
```

```

model = nerf_model(num_x_frequencies=num_x_frequencies, num_d_frequencies=num_d_frequencies).
        to(device)

def weights_init(m):
    if isinstance(m, torch.nn.Linear):
        torch.nn.init.xavier_uniform_(m.weight)
model.apply(weights_init)

optimizer = torch.optim.Adam(model.parameters(), lr=learning_rate)
criterion = nn.MSELoss()

psnrs = []
iternums = []

t = time.time()
t0 = time.time()

```

```

[13]: from matplotlib import image
for i in range(iterations+1):

    ##### TODO 2.6 BEGIN #####
    #choose randomly a picture for the forward pass
    # torch.cuda.empty_cache()
    pic_idx = np.random.randint(images.shape[0])
    random_image, random_pose = images[pic_idx], poses[pic_idx]

    # Run one iteration of NeRF and get the rendered RGB image.
    optimizer.zero_grad()
    rec_img = one_forward_pass(height, width, intrinsics, random_pose, near,
        far, samples, model, num_x_frequencies, num_d_frequencies)

    # Compute mean-squared error between the predicted and target images.
    # Backprop!
    loss = criterion(rec_img, random_image)
    # print("Loss: ", loss)
    loss.backward()

    optimizer.step()

    ##### TODO 2.6 END #####
    # Display images/plots/stats
    if i % display == 0:

```

```

    with torch.no_grad():
        ##### TODO 2.6 BEGIN #####
        # Render the held-out view
        test_rec_image = one_forward_pass(height, width, intrinsics,
        ↵test_pose, near, far, samples, model, num_x_frequencies, num_d_frequencies)
        loss = nn.MSELoss()(test_rec_image, test_image)

        #calculate the loss and the psnr between the original test image and
        ↵the reconstructed one.
        psnr = 10*torch.log10(1/loss)

        ##### TODO 2.6 END #####
        #####
        print("Iteration %d " % i, "Loss: %.4f " % loss.item(), "PSNR: %.2f " %
        ↵psnr.item(), \
            "Time: %.2f secs per iter, " % ((time.time() - t) / display), \
        ↵"%.2f mins in total" % ((time.time() - t0)/60))

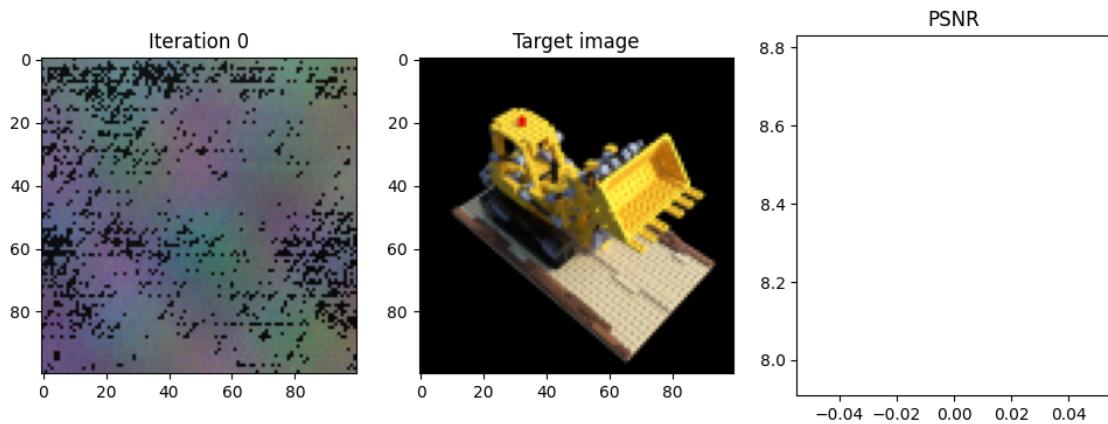
        t = time.time()
        psnrs.append(psnr.item())
        iternums.append(i)

        plt.figure(figsize=(16, 4))
        plt.subplot(141)
        plt.imshow(test_rec_image.detach().cpu().numpy())
        plt.title(f"Iteration {i}")
        plt.subplot(142)
        plt.imshow(test_image.detach().cpu().numpy())
        plt.title("Target image")
        plt.subplot(143)
        plt.plot(iternums, psnrs)
        plt.title("PSNR")
        plt.show()

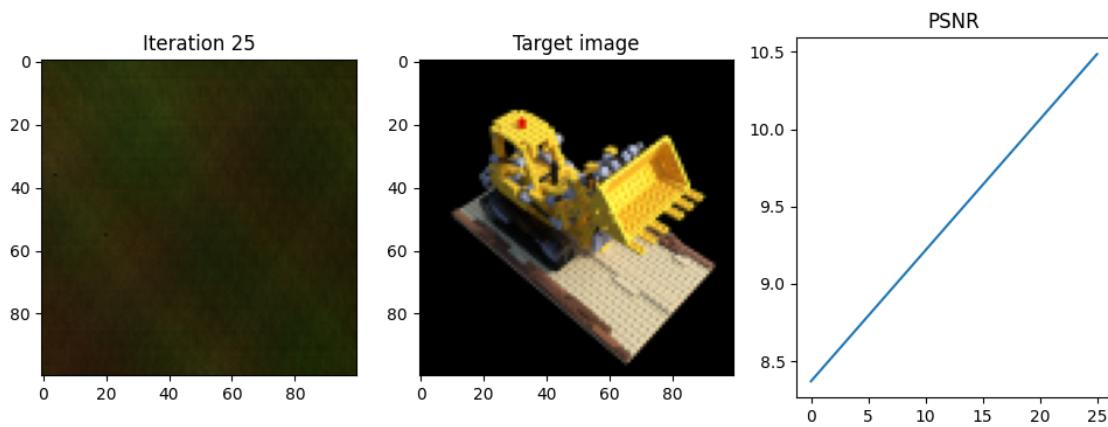
print('Done!')

```

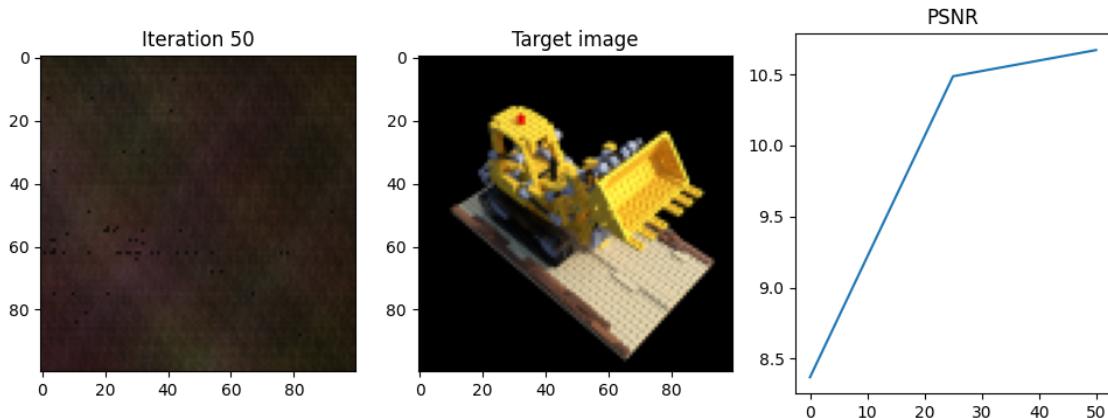
Iteration 0 Loss: 0.1456 PSNR: 8.37 Time: 0.70 secs per iter, 0.29 mins in total



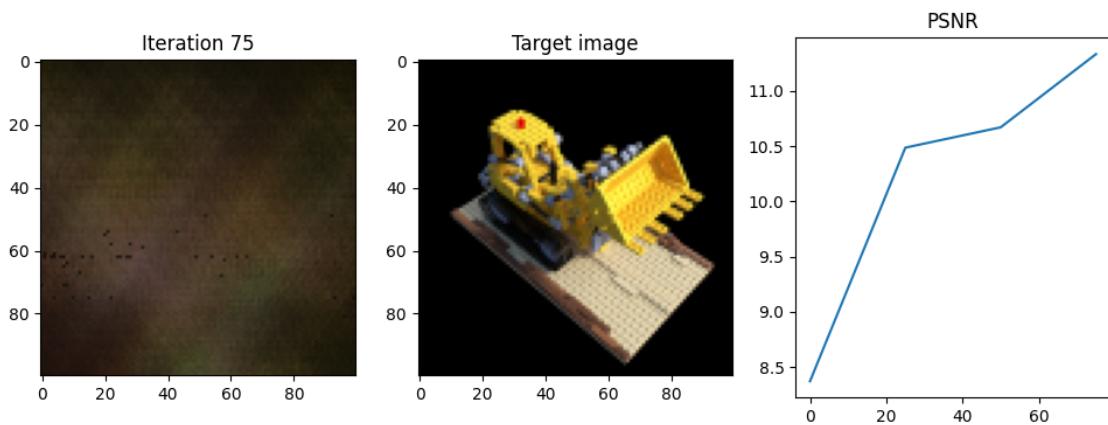
Iteration 25 Loss: 0.0894 PSNR: 10.49 Time: 0.79 secs per iter, 0.62 mins in total



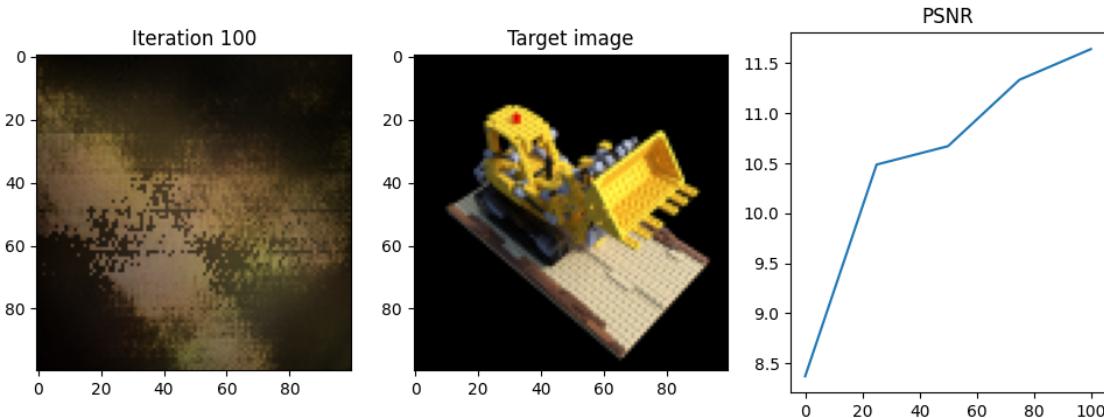
Iteration 50 Loss: 0.0857 PSNR: 10.67 Time: 0.80 secs per iter, 0.95 mins in total



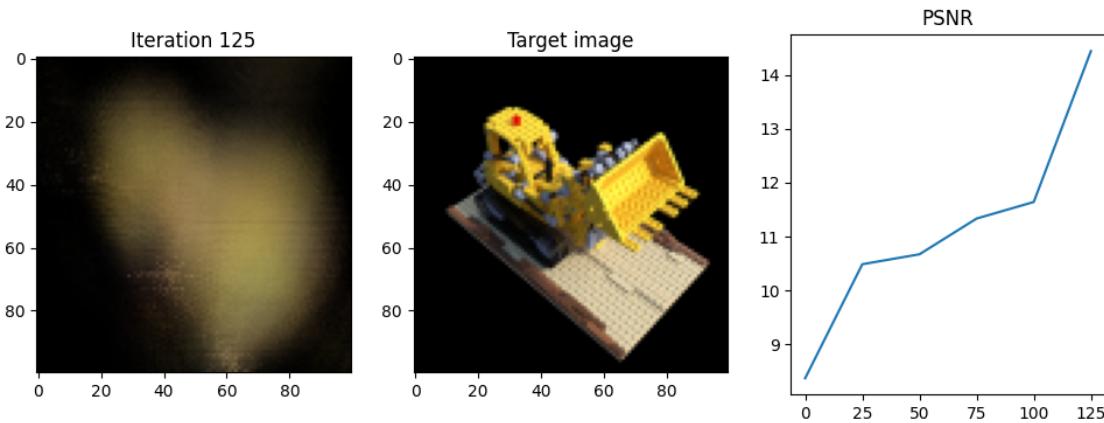
Iteration 75 Loss: 0.0735 PSNR: 11.33 Time: 0.80 secs per iter, 1.29 mins in total



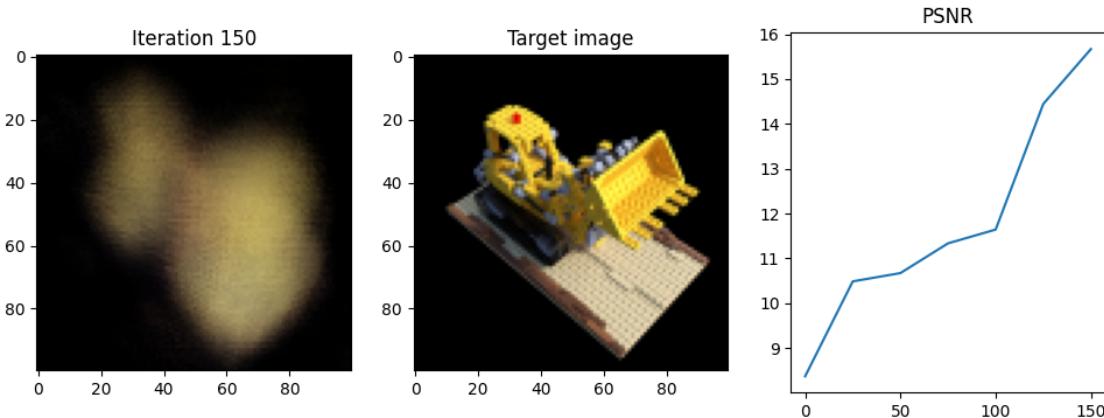
Iteration 100 Loss: 0.0685 PSNR: 11.64 Time: 0.82 secs per iter, 1.63 mins in total



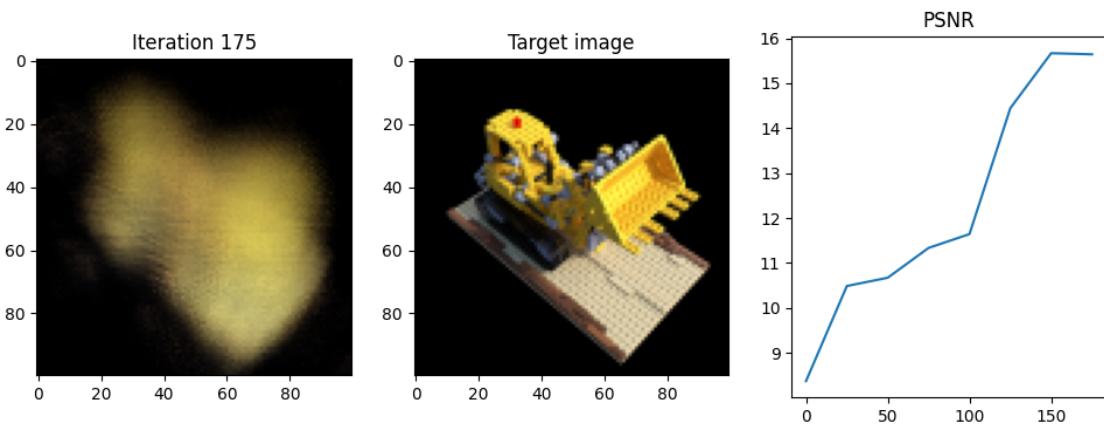
Iteration 125 Loss: 0.0359 PSNR: 14.44 Time: 0.82 secs per iter, 1.97 mins in total



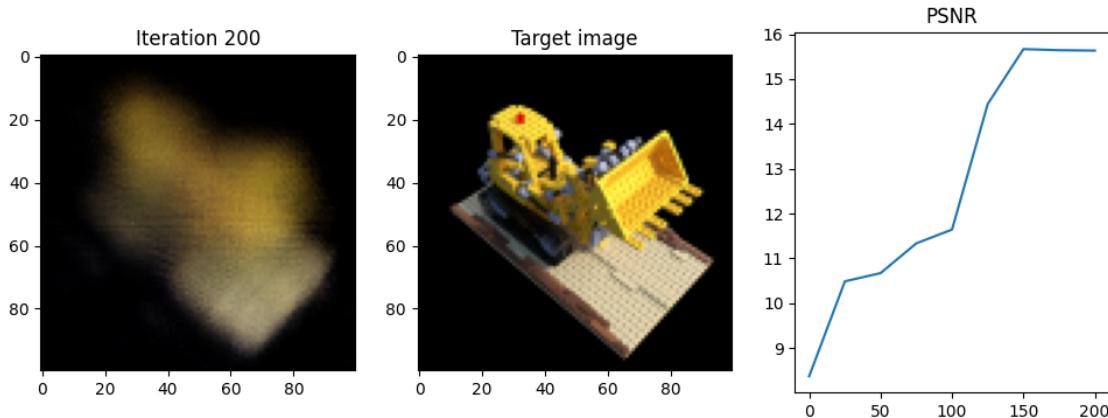
Iteration 150 Loss: 0.0271 PSNR: 15.67 Time: 0.81 secs per iter, 2.31 mins in total



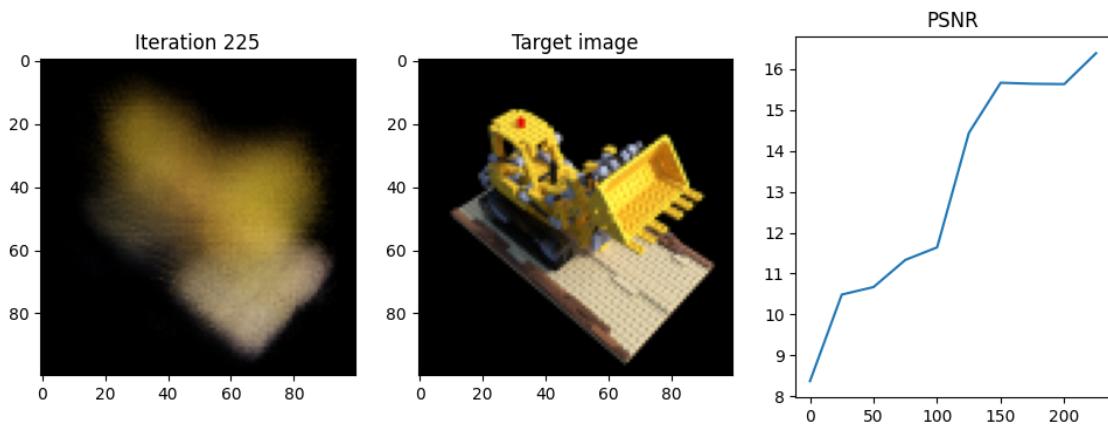
Iteration 175 Loss: 0.0273 PSNR: 15.64 Time: 0.81 secs per iter, 2.64 mins in total



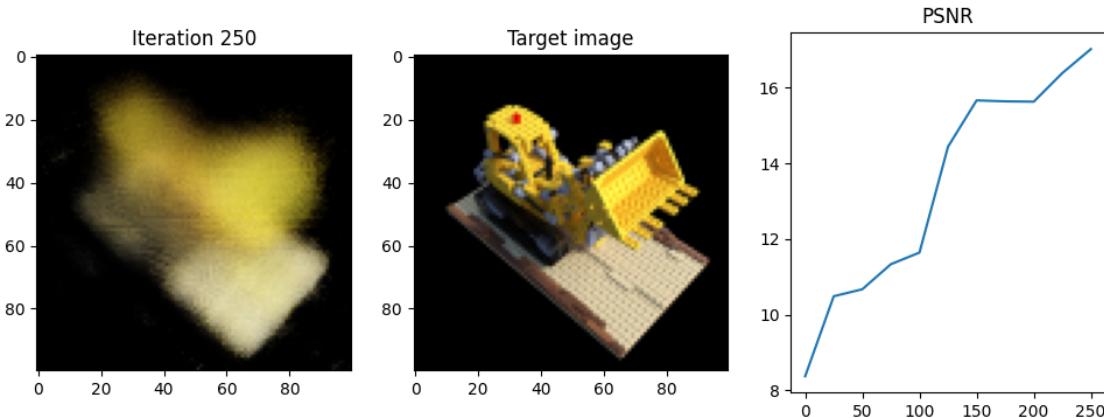
Iteration 200 Loss: 0.0273 PSNR: 15.63 Time: 0.82 secs per iter, 2.99 mins in total



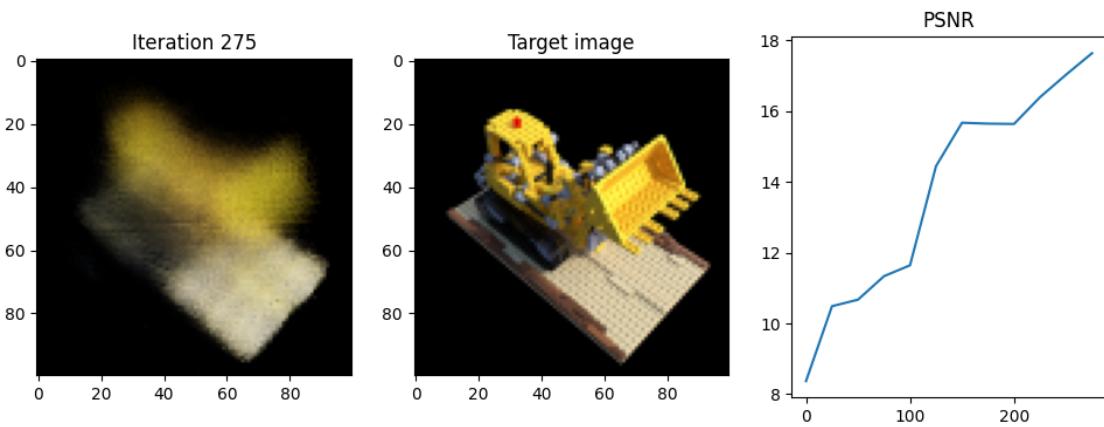
Iteration 225 Loss: 0.0230 PSNR: 16.39 Time: 0.81 secs per iter, 3.32 mins in total



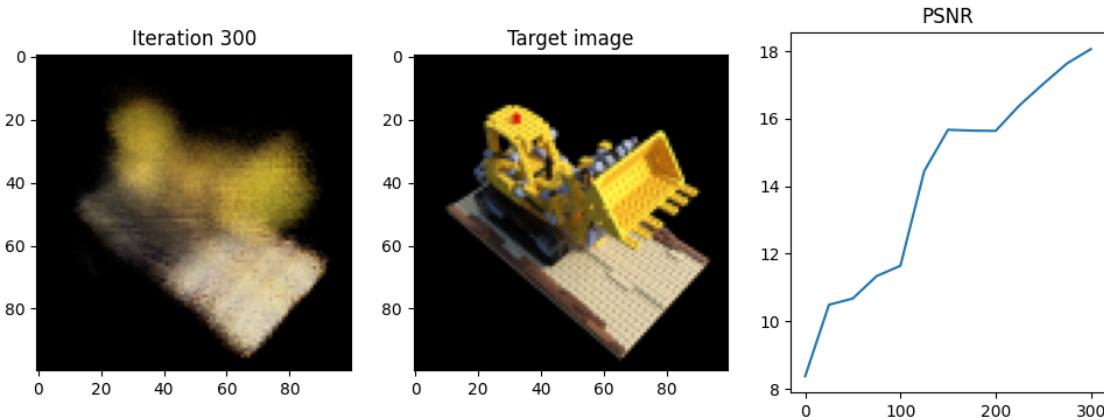
Iteration 250 Loss: 0.0198 PSNR: 17.03 Time: 0.82 secs per iter, 3.67 mins in total



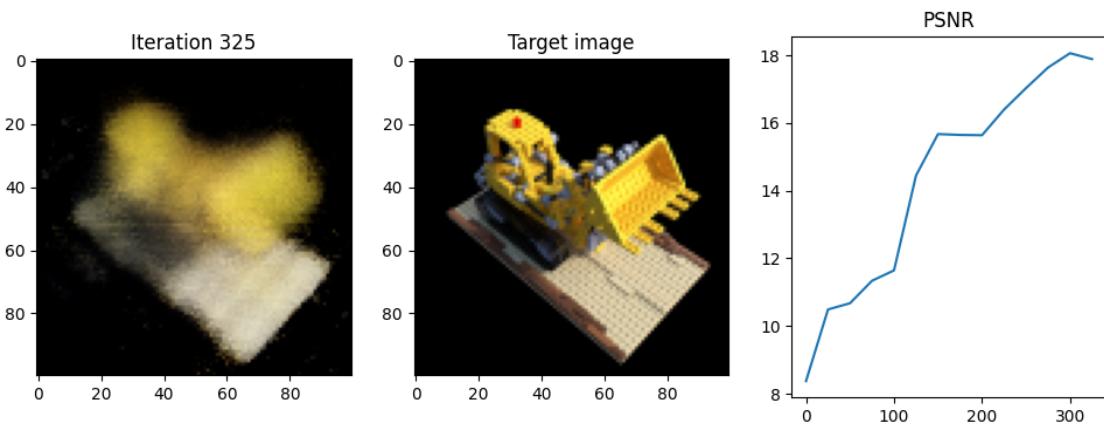
Iteration 275 Loss: 0.0172 PSNR: 17.63 Time: 0.82 secs per iter, 4.01 mins in total



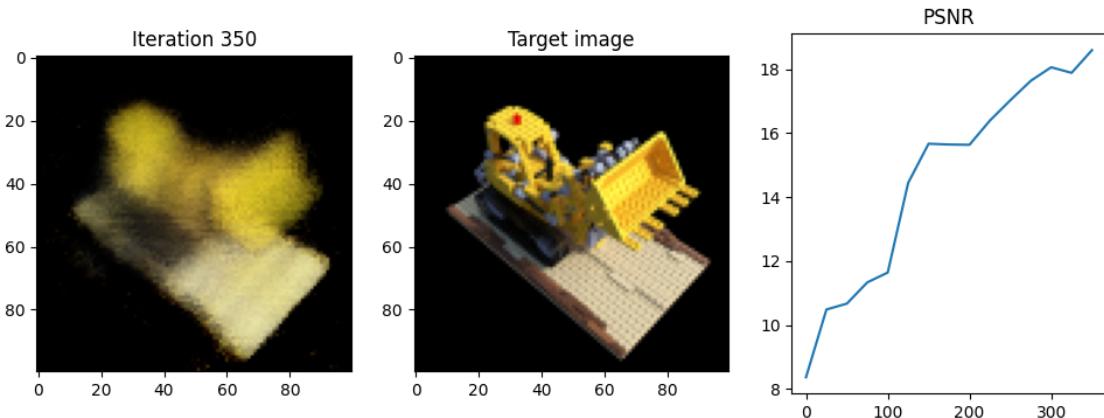
Iteration 300 Loss: 0.0156 PSNR: 18.06 Time: 0.83 secs per iter, 4.35 mins in total



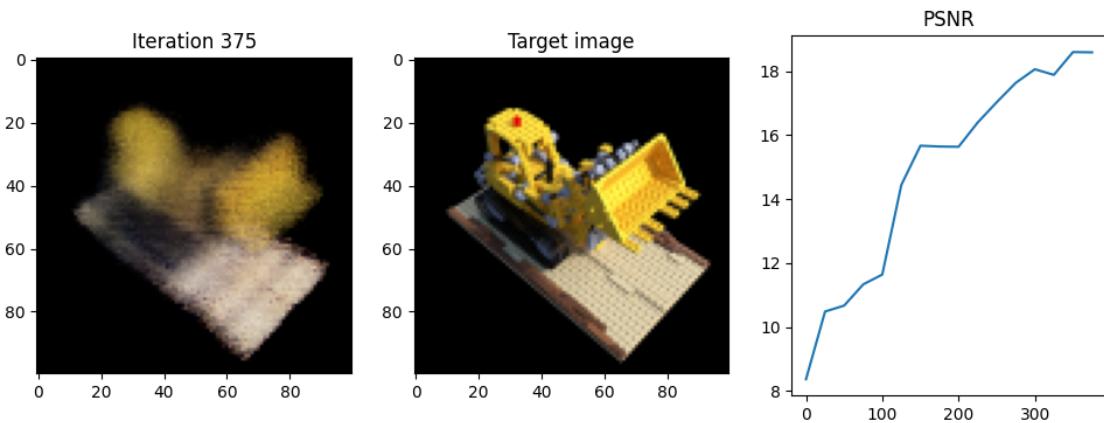
Iteration 325 Loss: 0.0163 PSNR: 17.88 Time: 0.83 secs per iter, 4.70 mins in total



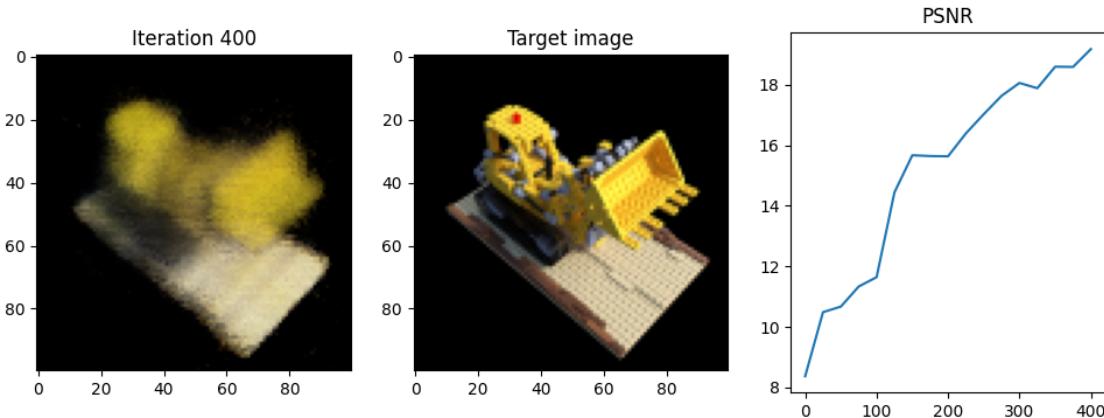
Iteration 350 Loss: 0.0138 PSNR: 18.60 Time: 0.82 secs per iter, 5.04 mins in total



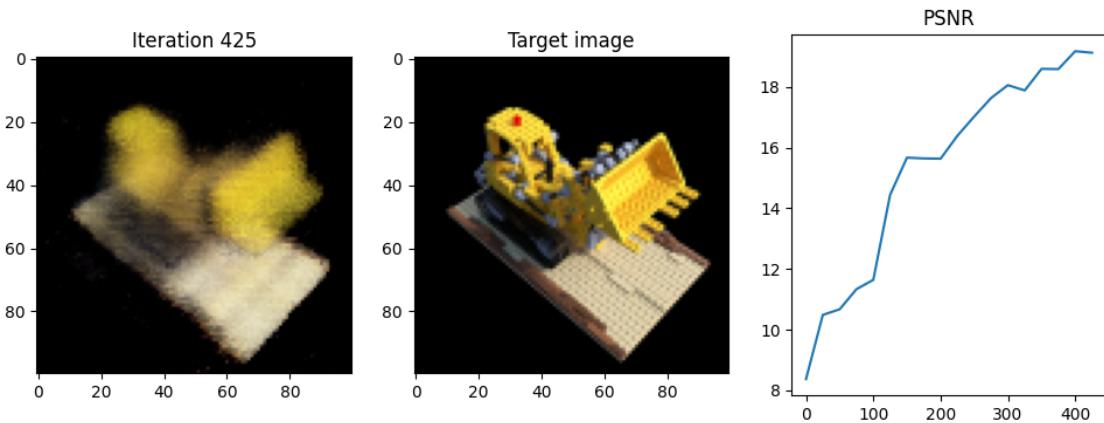
Iteration 375 Loss: 0.0138 PSNR: 18.59 Time: 0.82 secs per iter, 5.38 mins in total



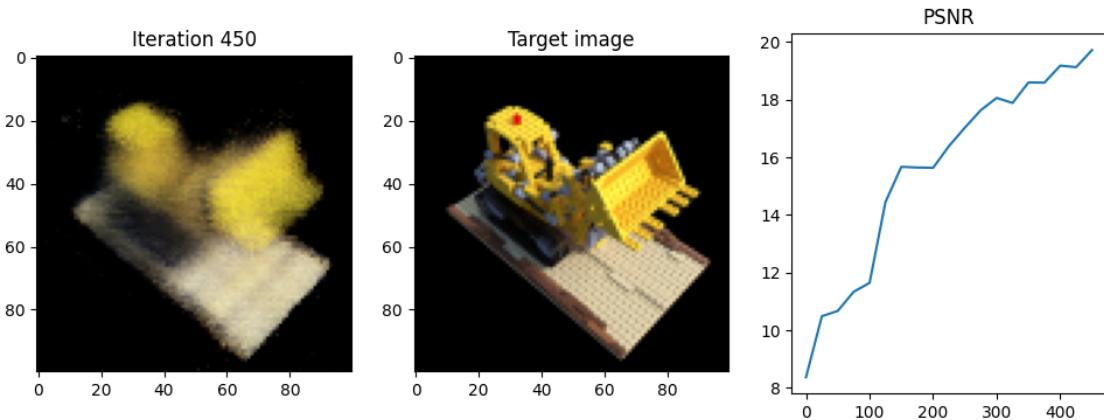
Iteration 400 Loss: 0.0121 PSNR: 19.18 Time: 0.82 secs per iter, 5.72 mins in total



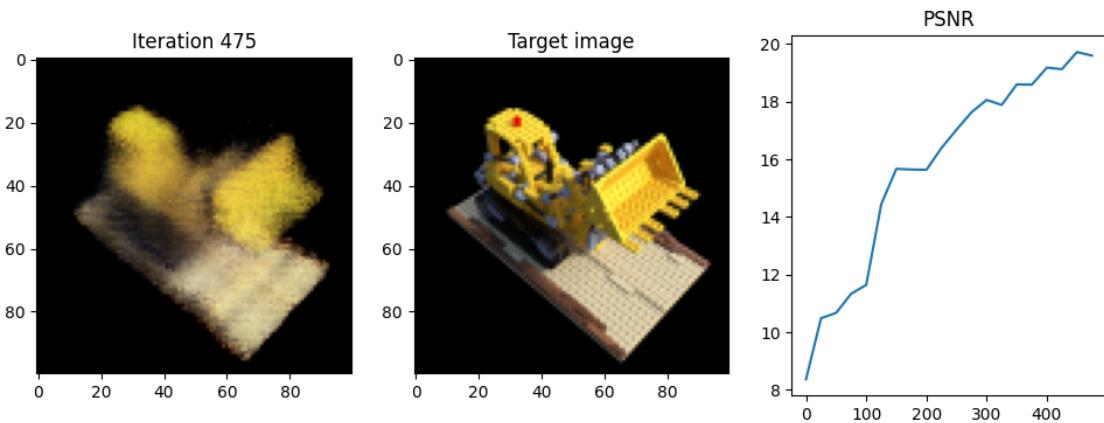
Iteration 425 Loss: 0.0122 PSNR: 19.13 Time: 0.82 secs per iter, 6.06 mins in total



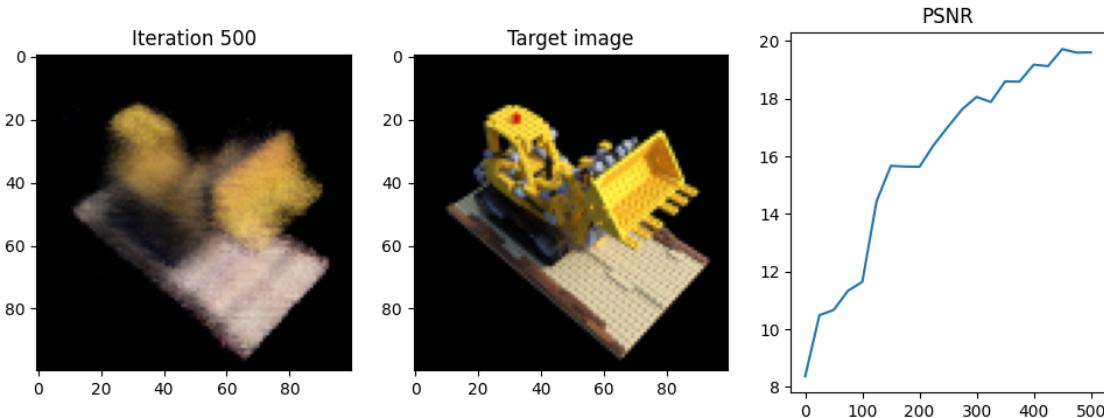
Iteration 450 Loss: 0.0107 PSNR: 19.72 Time: 0.81 secs per iter, 6.40 mins in total



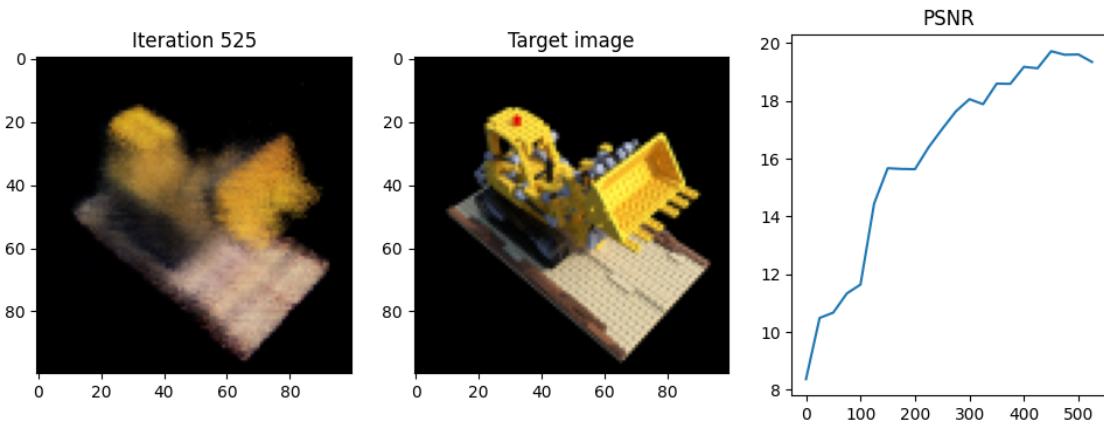
Iteration 475 Loss: 0.0110 PSNR: 19.59 Time: 0.82 secs per iter, 6.74 mins in total



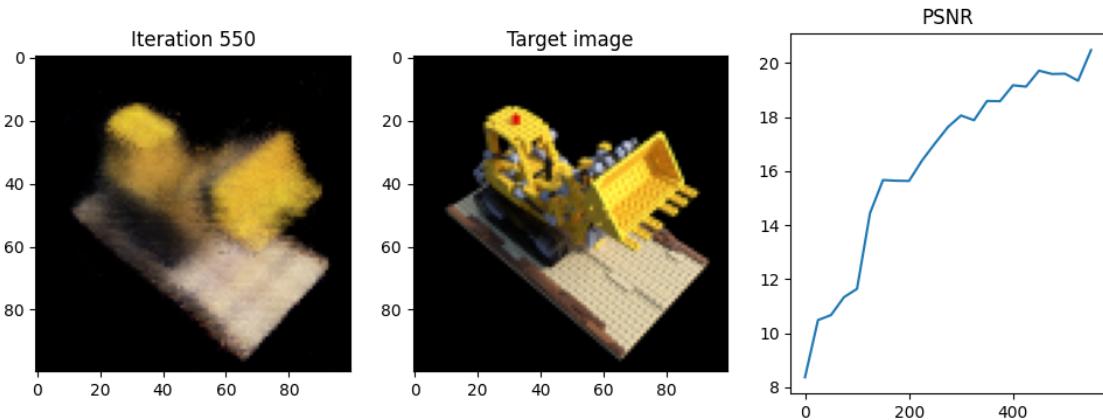
Iteration 500 Loss: 0.0110 PSNR: 19.61 Time: 0.81 secs per iter, 7.08 mins in total



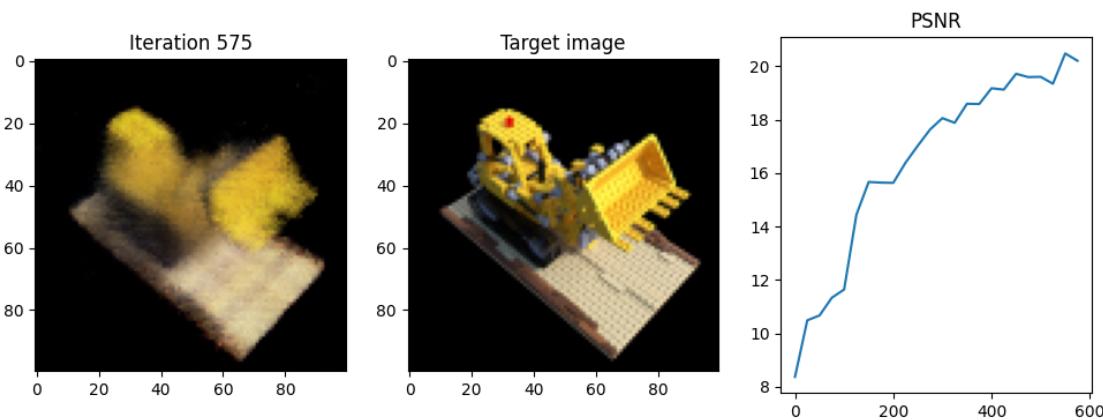
Iteration 525 Loss: 0.0116 PSNR: 19.35 Time: 0.82 secs per iter, 7.43 mins in total



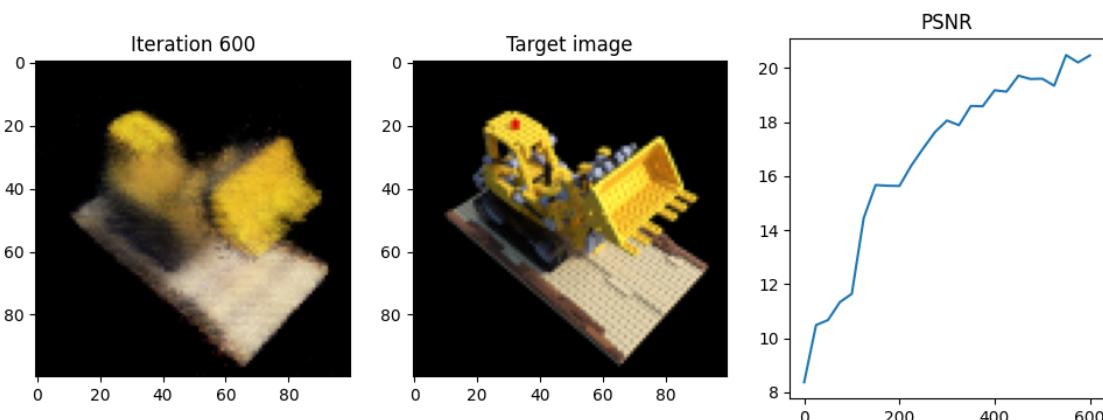
Iteration 550 Loss: 0.0090 PSNR: 20.48 Time: 0.82 secs per iter, 7.77 mins in total



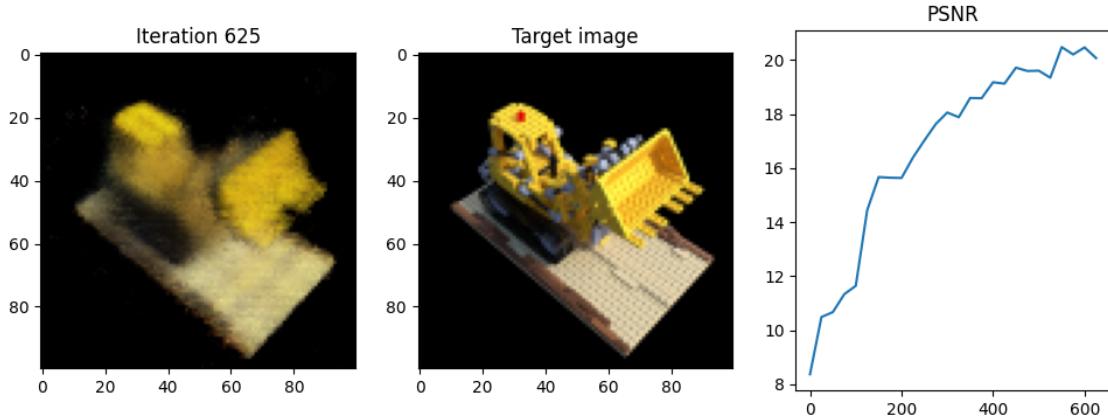
Iteration 575 Loss: 0.0095 PSNR: 20.21 Time: 0.82 secs per iter, 8.11 mins in total



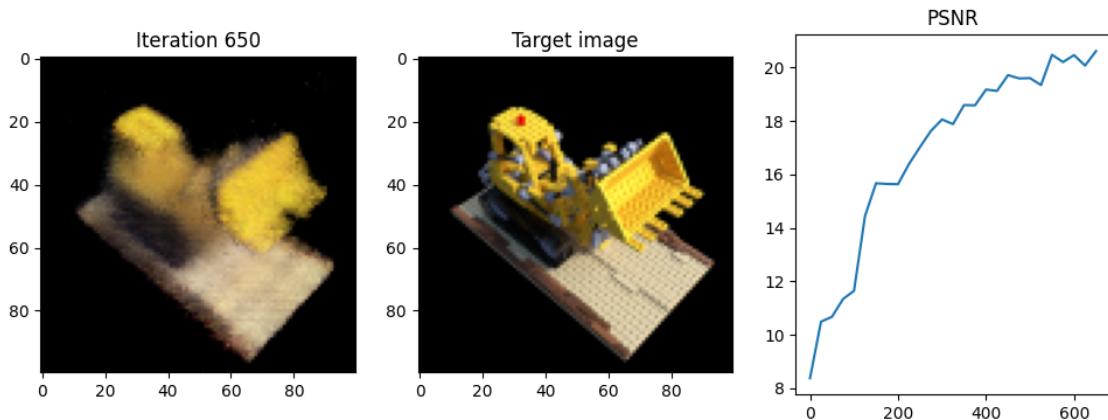
Iteration 600 Loss: 0.0090 PSNR: 20.47 Time: 0.81 secs per iter, 8.45 mins in total



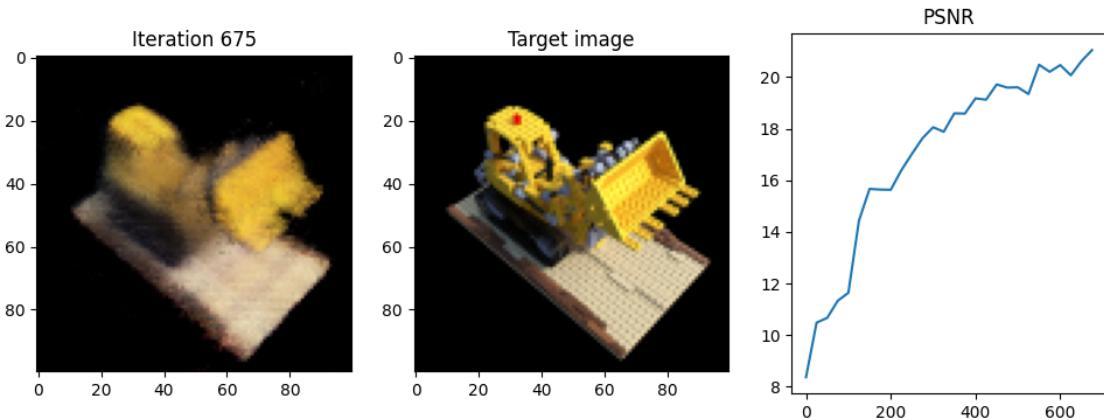
Iteration 625 Loss: 0.0098 PSNR: 20.07 Time: 0.81 secs per iter, 8.79 mins in total



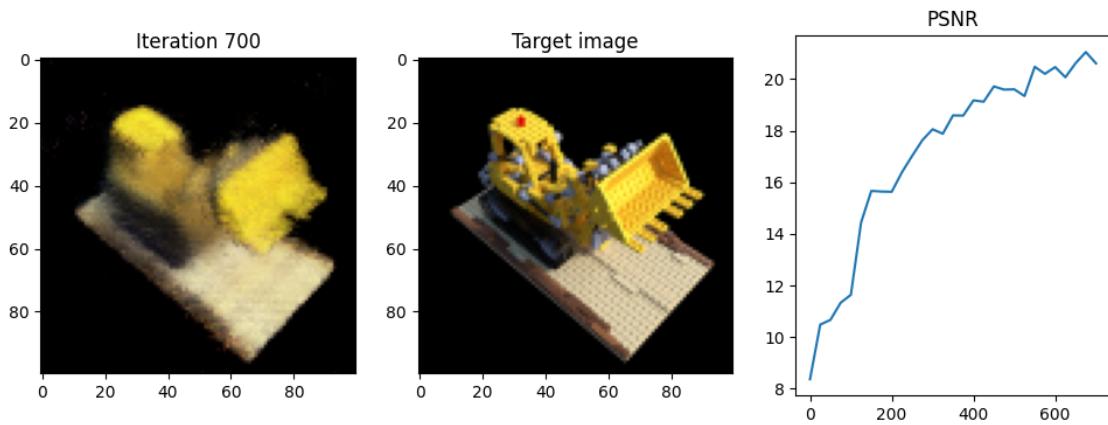
Iteration 650 Loss: 0.0087 PSNR: 20.62 Time: 0.81 secs per iter, 9.12 mins in total



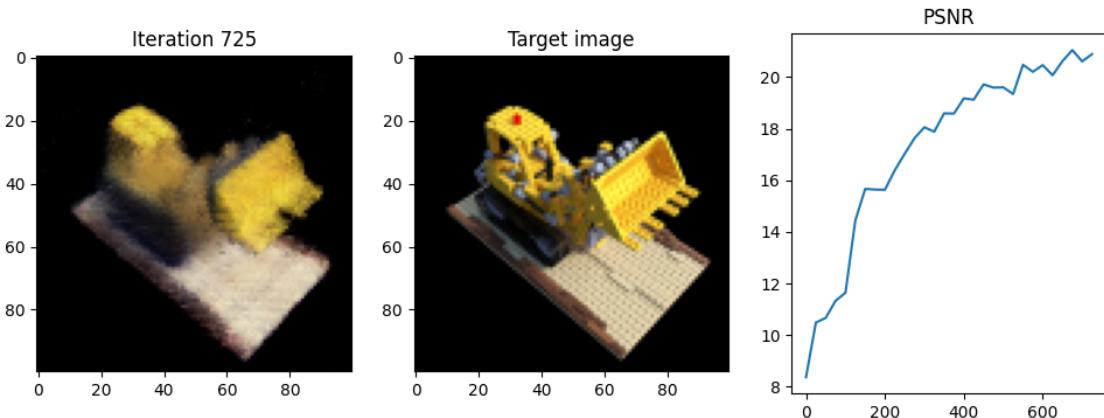
Iteration 675 Loss: 0.0079 PSNR: 21.05 Time: 0.81 secs per iter, 9.46 mins in total



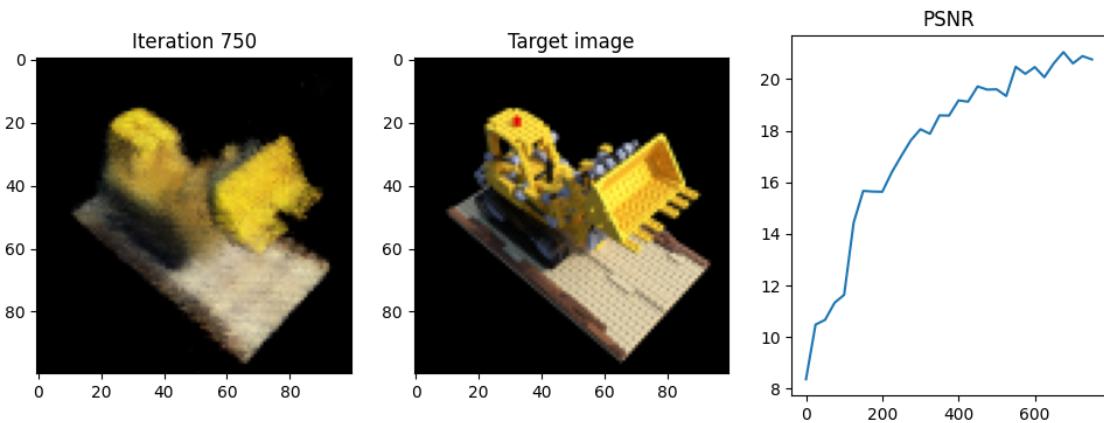
Iteration 700 Loss: 0.0087 PSNR: 20.61 Time: 0.81 secs per iter, 9.80 mins in total



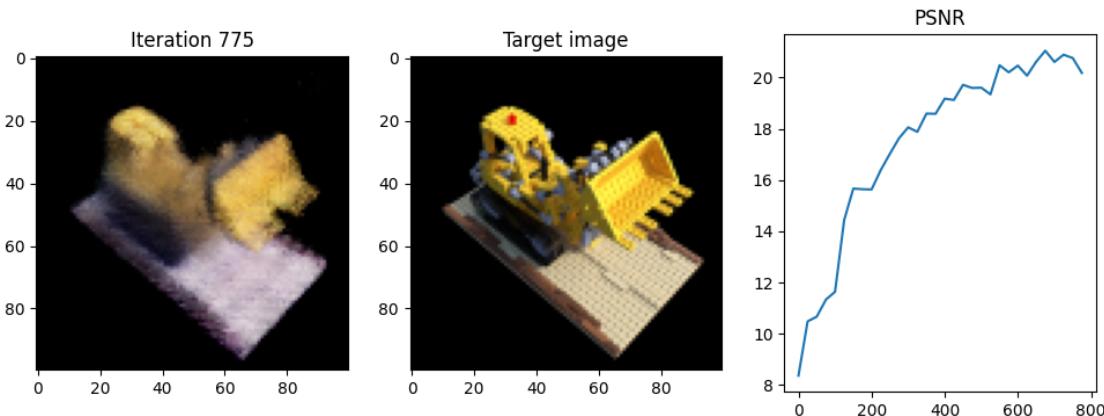
Iteration 725 Loss: 0.0081 PSNR: 20.89 Time: 0.82 secs per iter, 10.14 mins in total



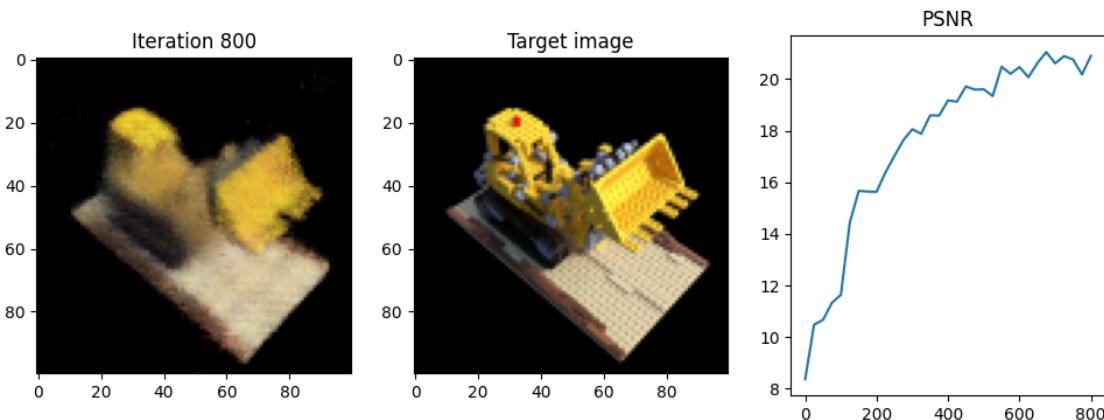
Iteration 750 Loss: 0.0084 PSNR: 20.76 Time: 0.81 secs per iter, 10.48 mins in total



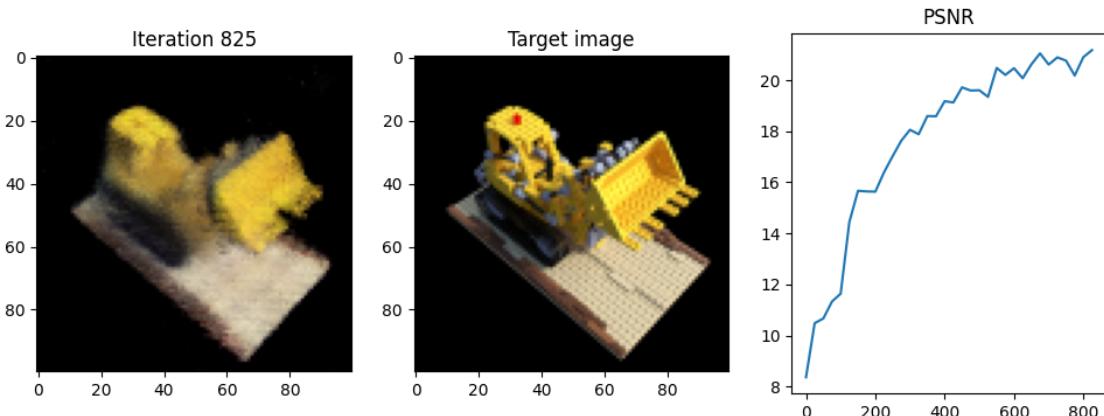
Iteration 775 Loss: 0.0096 PSNR: 20.18 Time: 0.81 secs per iter, 10.82 mins in total



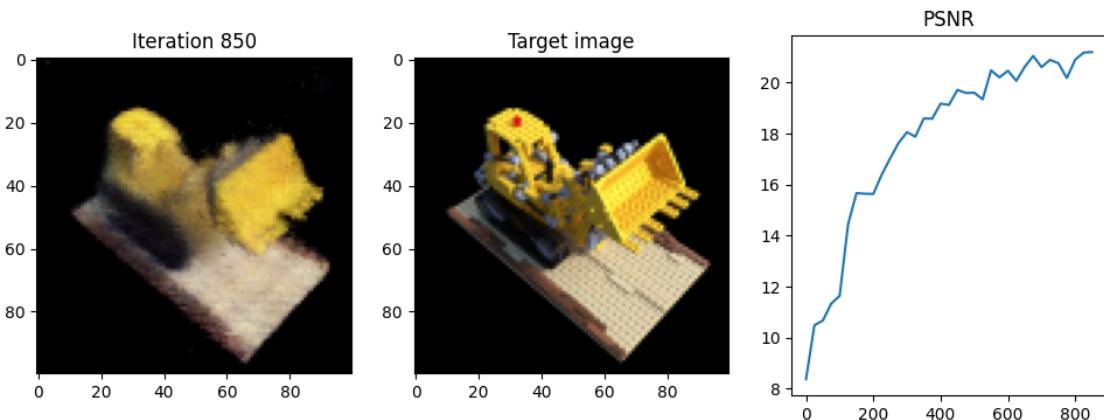
Iteration 800 Loss: 0.0081 PSNR: 20.90 Time: 0.81 secs per iter, 11.16 mins in total



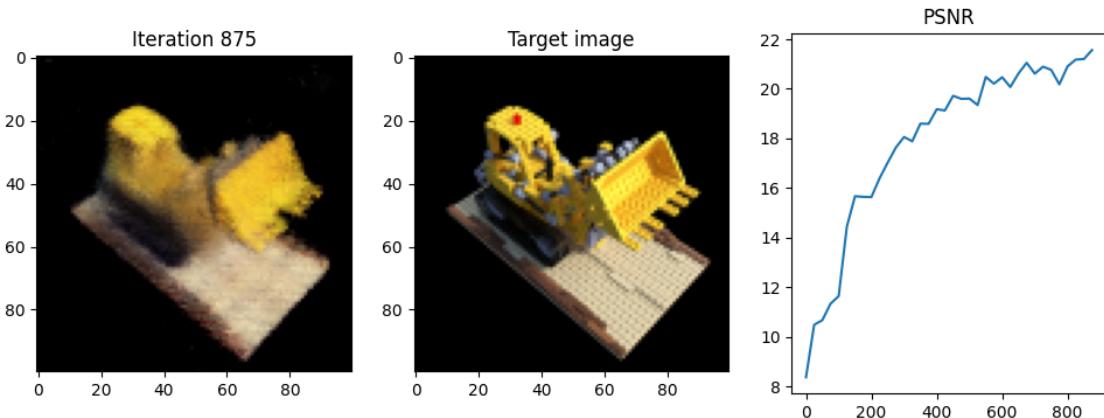
Iteration 825 Loss: 0.0076 PSNR: 21.18 Time: 0.81 secs per iter, 11.50 mins in total



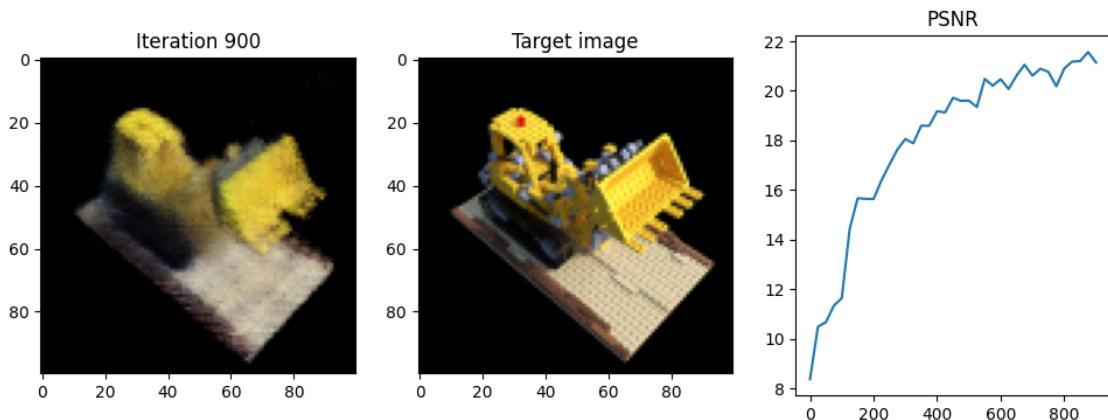
Iteration 850 Loss: 0.0076 PSNR: 21.20 Time: 0.81 secs per iter, 11.83 mins in total



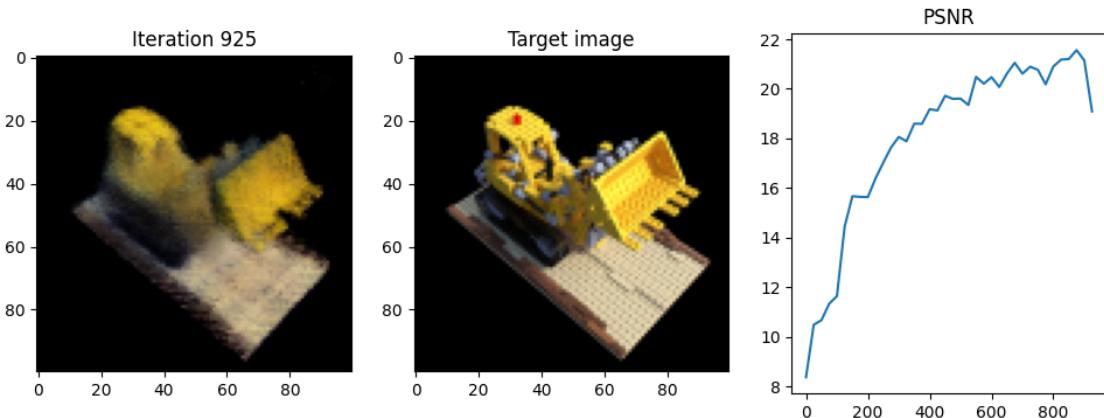
Iteration 875 Loss: 0.0070 PSNR: 21.56 Time: 0.81 secs per iter, 12.17 mins in total



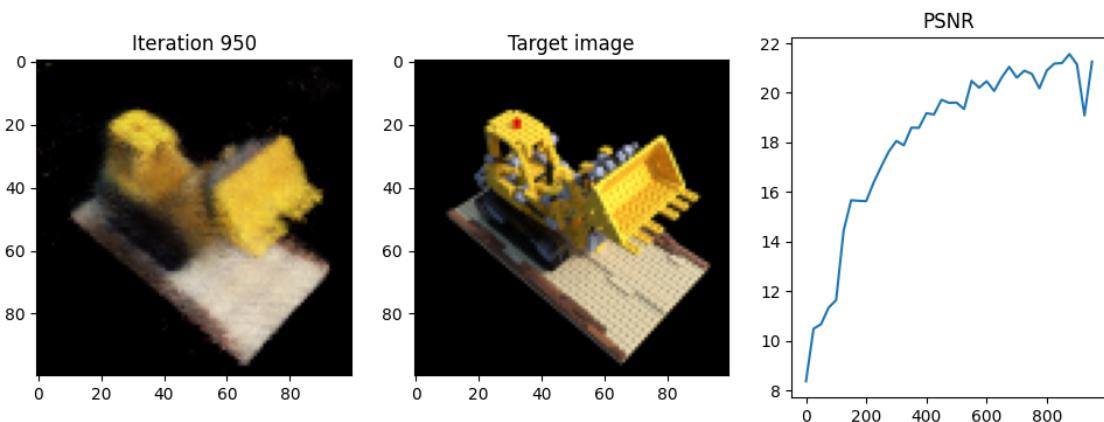
Iteration 900 Loss: 0.0077 PSNR: 21.14 Time: 0.82 secs per iter, 12.51 mins in total



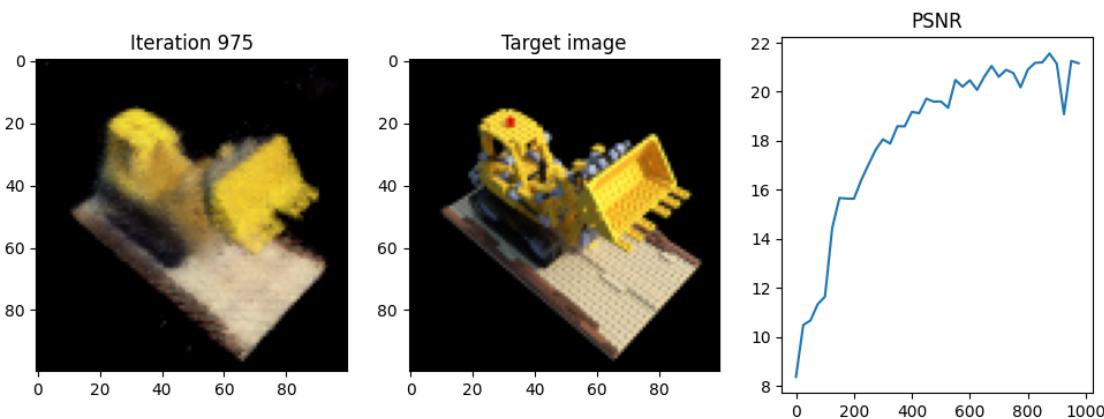
Iteration 925 Loss: 0.0123 PSNR: 19.09 Time: 0.81 secs per iter, 12.85 mins in total



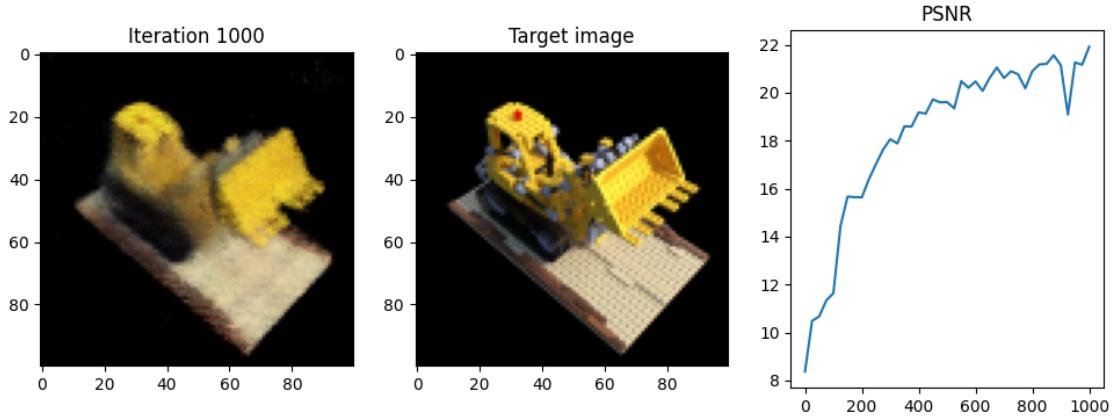
Iteration 950 Loss: 0.0075 PSNR: 21.26 Time: 0.81 secs per iter, 13.19 mins in total



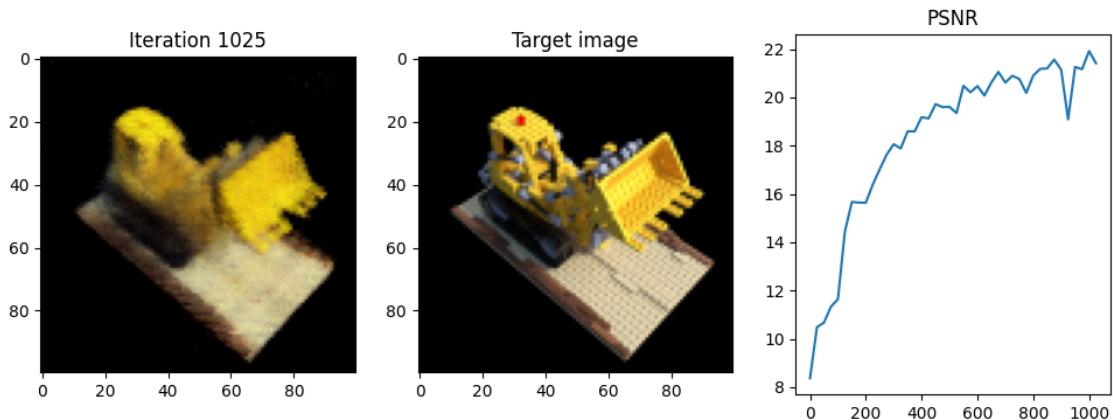
Iteration 975 Loss: 0.0077 PSNR: 21.16 Time: 0.82 secs per iter, 13.53 mins in total



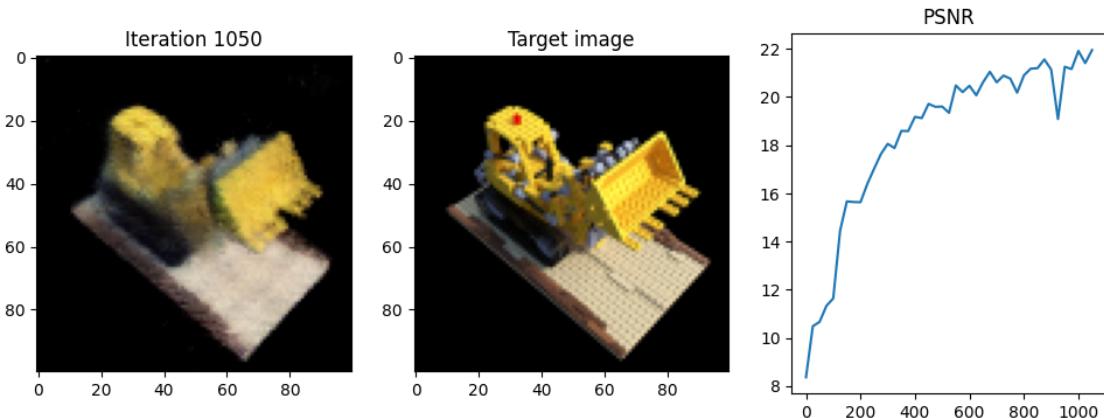
Iteration 1000 Loss: 0.0064 PSNR: 21.91 Time: 0.81 secs per iter, 13.87 mins in total



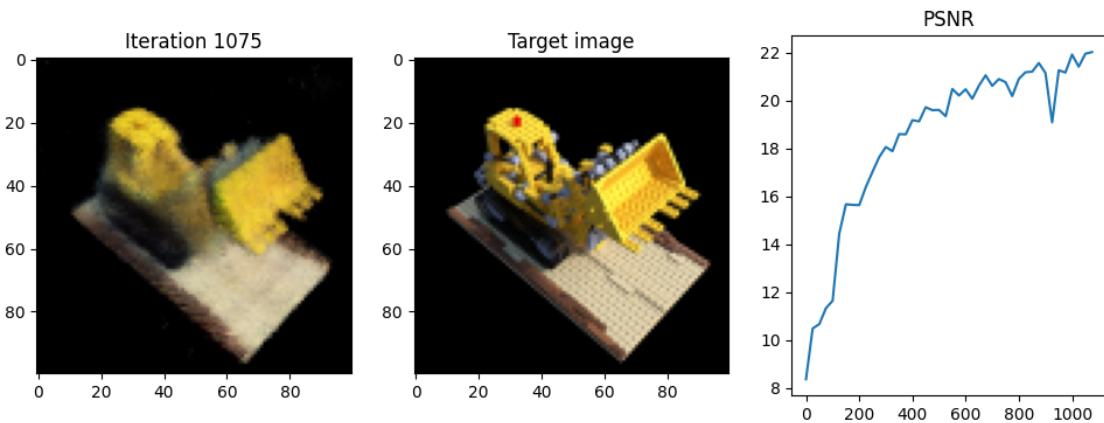
Iteration 1025 Loss: 0.0072 PSNR: 21.41 Time: 0.81 secs per iter, 14.20 mins in total



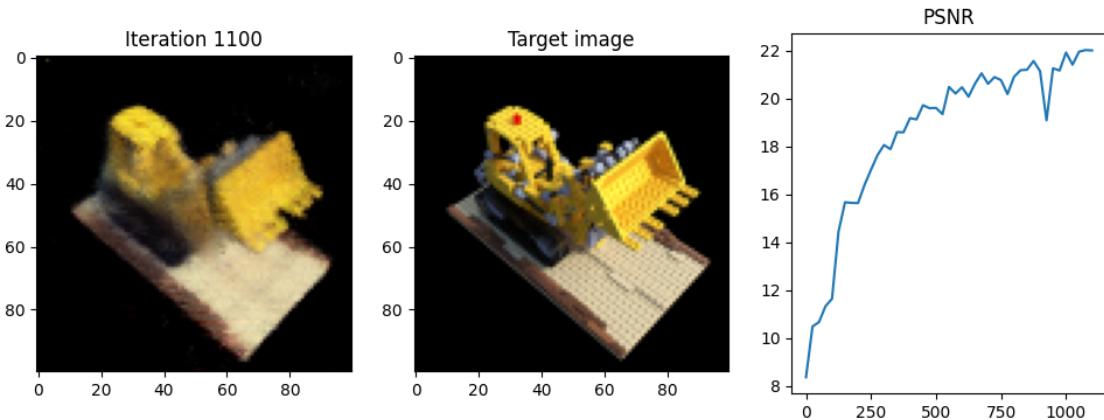
Iteration 1050 Loss: 0.0064 PSNR: 21.95 Time: 0.82 secs per iter, 14.55 mins in total



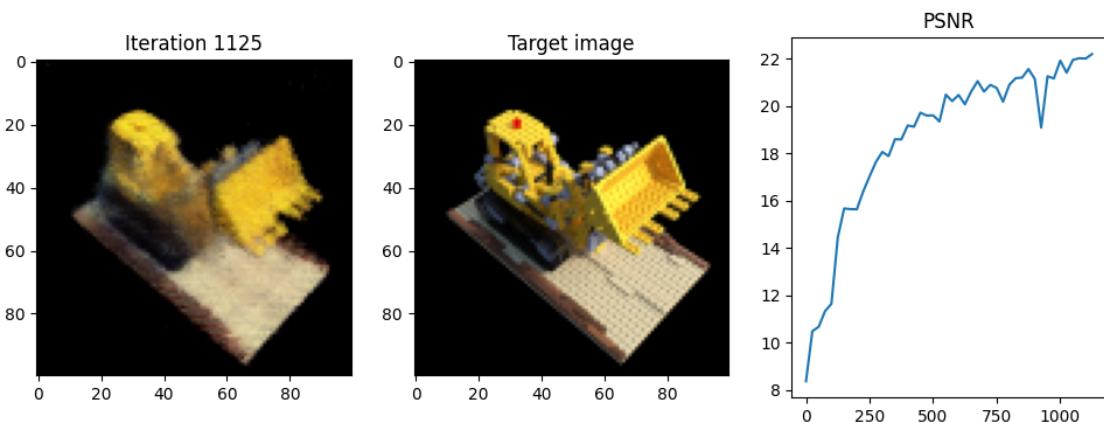
Iteration 1075 Loss: 0.0063 PSNR: 22.02 Time: 0.81 secs per iter, 14.88 mins in total



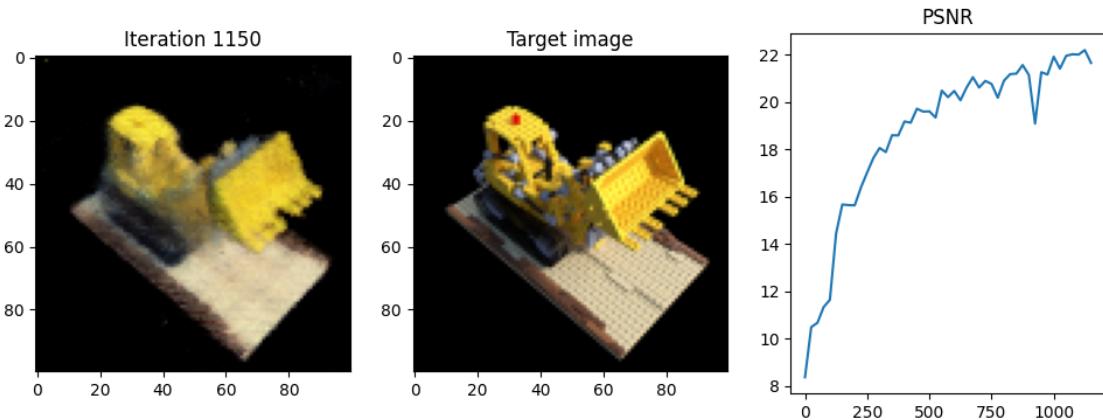
Iteration 1100 Loss: 0.0063 PSNR: 22.00 Time: 0.82 secs per iter, 15.22 mins in total



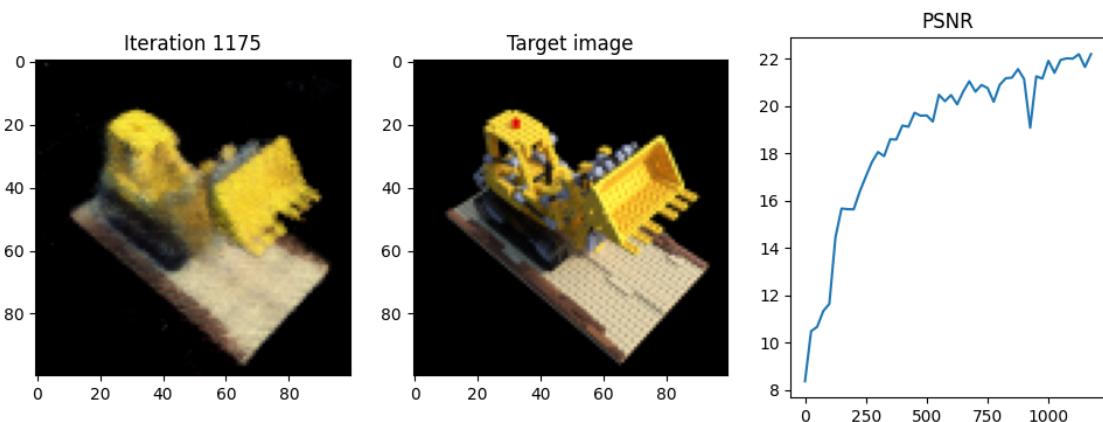
Iteration 1125 Loss: 0.0060 PSNR: 22.19 Time: 0.82 secs per iter, 15.56 mins in total



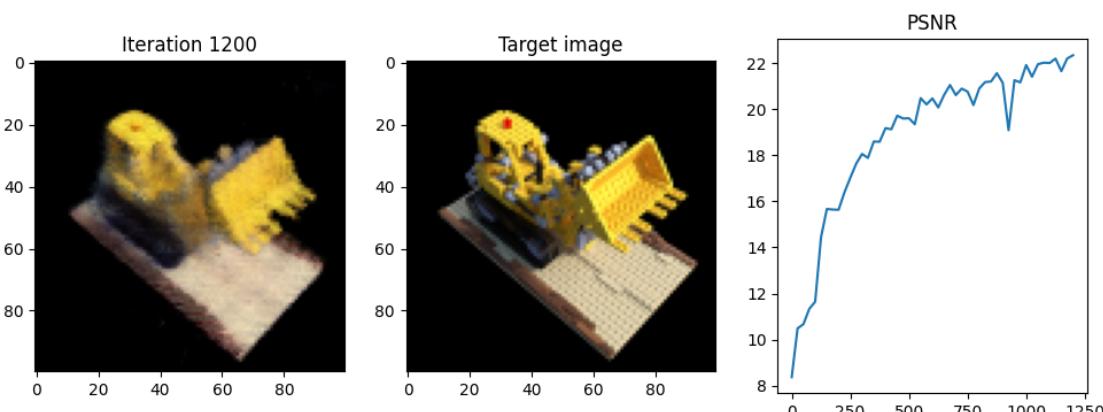
Iteration 1150 Loss: 0.0068 PSNR: 21.65 Time: 0.81 secs per iter, 15.90 mins in total



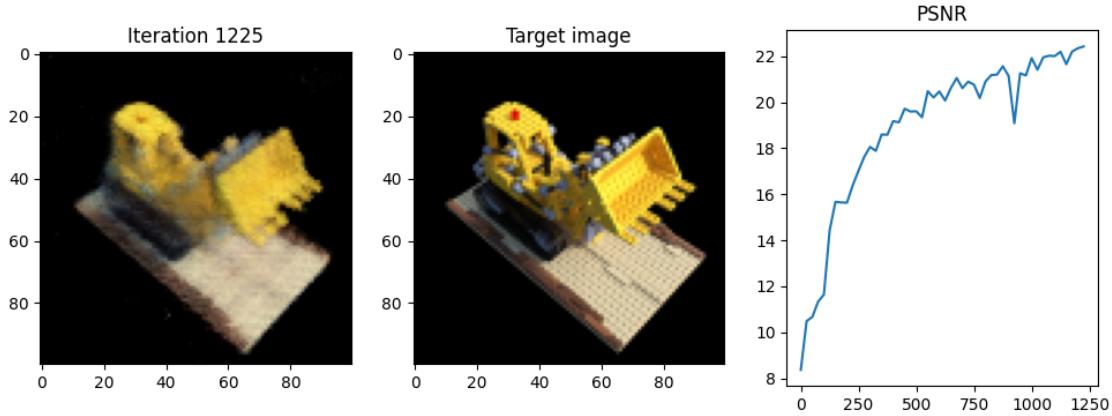
Iteration 1175 Loss: 0.0060 PSNR: 22.20 Time: 0.81 secs per iter, 16.24 mins in total



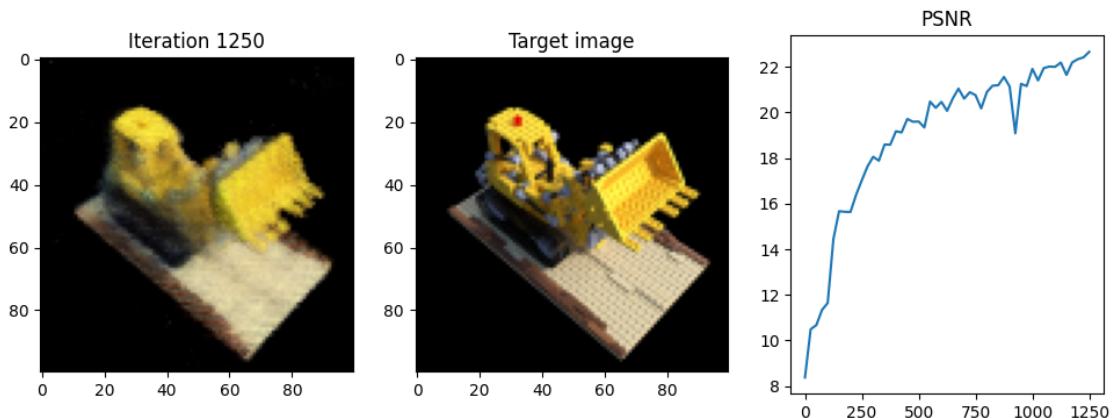
Iteration 1200 Loss: 0.0058 PSNR: 22.34 Time: 0.82 secs per iter, 16.58 mins in total



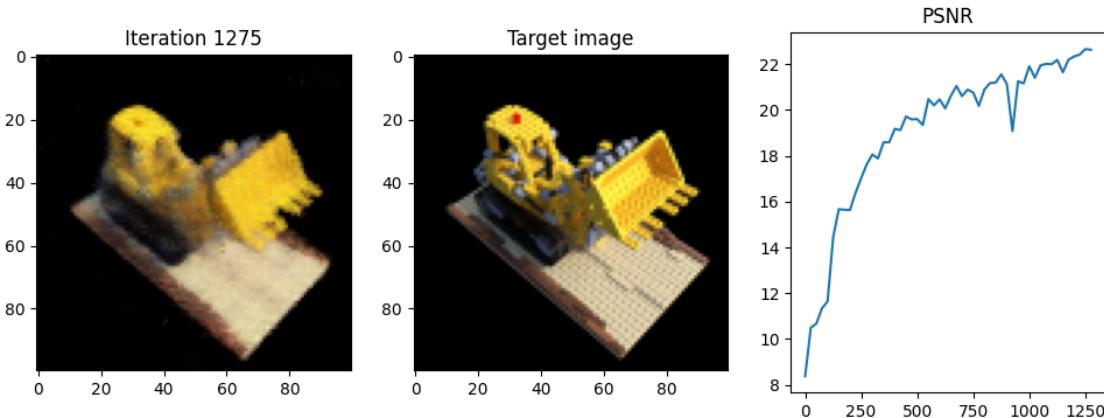
Iteration 1225 Loss: 0.0057 PSNR: 22.42 Time: 0.81 secs per iter, 16.92 mins in total



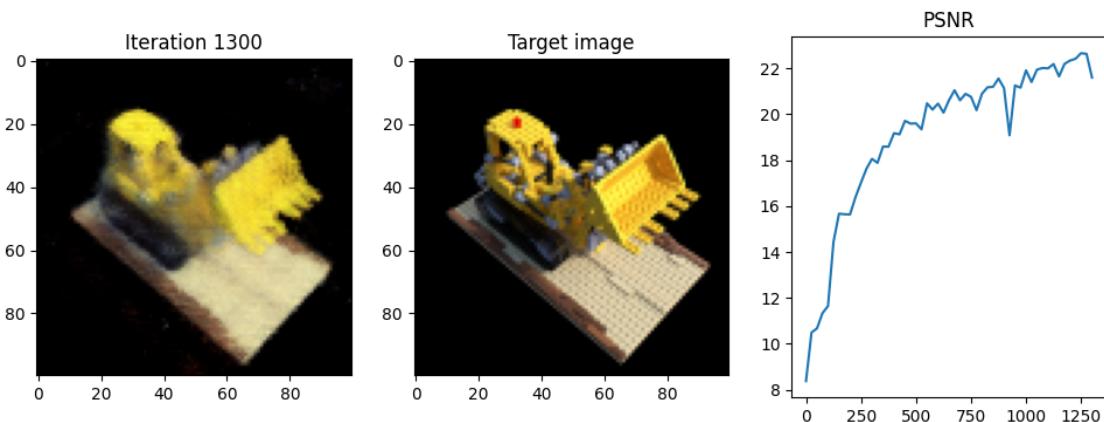
Iteration 1250 Loss: 0.0054 PSNR: 22.66 Time: 0.81 secs per iter, 17.26 mins in total



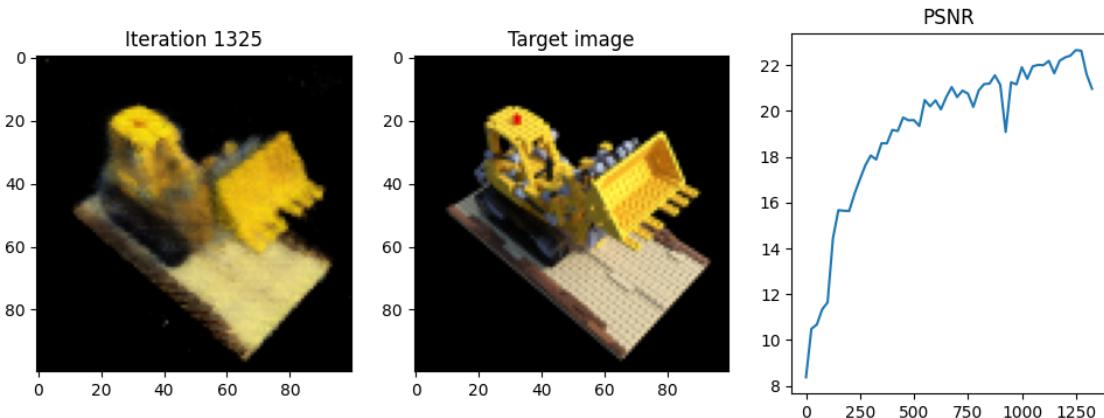
Iteration 1275 Loss: 0.0055 PSNR: 22.64 Time: 0.83 secs per iter, 17.60 mins in total



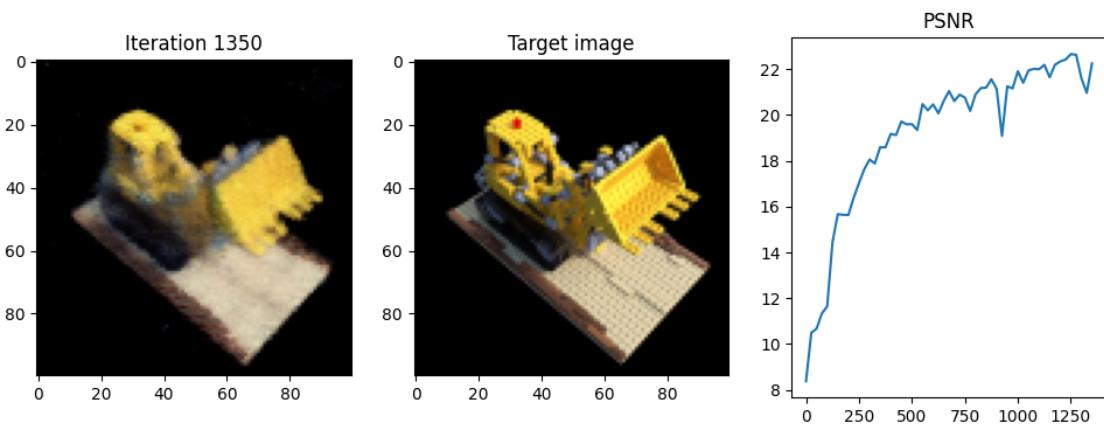
Iteration 1300 Loss: 0.0069 PSNR: 21.60 Time: 0.81 secs per iter, 17.94 mins in total



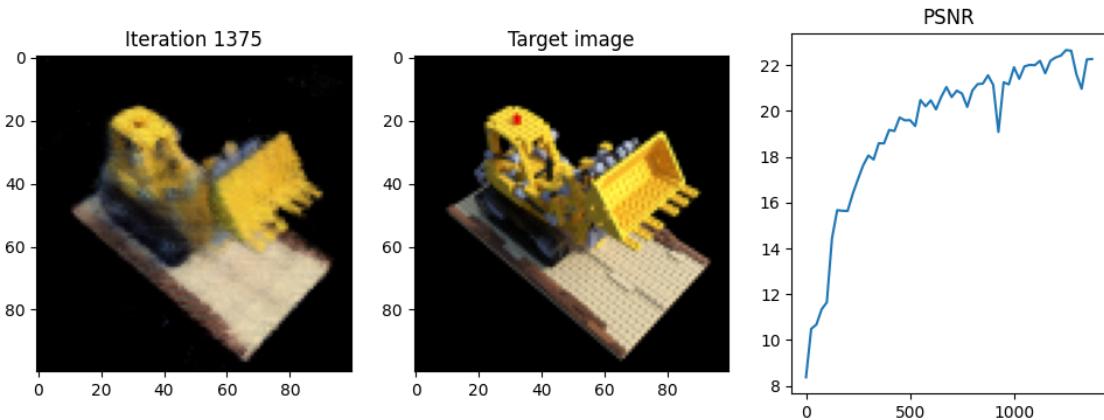
Iteration 1325 Loss: 0.0080 PSNR: 20.97 Time: 0.81 secs per iter, 18.28 mins in total



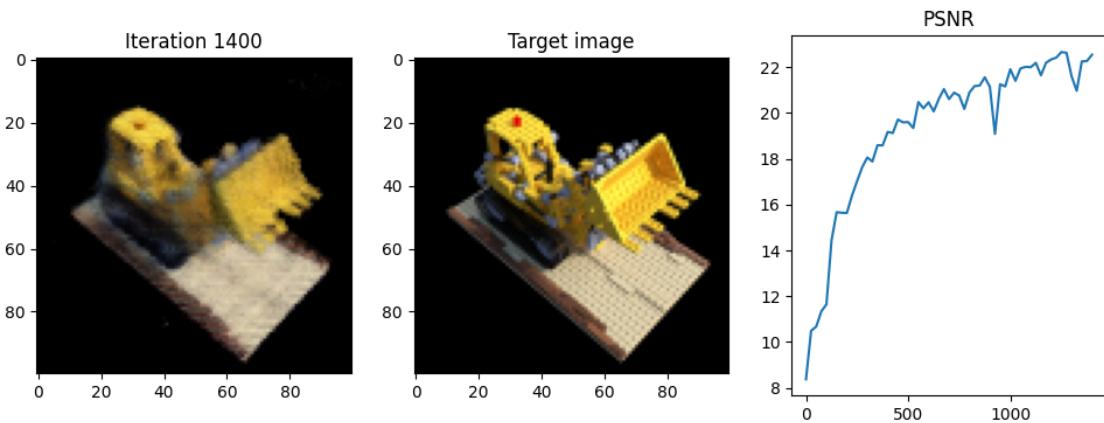
Iteration 1350 Loss: 0.0059 PSNR: 22.26 Time: 0.82 secs per iter, 18.62 mins in total



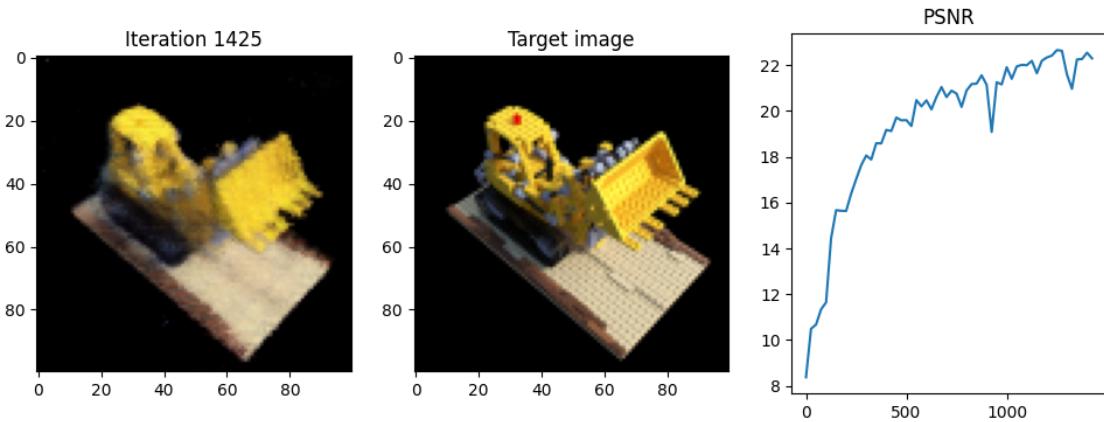
Iteration 1375 Loss: 0.0059 PSNR: 22.27 Time: 0.81 secs per iter, 18.95 mins in total



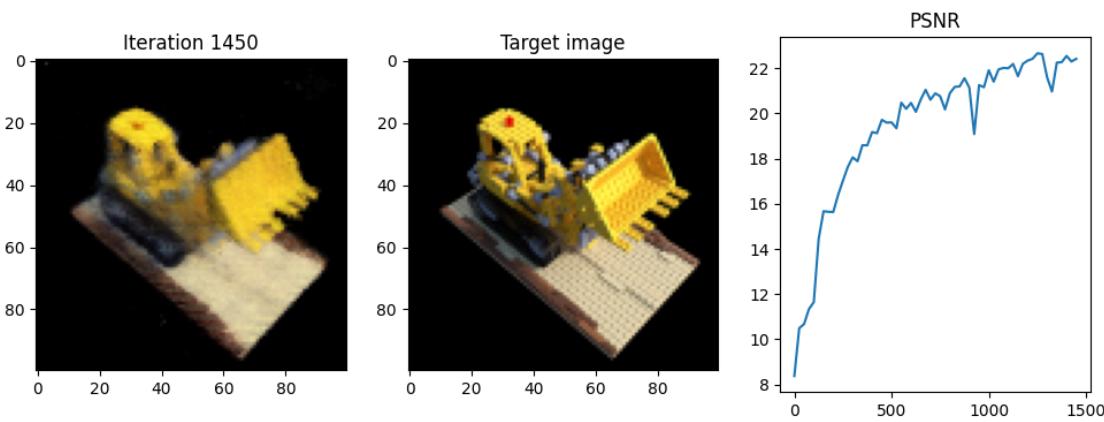
Iteration 1400 Loss: 0.0056 PSNR: 22.55 Time: 0.81 secs per iter, 19.29 mins in total



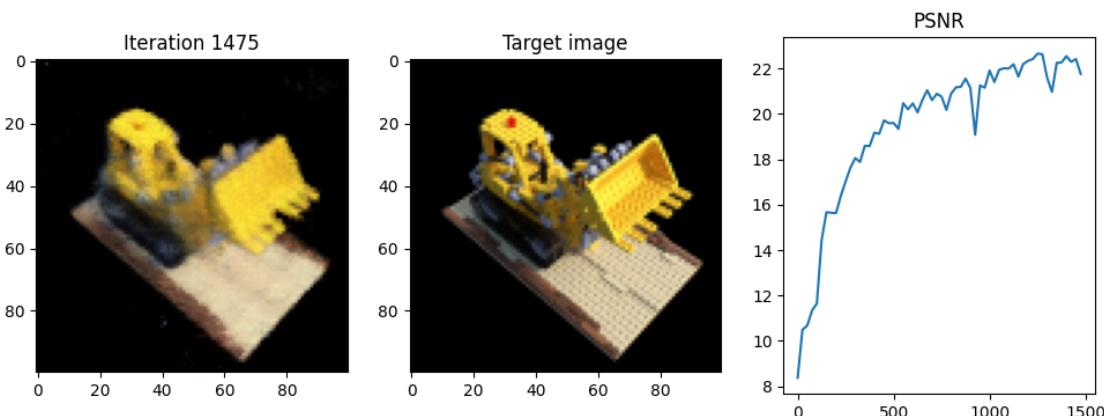
Iteration 1425 Loss: 0.0059 PSNR: 22.30 Time: 0.82 secs per iter, 19.63 mins in total



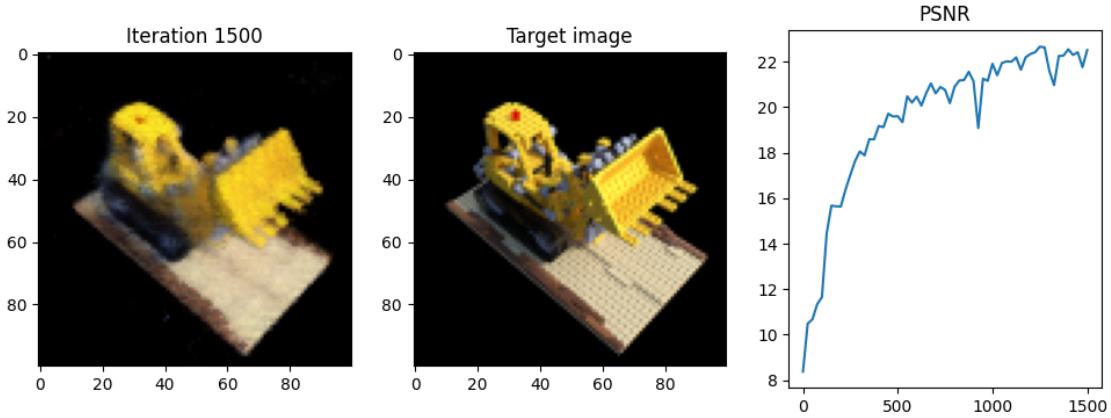
Iteration 1450 Loss: 0.0057 PSNR: 22.42 Time: 0.81 secs per iter, 19.97 mins in total



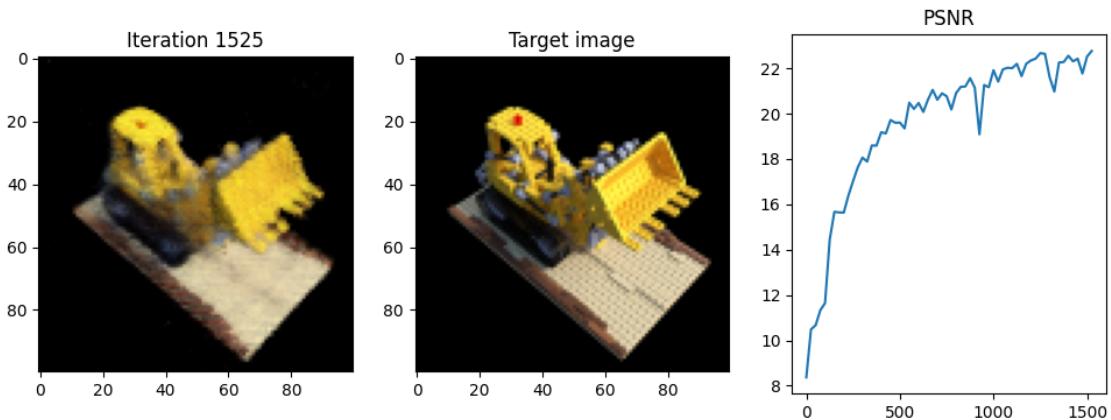
Iteration 1475 Loss: 0.0067 PSNR: 21.77 Time: 0.82 secs per iter, 20.31 mins in total



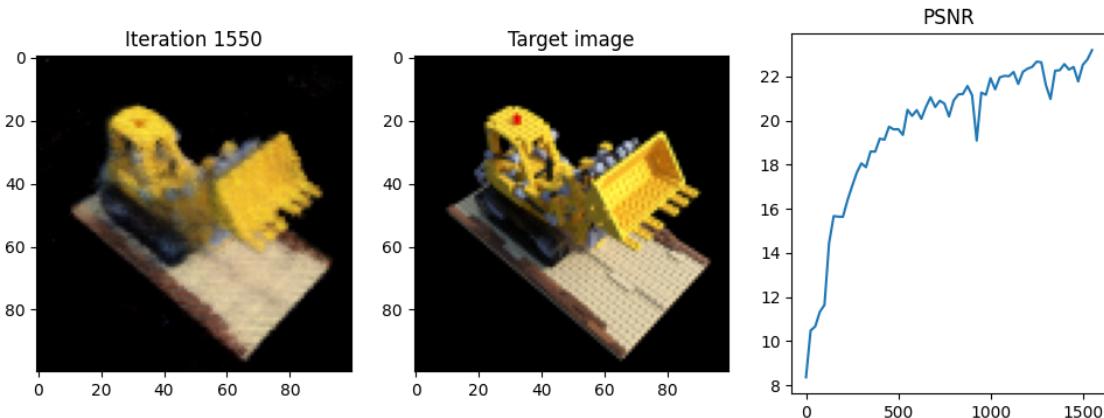
Iteration 1500 Loss: 0.0056 PSNR: 22.52 Time: 0.81 secs per iter, 20.65 mins in total



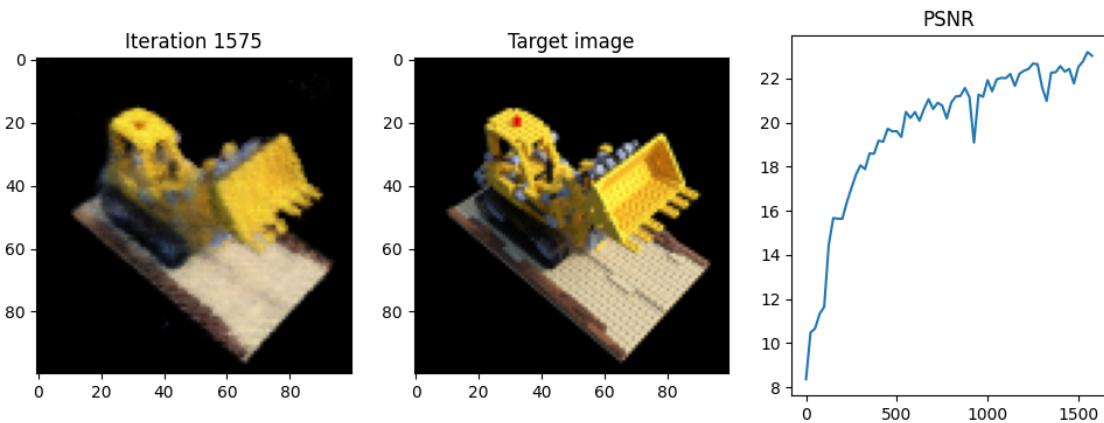
Iteration 1525 Loss: 0.0053 PSNR: 22.76 Time: 0.81 secs per iter, 20.98 mins in total



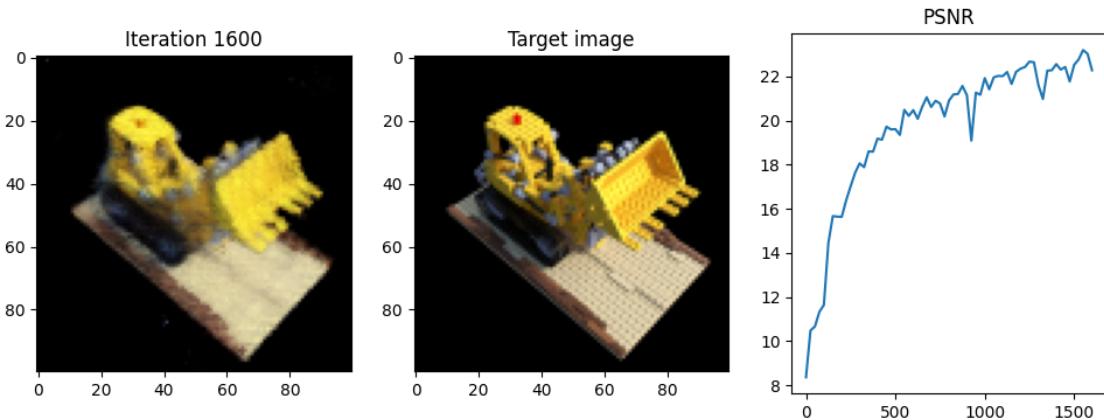
Iteration 1550 Loss: 0.0048 PSNR: 23.19 Time: 0.81 secs per iter, 21.32 mins in total



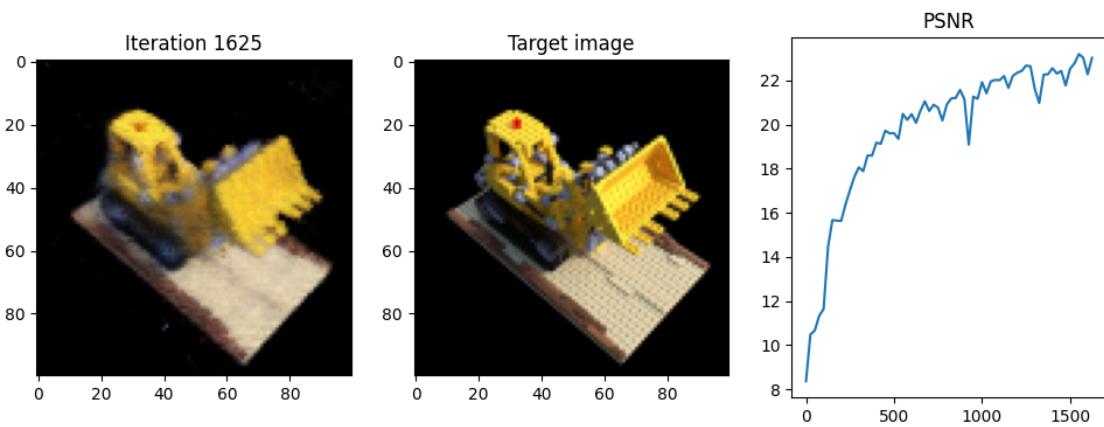
Iteration 1575 Loss: 0.0050 PSNR: 23.02 Time: 0.81 secs per iter, 21.66 mins in total



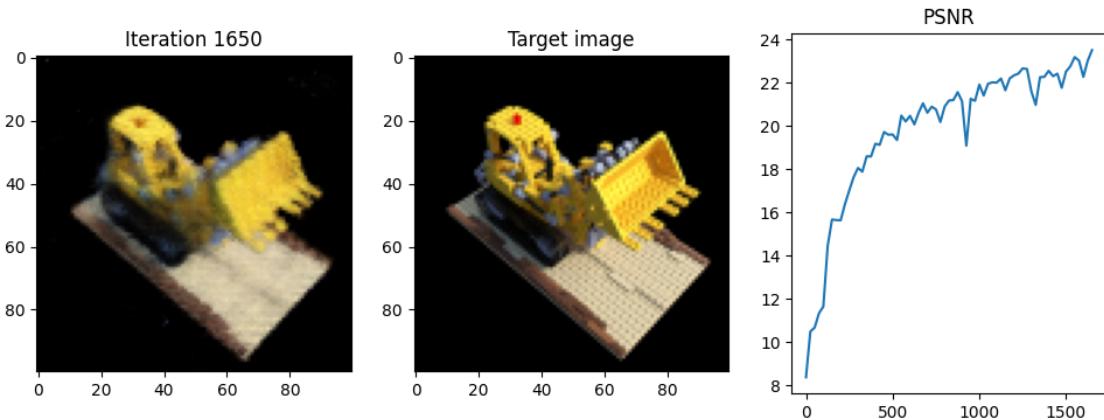
Iteration 1600 Loss: 0.0059 PSNR: 22.27 Time: 0.81 secs per iter, 22.00 mins in total



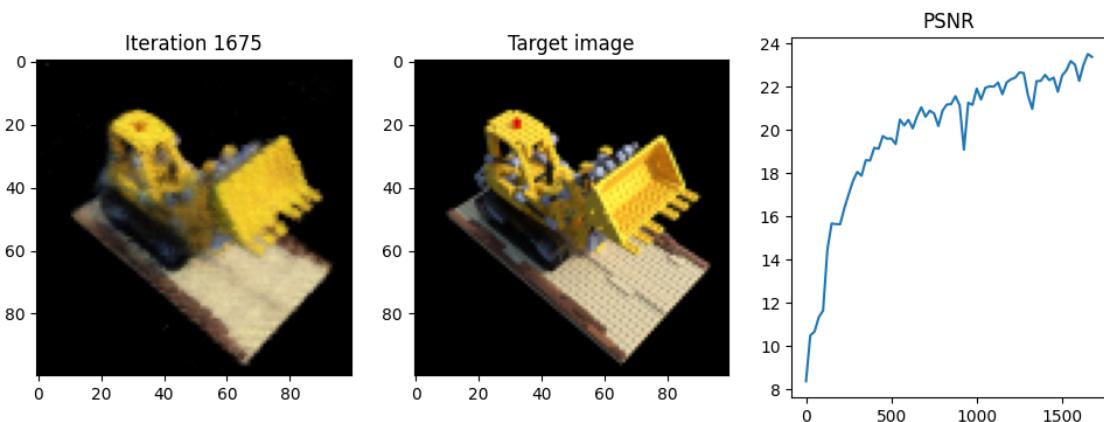
Iteration 1625 Loss: 0.0050 PSNR: 23.02 Time: 0.81 secs per iter, 22.33 mins in total



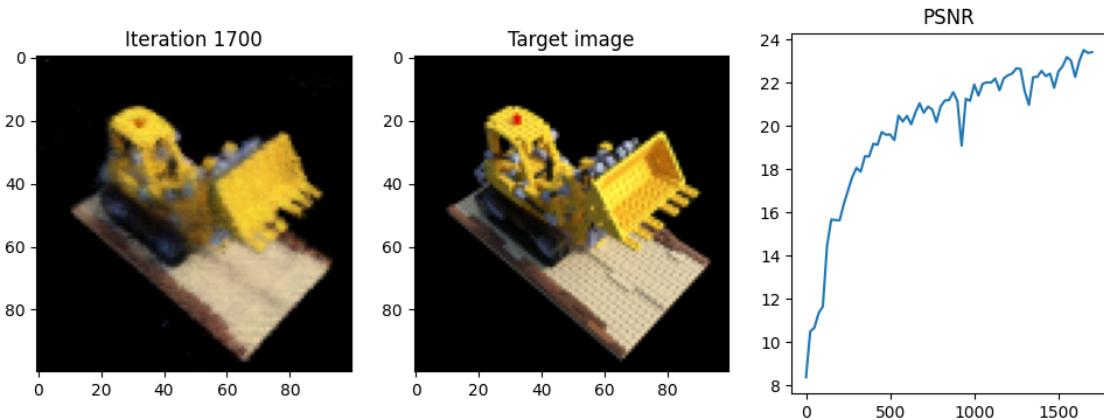
Iteration 1650 Loss: 0.0045 PSNR: 23.51 Time: 0.81 secs per iter, 22.67 mins in total



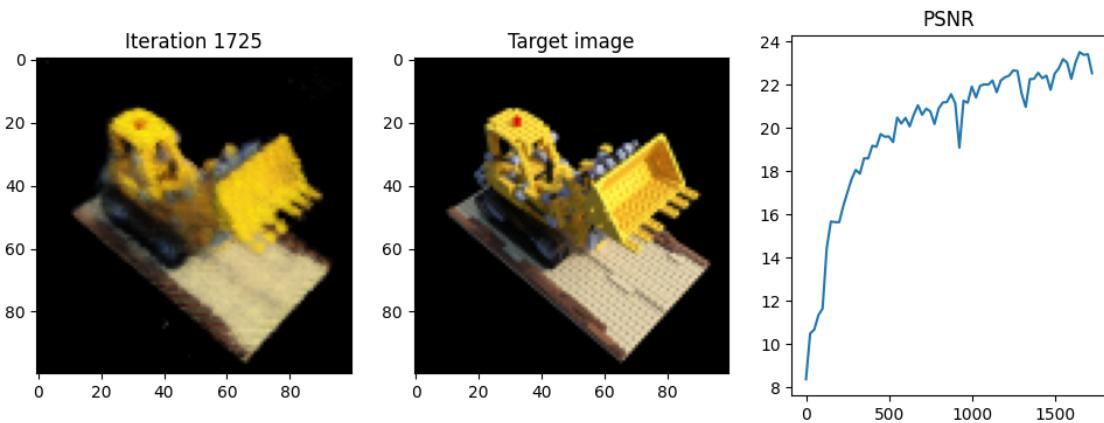
Iteration 1675 Loss: 0.0046 PSNR: 23.38 Time: 0.81 secs per iter, 23.01 mins in total



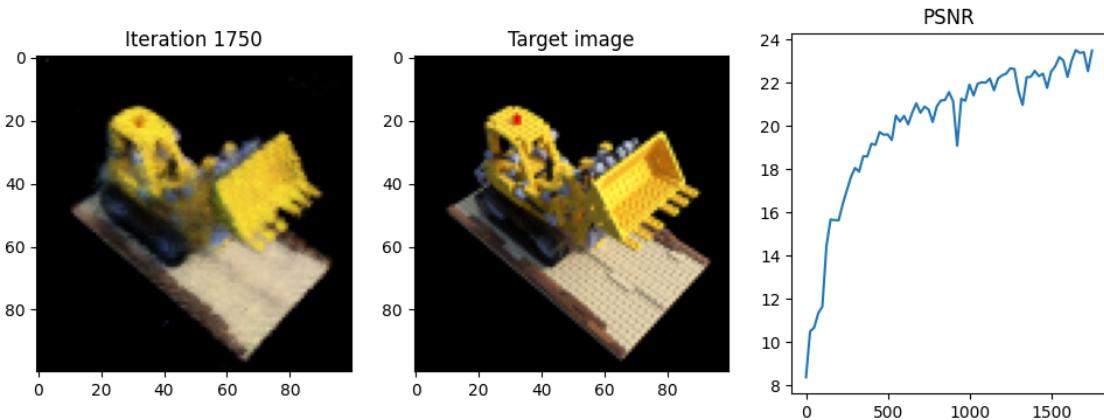
Iteration 1700 Loss: 0.0046 PSNR: 23.41 Time: 0.81 secs per iter, 23.35 mins in total



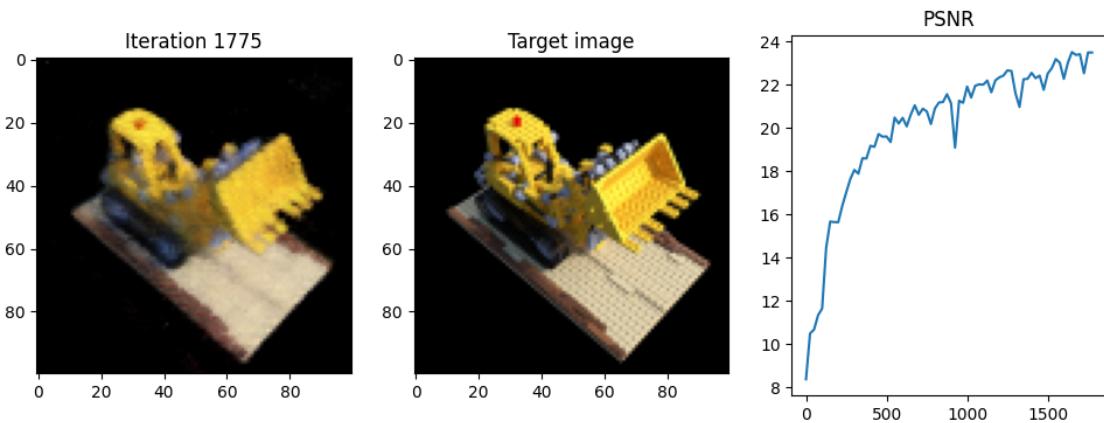
Iteration 1725 Loss: 0.0056 PSNR: 22.53 Time: 0.81 secs per iter, 23.69 mins in total



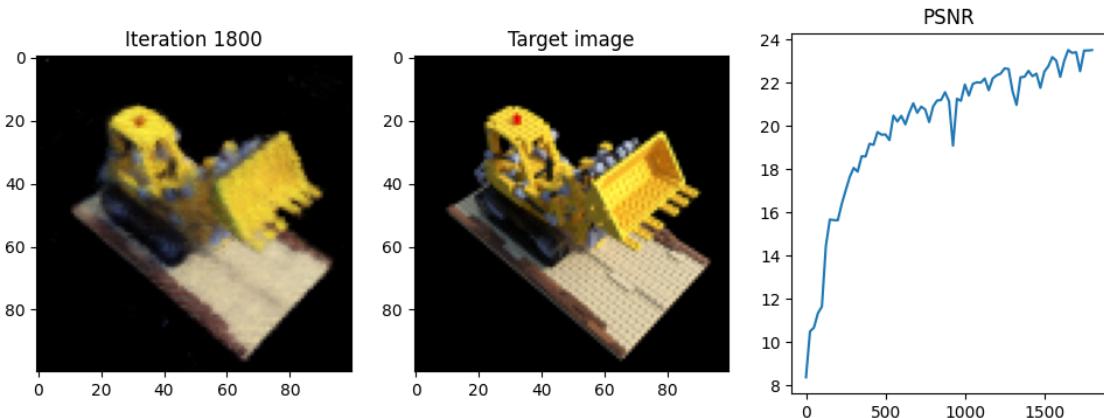
Iteration 1750 Loss: 0.0045 PSNR: 23.49 Time: 0.81 secs per iter, 24.03 mins in total



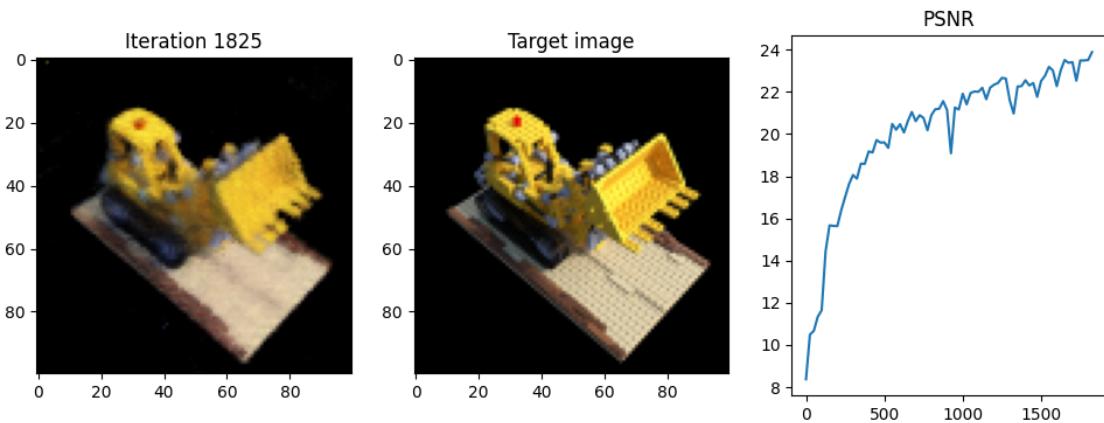
Iteration 1775 Loss: 0.0045 PSNR: 23.49 Time: 0.81 secs per iter, 24.36 mins in total



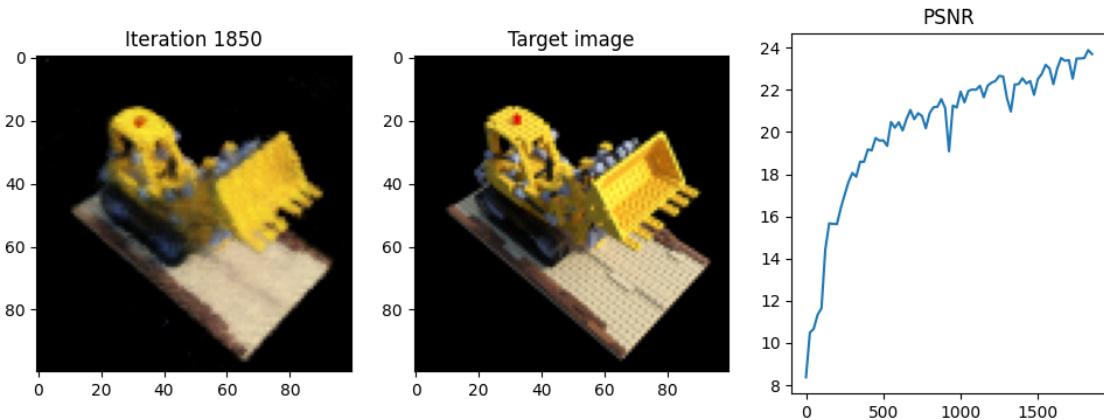
Iteration 1800 Loss: 0.0045 PSNR: 23.51 Time: 0.81 secs per iter, 24.70 mins in total



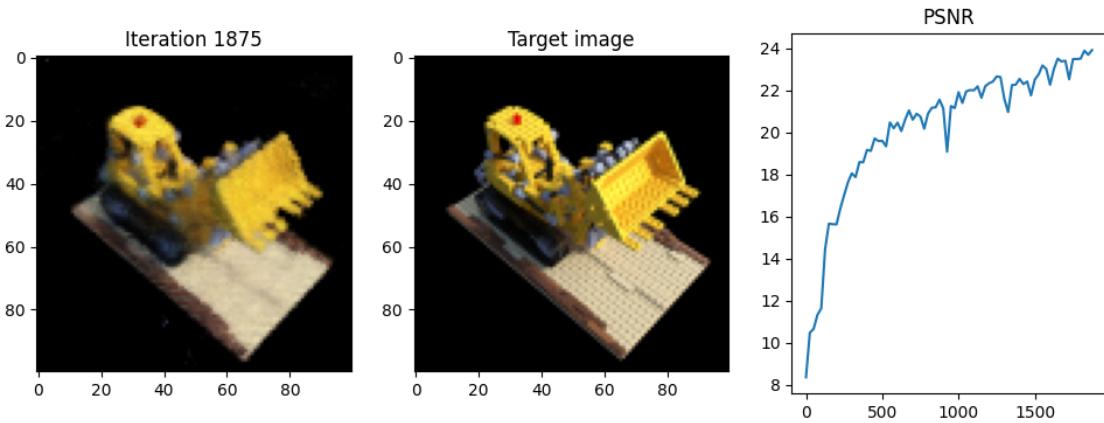
Iteration 1825 Loss: 0.0041 PSNR: 23.89 Time: 0.81 secs per iter, 25.04 mins in total



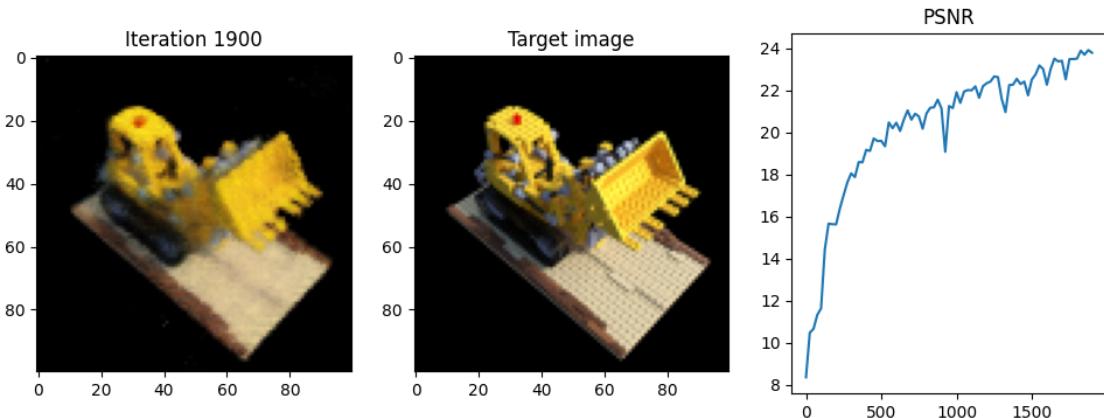
Iteration 1850 Loss: 0.0043 PSNR: 23.70 Time: 0.82 secs per iter, 25.38 mins in total



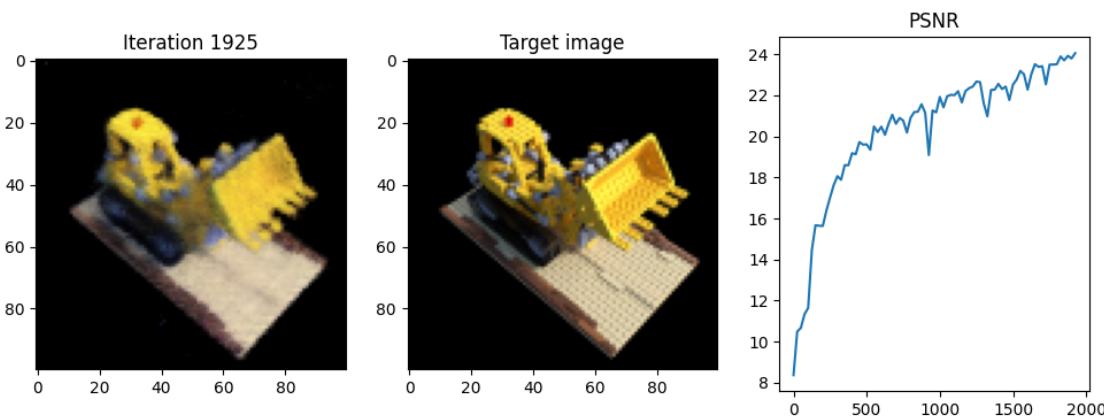
Iteration 1875 Loss: 0.0041 PSNR: 23.92 Time: 0.81 secs per iter, 25.72 mins in total



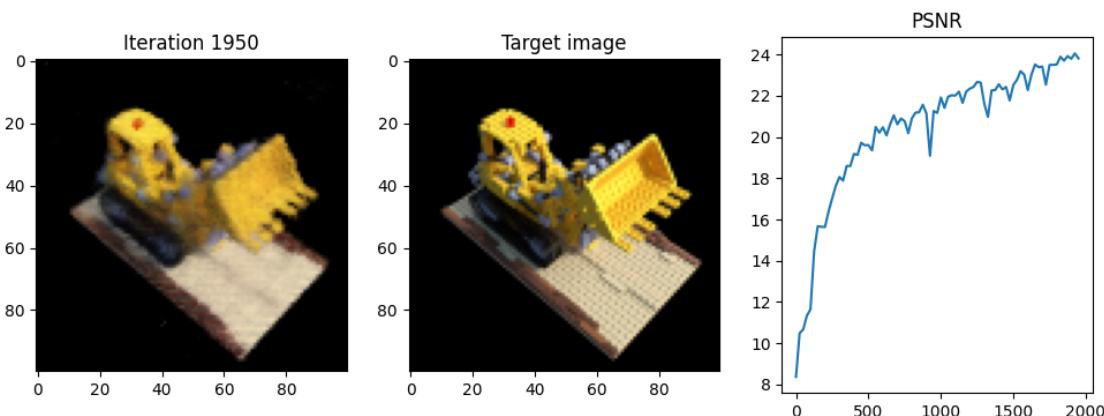
Iteration 1900 Loss: 0.0042 PSNR: 23.79 Time: 0.81 secs per iter, 26.05 mins in total



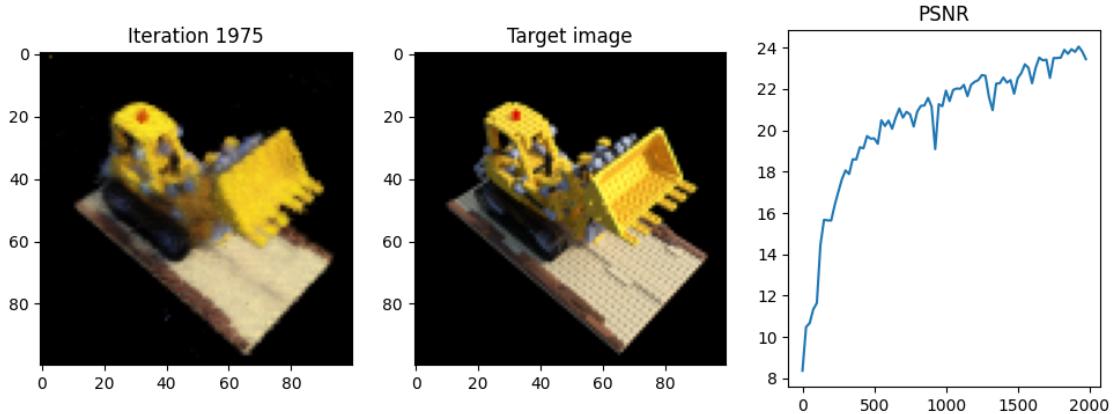
Iteration 1925 Loss: 0.0039 PSNR: 24.05 Time: 0.81 secs per iter, 26.39 mins in total



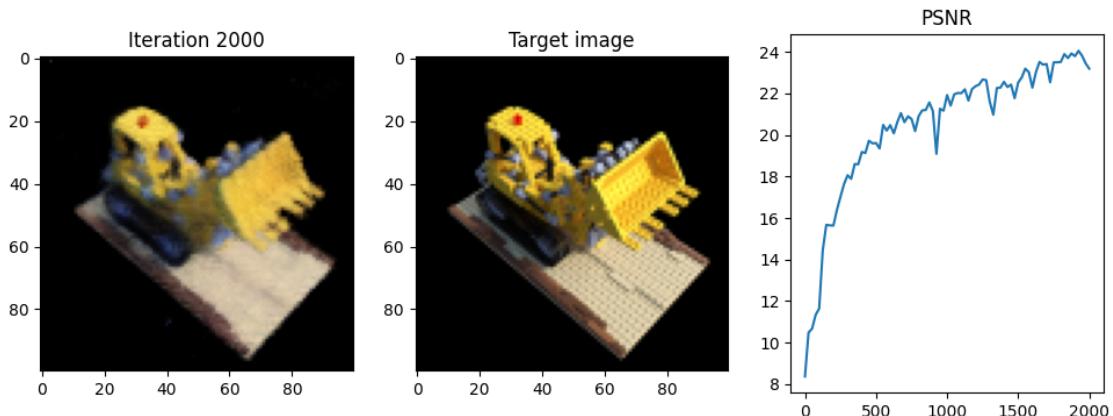
Iteration 1950 Loss: 0.0042 PSNR: 23.80 Time: 0.81 secs per iter, 26.73 mins in total



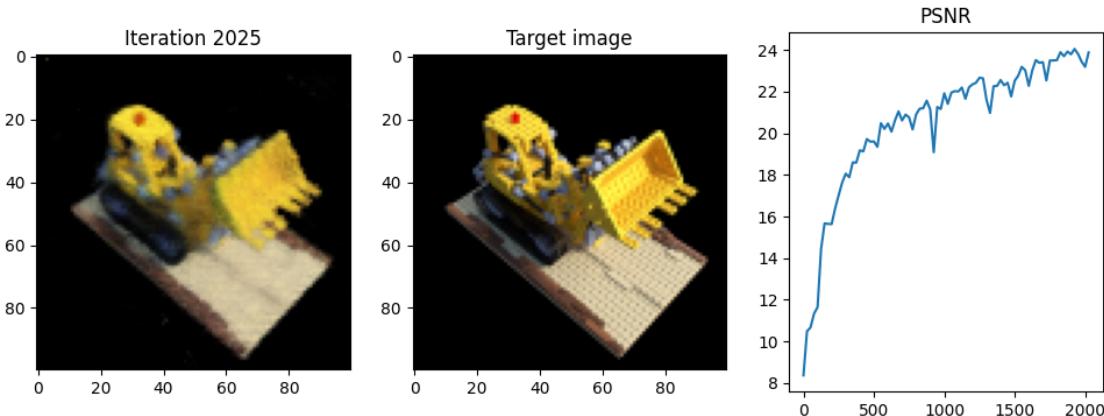
Iteration 1975 Loss: 0.0045 PSNR: 23.44 Time: 0.81 secs per iter, 27.07 mins in total



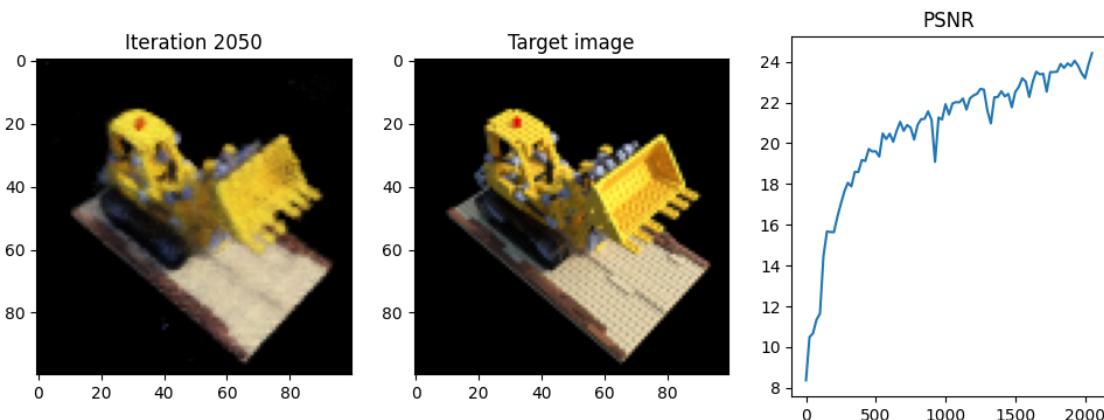
Iteration 2000 Loss: 0.0048 PSNR: 23.19 Time: 0.82 secs per iter, 27.41 mins in total



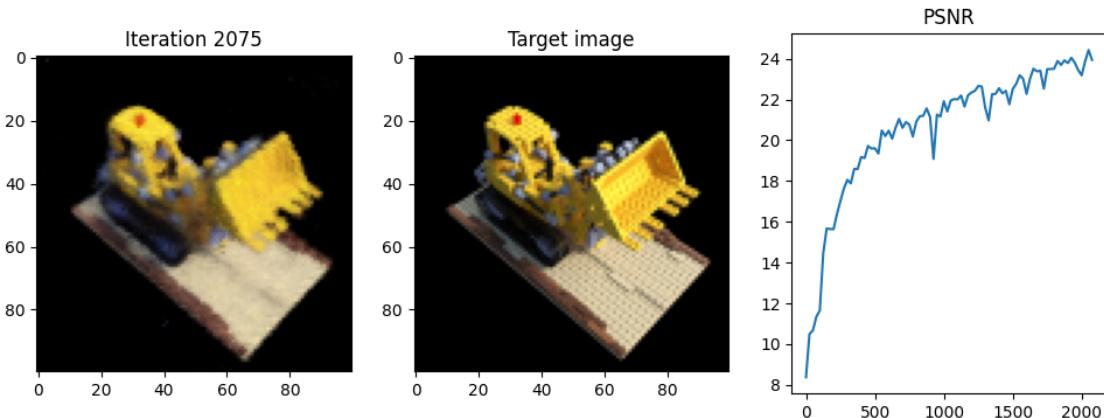
Iteration 2025 Loss: 0.0041 PSNR: 23.88 Time: 0.81 secs per iter, 27.75 mins in total



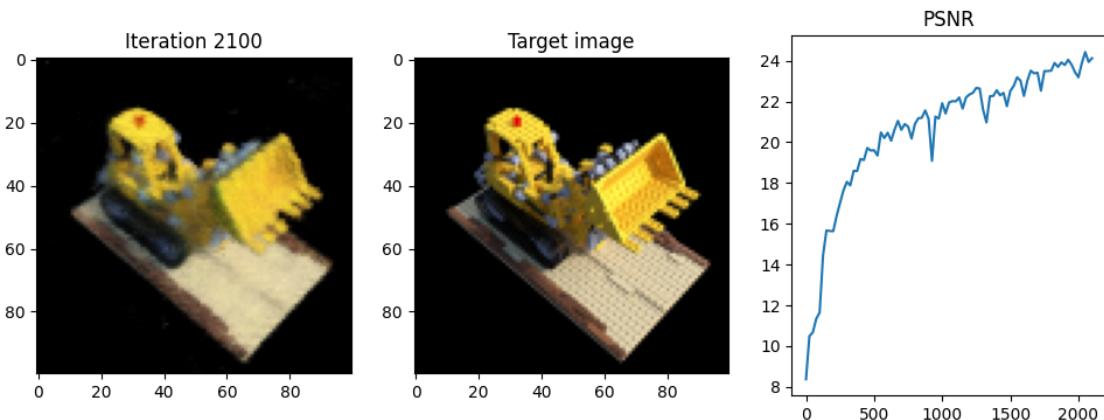
Iteration 2050 Loss: 0.0036 PSNR: 24.42 Time: 0.81 secs per iter, 28.08 mins in total



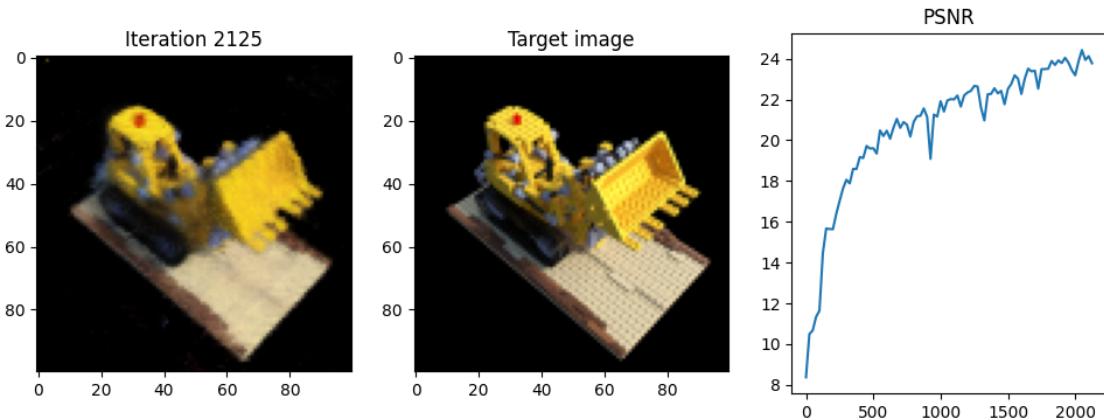
Iteration 2075 Loss: 0.0040 PSNR: 23.94 Time: 0.82 secs per iter, 28.42 mins in total



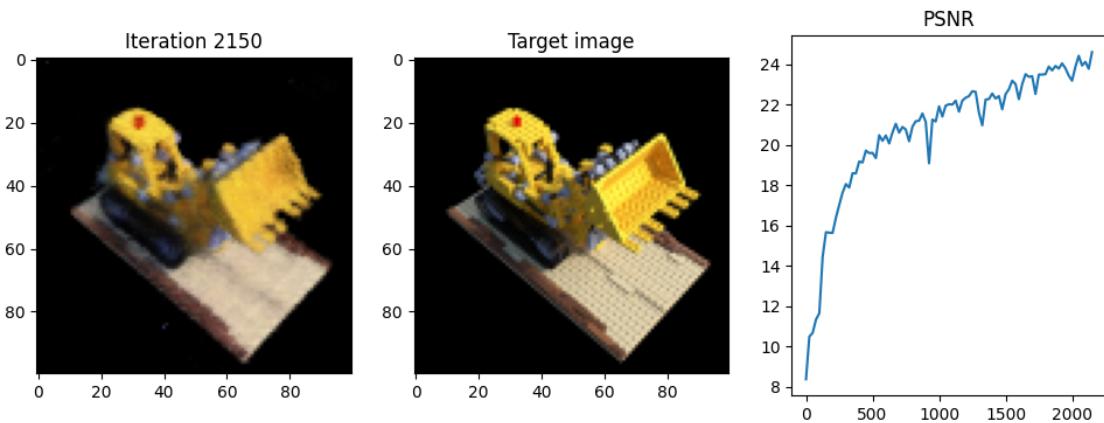
Iteration 2100 Loss: 0.0039 PSNR: 24.12 Time: 0.81 secs per iter, 28.76 mins in total



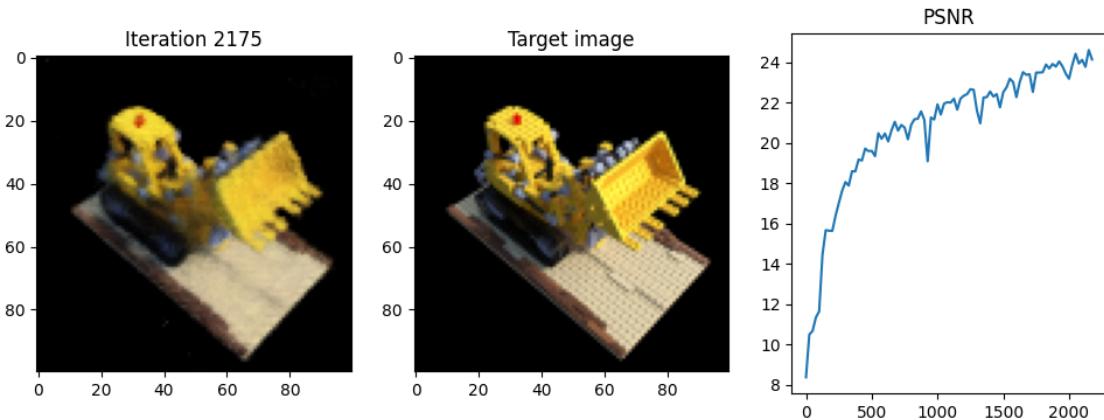
Iteration 2125 Loss: 0.0042 PSNR: 23.78 Time: 0.81 secs per iter, 29.10 mins in total



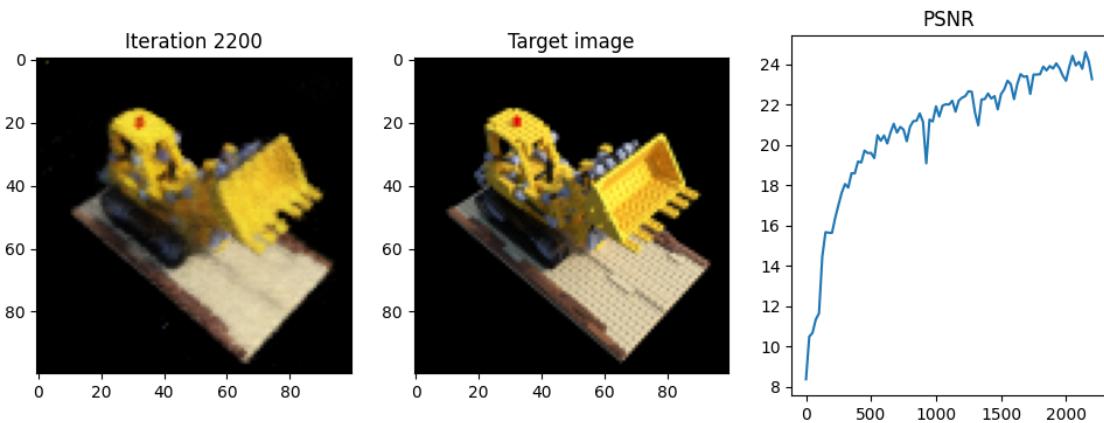
Iteration 2150 Loss: 0.0035 PSNR: 24.61 Time: 0.82 secs per iter, 29.44 mins in total



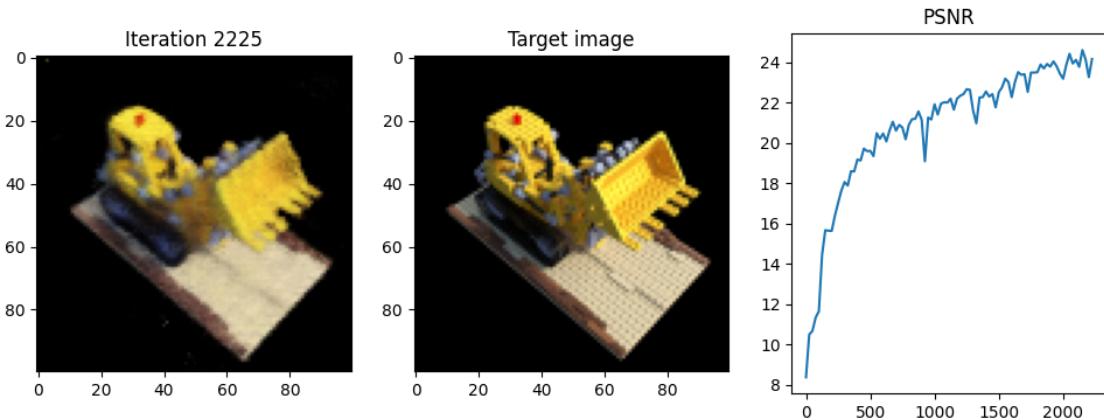
Iteration 2175 Loss: 0.0039 PSNR: 24.14 Time: 0.81 secs per iter, 29.77 mins in total



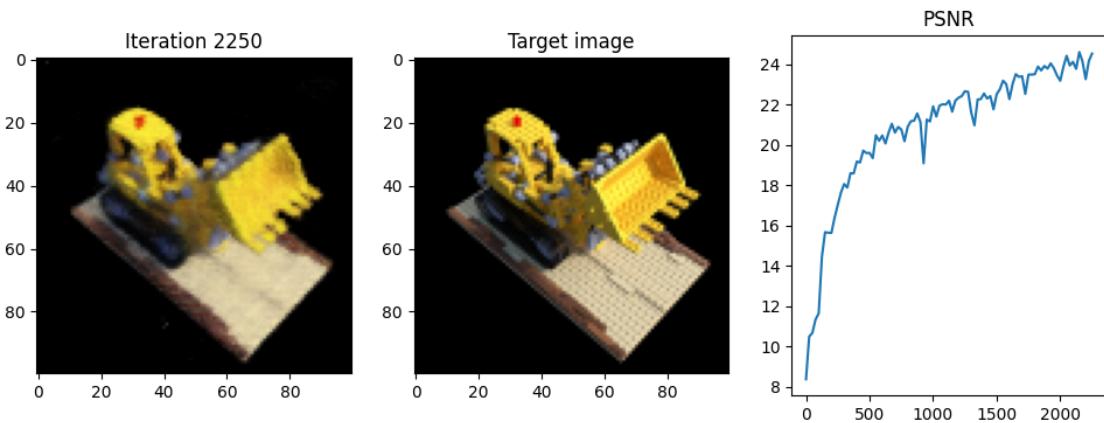
Iteration 2200 Loss: 0.0047 PSNR: 23.26 Time: 0.81 secs per iter, 30.11 mins in total



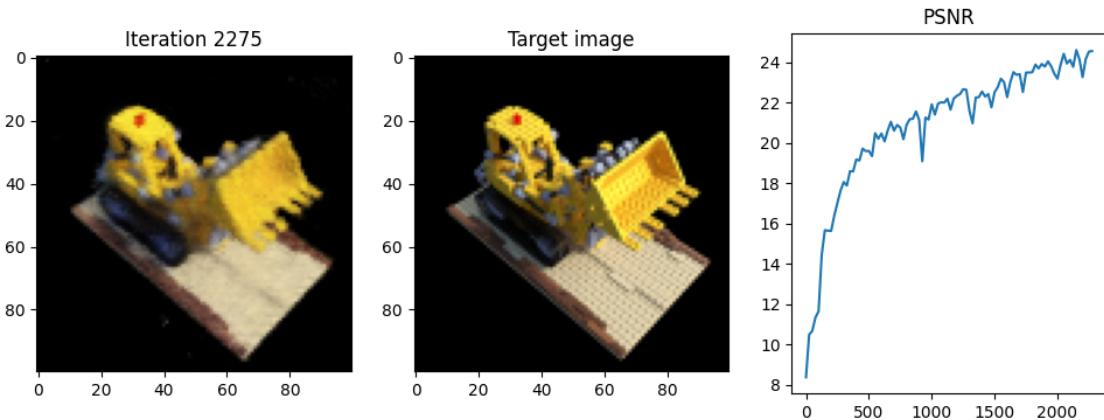
Iteration 2225 Loss: 0.0038 PSNR: 24.16 Time: 0.82 secs per iter, 30.45 mins in total



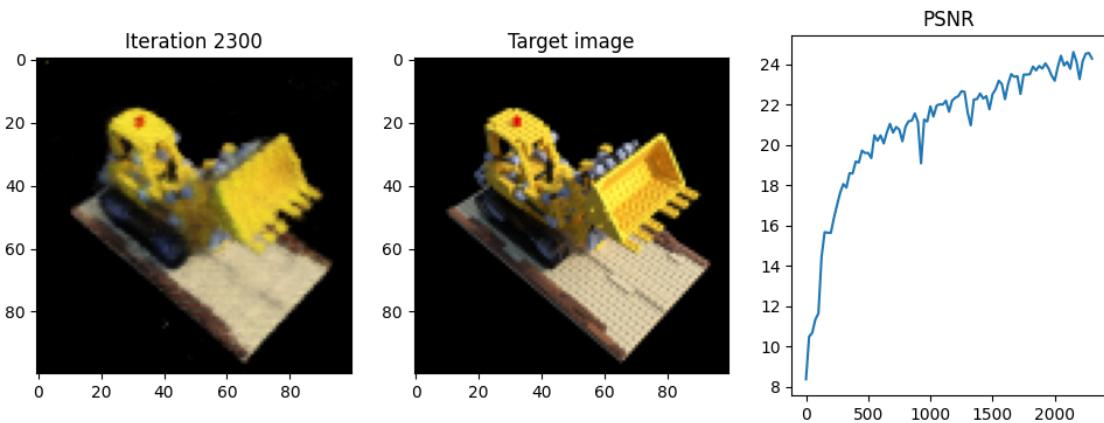
Iteration 2250 Loss: 0.0035 PSNR: 24.53 Time: 0.81 secs per iter, 30.79 mins in total



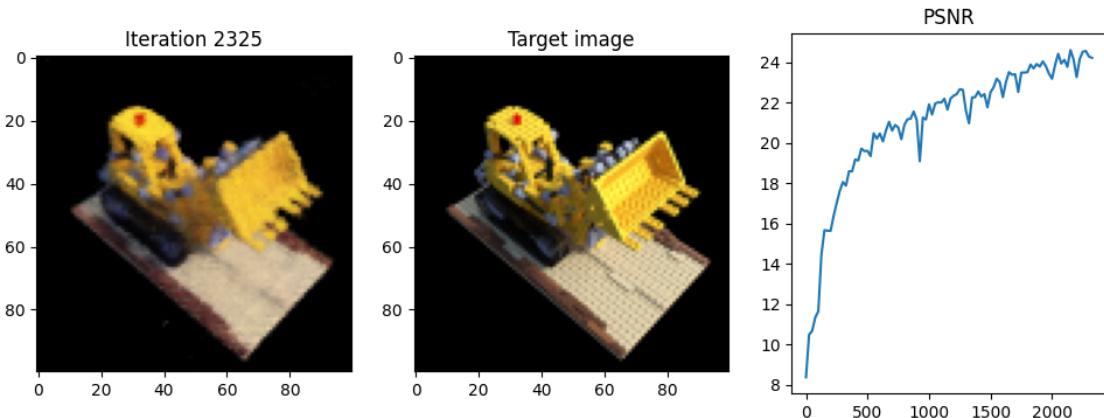
Iteration 2275 Loss: 0.0035 PSNR: 24.56 Time: 0.81 secs per iter, 31.13 mins in total



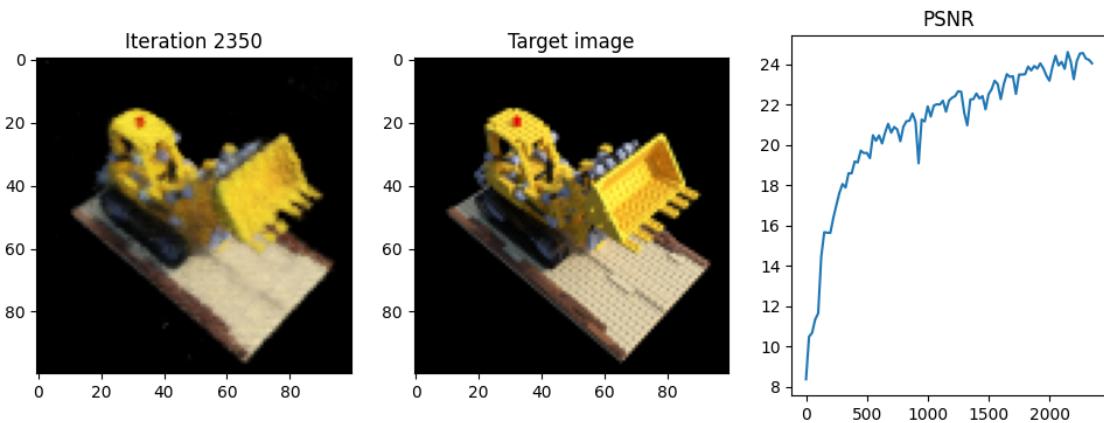
Iteration 2300 Loss: 0.0037 PSNR: 24.28 Time: 0.81 secs per iter, 31.46 mins in total



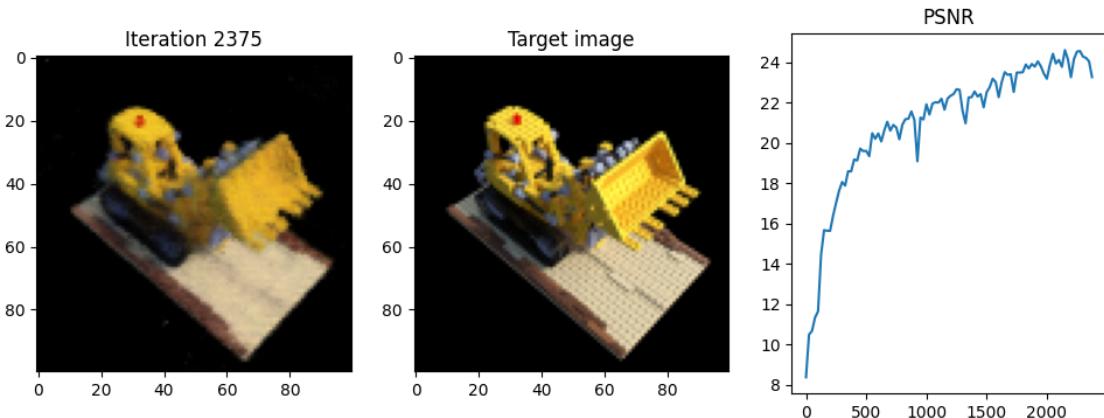
Iteration 2325 Loss: 0.0038 PSNR: 24.22 Time: 0.81 secs per iter, 31.80 mins in total



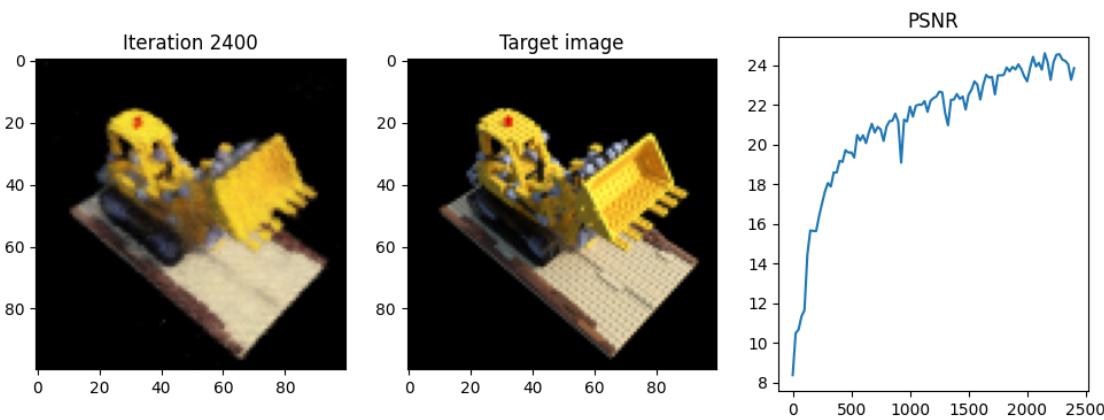
Iteration 2350 Loss: 0.0039 PSNR: 24.05 Time: 0.81 secs per iter, 32.14 mins in total



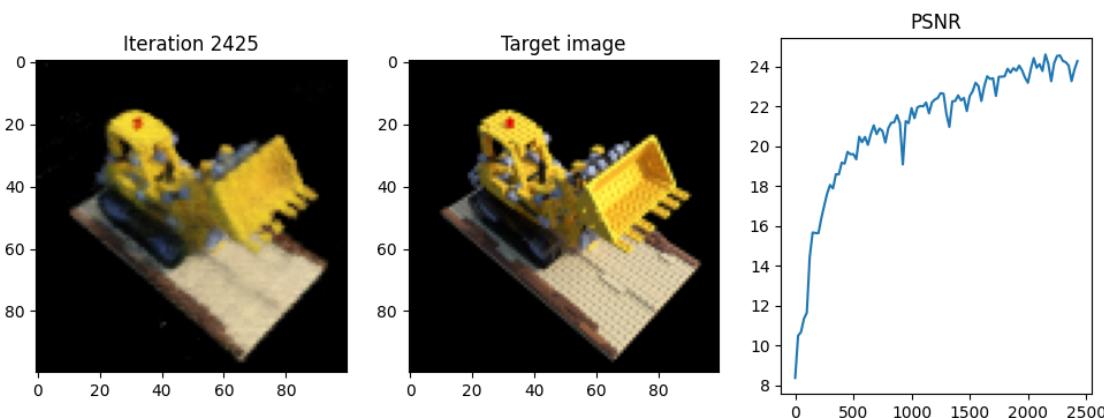
Iteration 2375 Loss: 0.0047 PSNR: 23.27 Time: 0.81 secs per iter, 32.48 mins in total



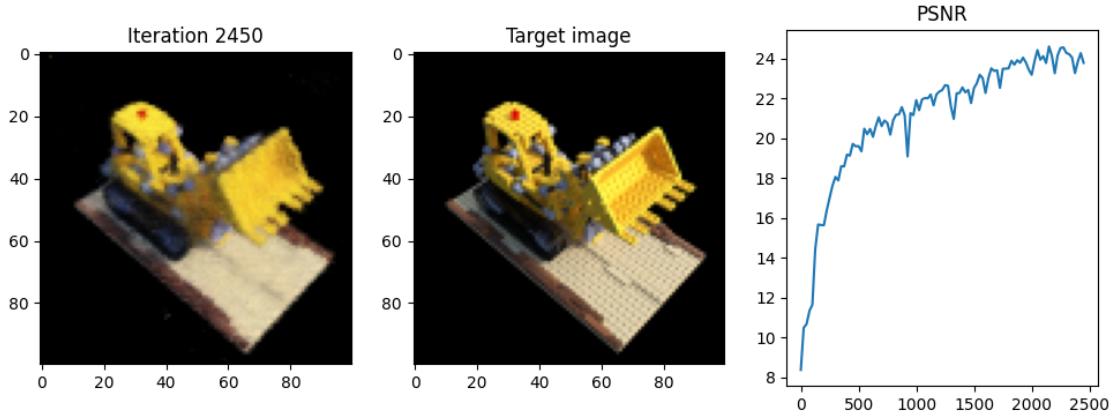
Iteration 2400 Loss: 0.0041 PSNR: 23.85 Time: 0.82 secs per iter, 32.82 mins in total



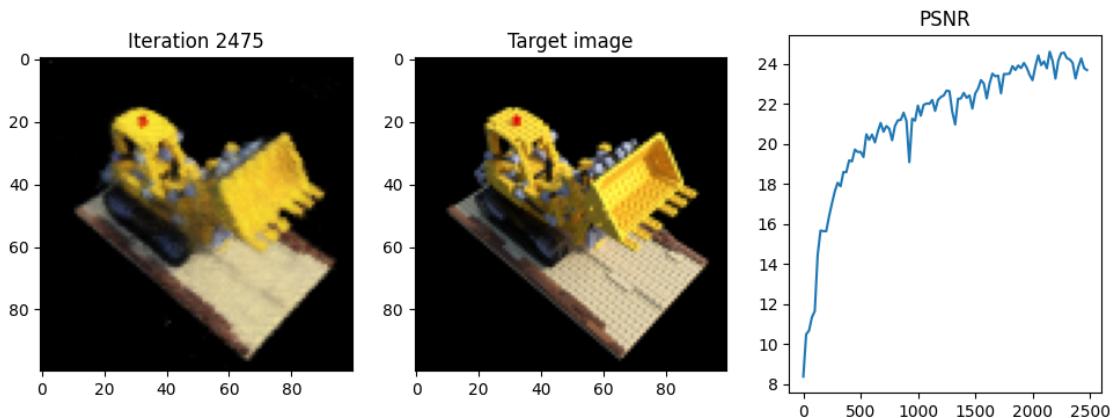
Iteration 2425 Loss: 0.0037 PSNR: 24.27 Time: 0.81 secs per iter, 33.15 mins in total



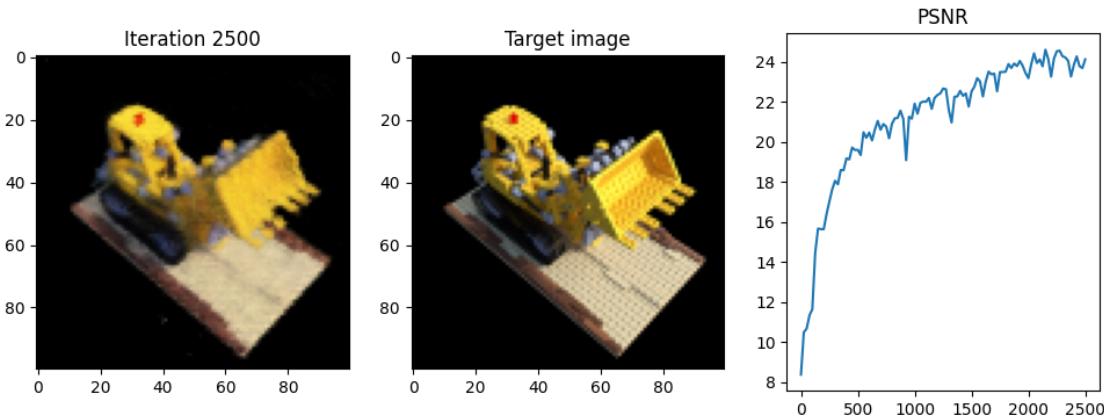
Iteration 2450 Loss: 0.0042 PSNR: 23.79 Time: 0.81 secs per iter, 33.49 mins in total



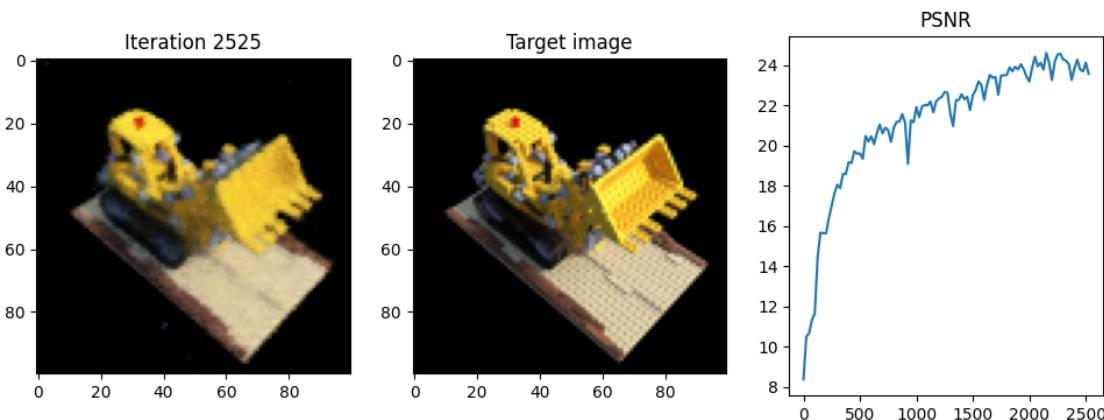
Iteration 2475 Loss: 0.0043 PSNR: 23.69 Time: 0.81 secs per iter, 33.83 mins in total



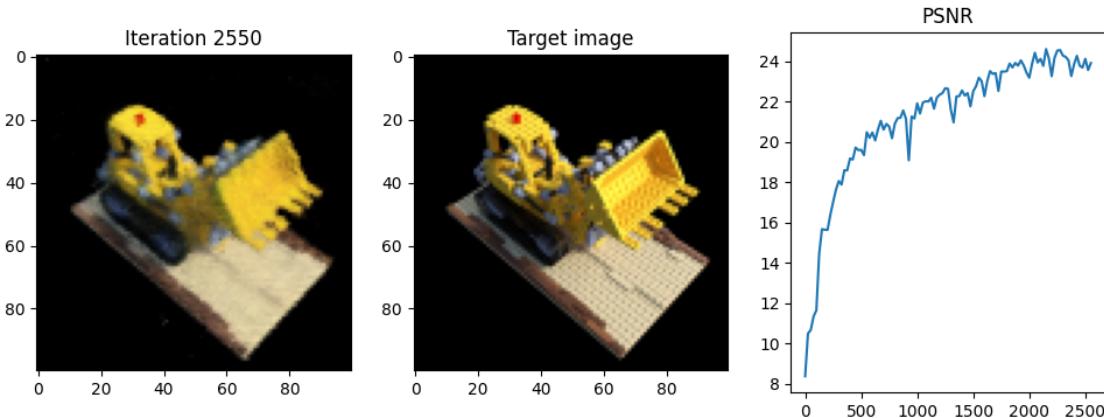
Iteration 2500 Loss: 0.0039 PSNR: 24.13 Time: 0.81 secs per iter, 34.17 mins in total



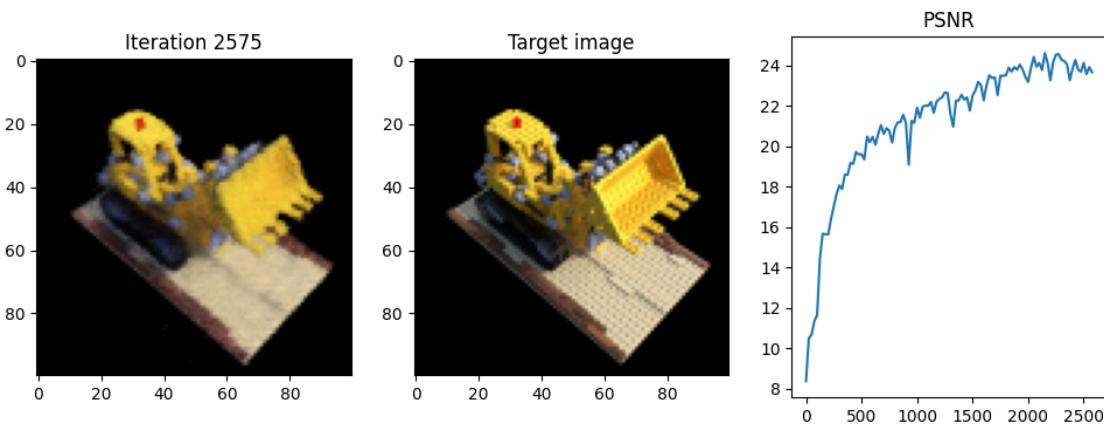
Iteration 2525 Loss: 0.0044 PSNR: 23.57 Time: 0.81 secs per iter, 34.51 mins in total



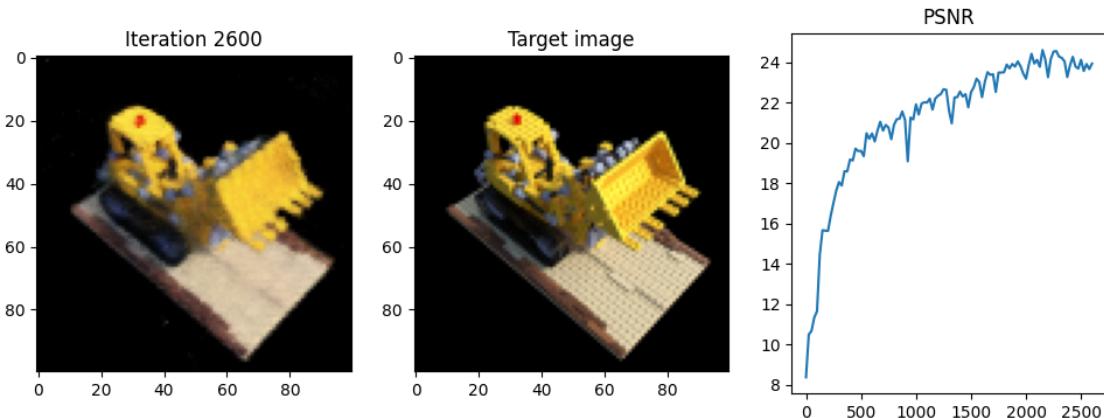
Iteration 2550 Loss: 0.0041 PSNR: 23.92 Time: 0.81 secs per iter, 34.84 mins in total



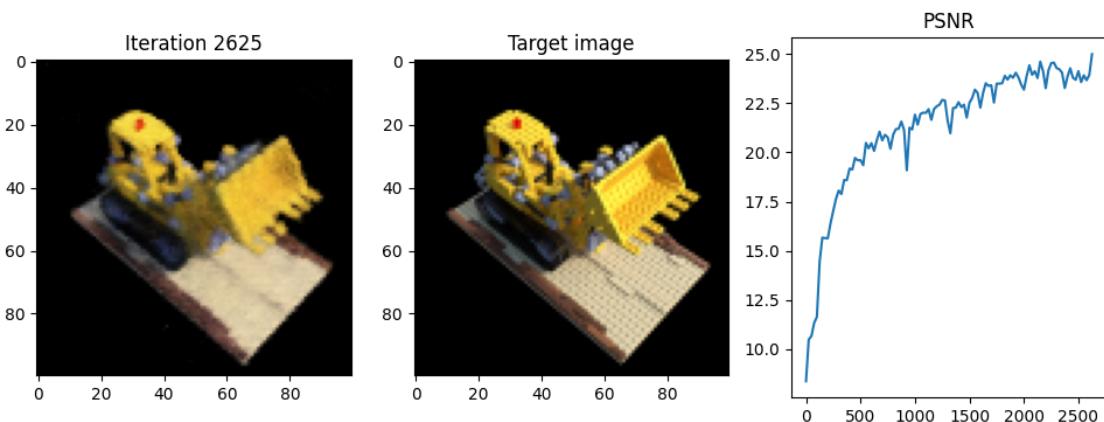
Iteration 2575 Loss: 0.0043 PSNR: 23.66 Time: 0.83 secs per iter, 35.19 mins in total



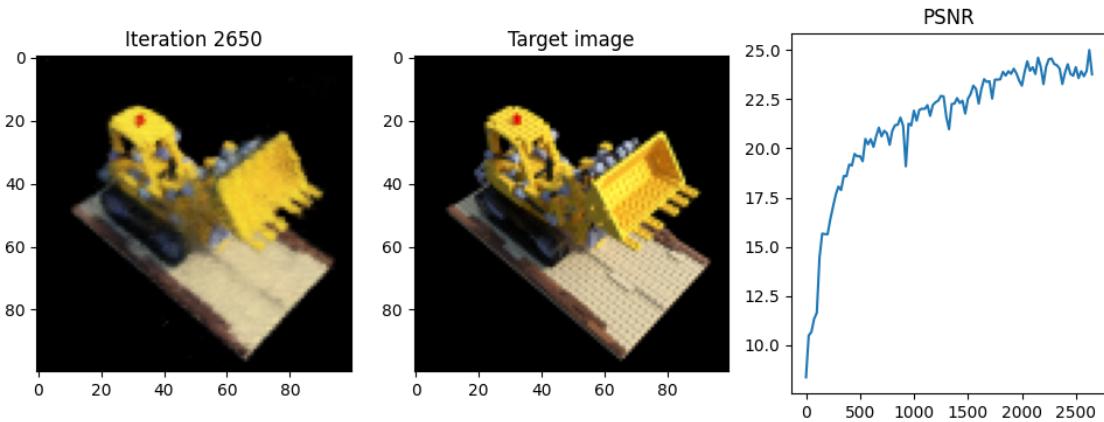
Iteration 2600 Loss: 0.0040 PSNR: 23.94 Time: 0.81 secs per iter, 35.53 mins in total



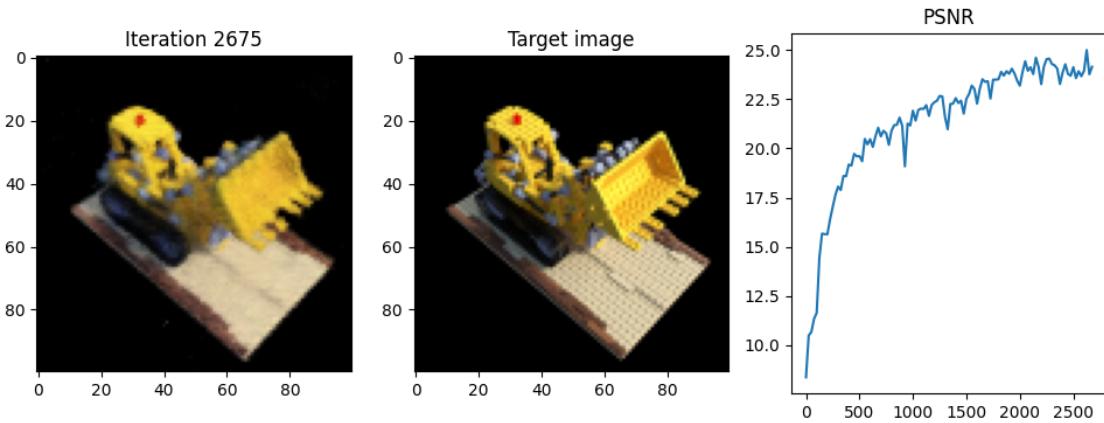
Iteration 2625 Loss: 0.0032 PSNR: 24.99 Time: 0.81 secs per iter, 35.87 mins in total



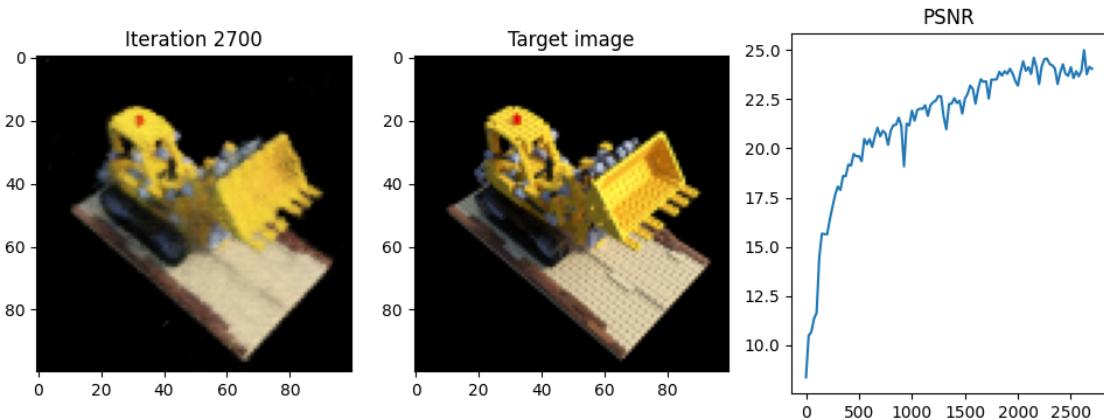
Iteration 2650 Loss: 0.0042 PSNR: 23.76 Time: 0.82 secs per iter, 36.21 mins in total



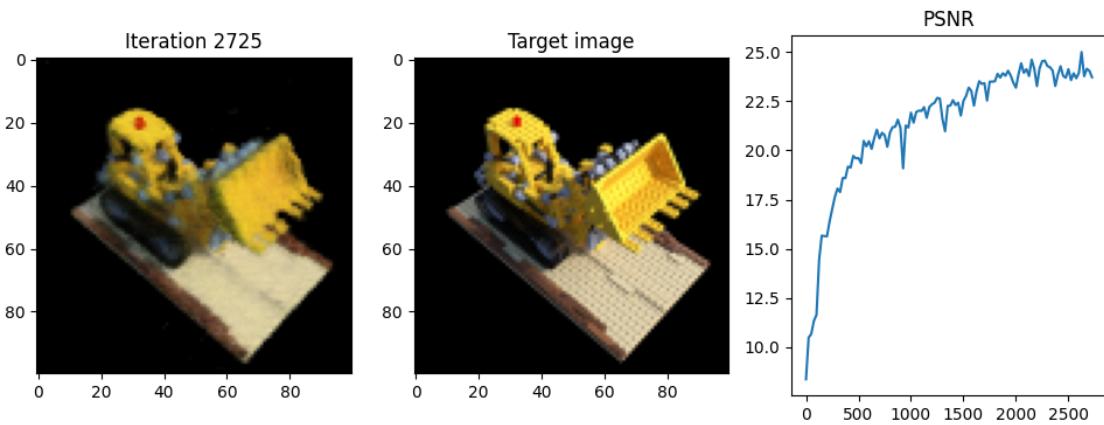
Iteration 2675 Loss: 0.0039 PSNR: 24.14 Time: 0.81 secs per iter, 36.55 mins in total



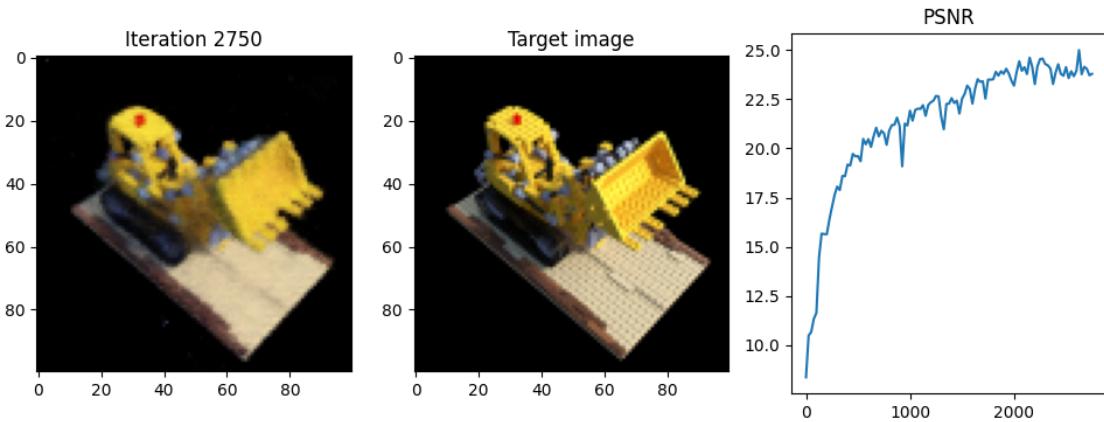
Iteration 2700 Loss: 0.0039 PSNR: 24.04 Time: 0.81 secs per iter, 36.89 mins in total



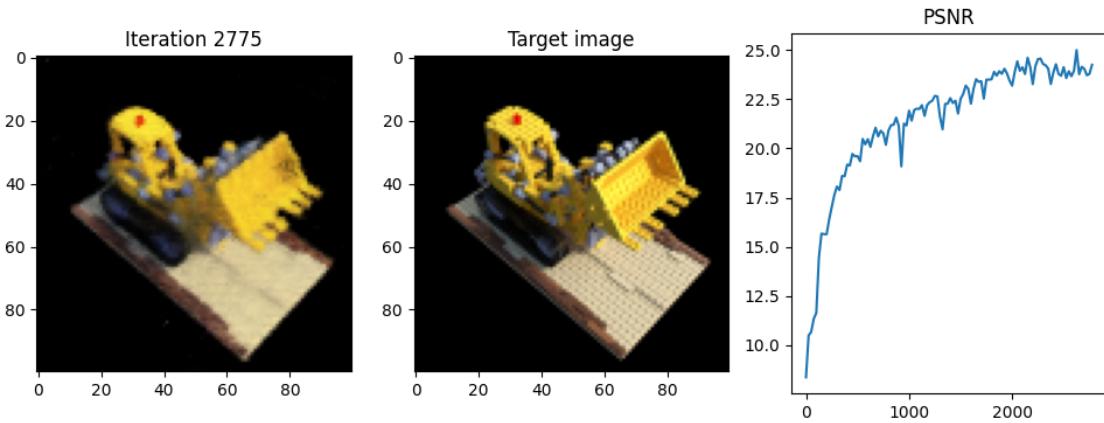
Iteration 2725 Loss: 0.0043 PSNR: 23.71 Time: 0.81 secs per iter, 37.22 mins in total



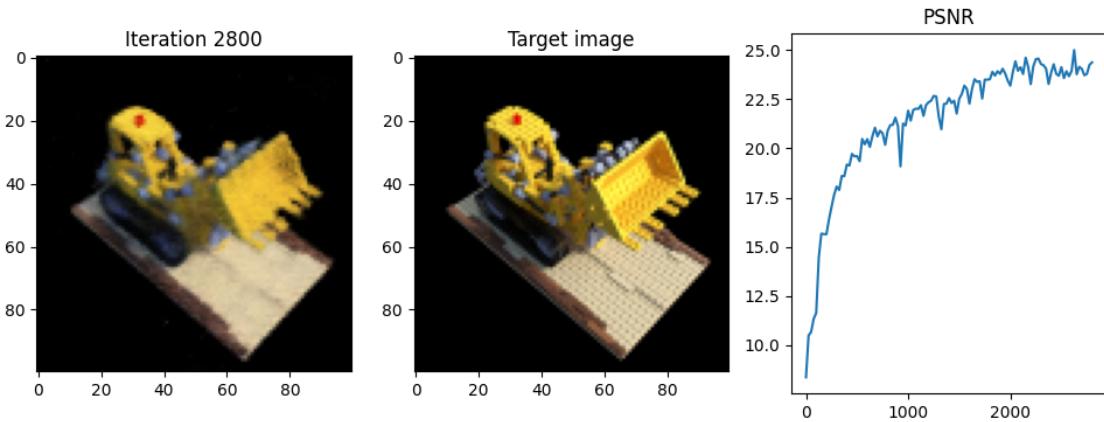
Iteration 2750 Loss: 0.0042 PSNR: 23.79 Time: 0.81 secs per iter, 37.56 mins in total



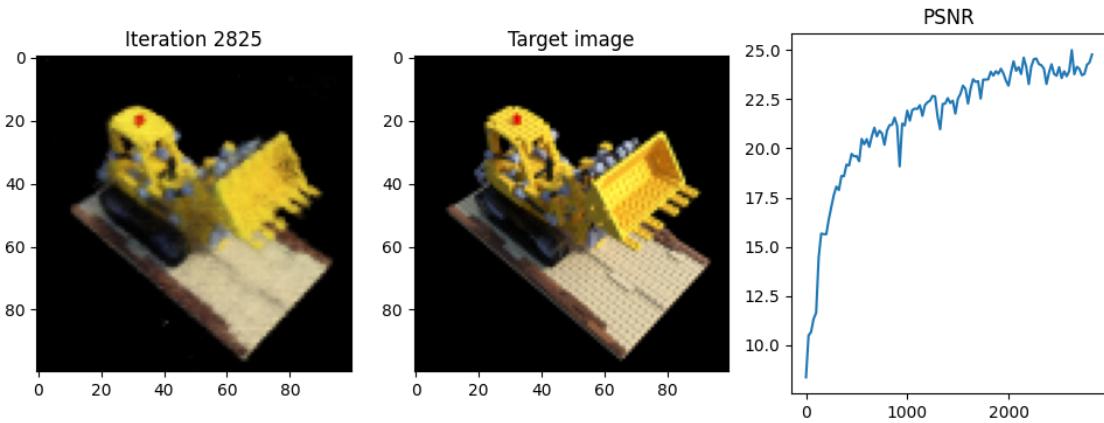
Iteration 2775 Loss: 0.0038 PSNR: 24.24 Time: 0.81 secs per iter, 37.90 mins in total



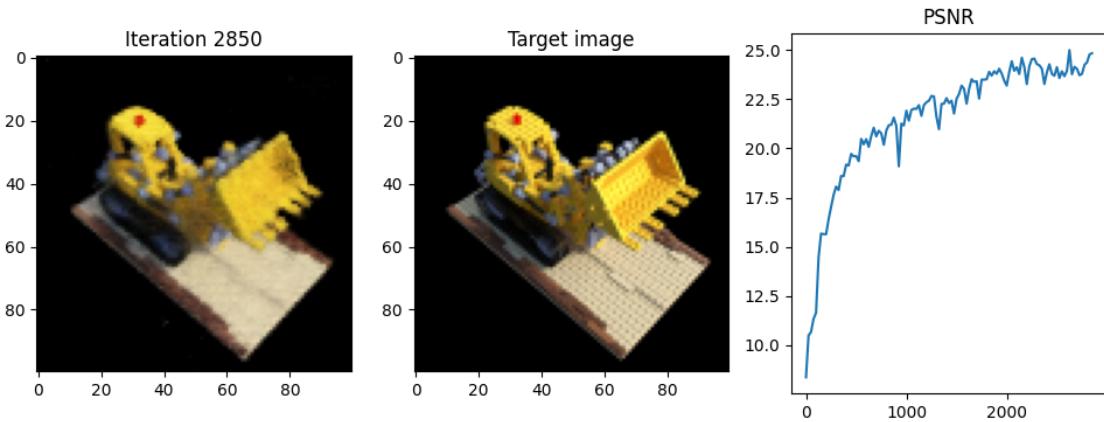
Iteration 2800 Loss: 0.0037 PSNR: 24.37 Time: 0.81 secs per iter, 38.23 mins in total



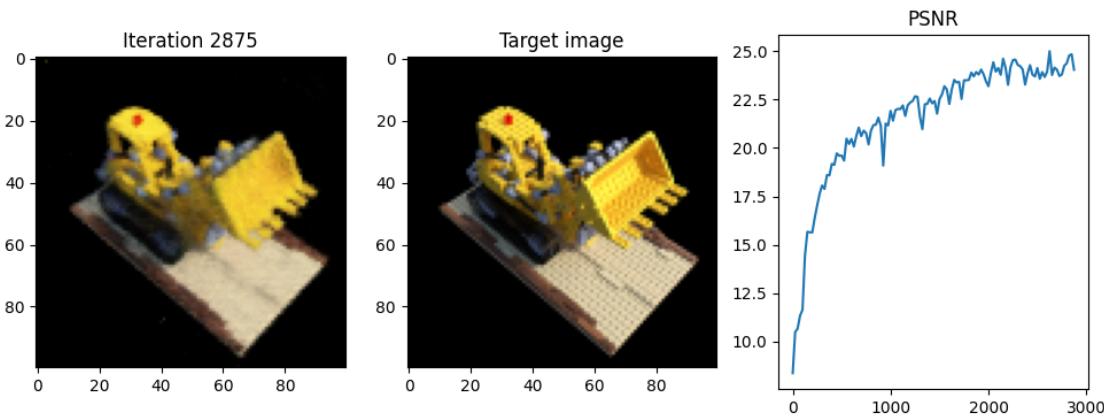
Iteration 2825 Loss: 0.0033 PSNR: 24.77 Time: 0.81 secs per iter, 38.57 mins in total



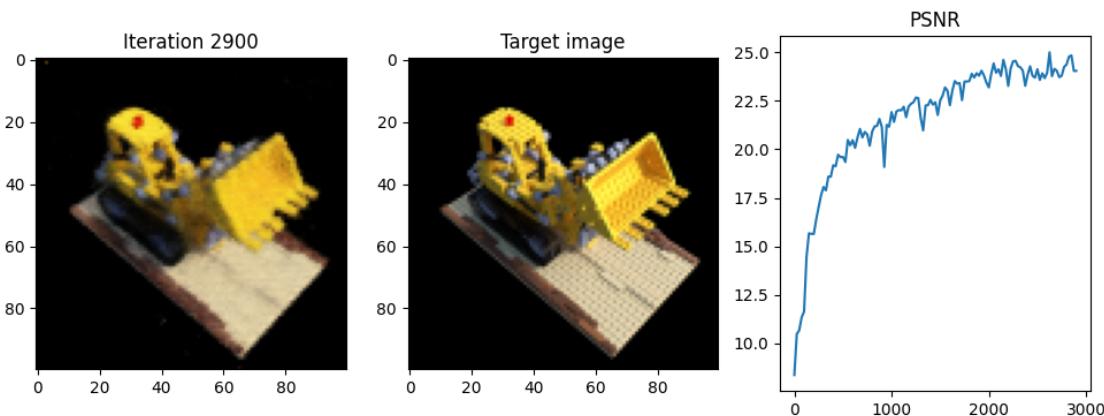
Iteration 2850 Loss: 0.0033 PSNR: 24.83 Time: 0.81 secs per iter, 38.91 mins in total



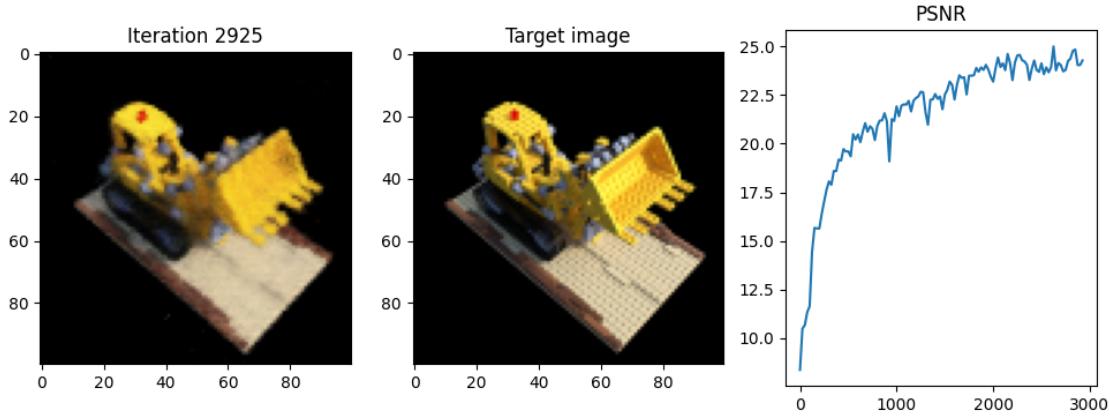
Iteration 2875 Loss: 0.0039 PSNR: 24.04 Time: 0.81 secs per iter, 39.25 mins in total



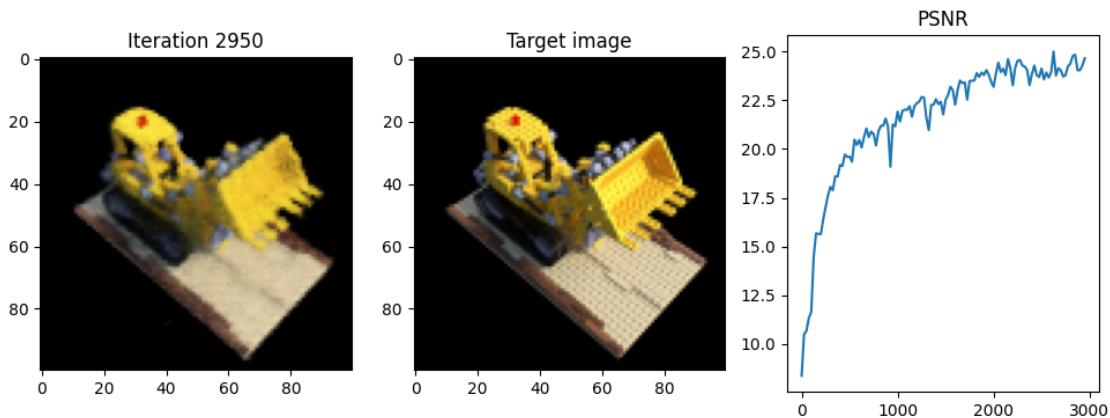
Iteration 2900 Loss: 0.0039 PSNR: 24.04 Time: 0.81 secs per iter, 39.59 mins in total



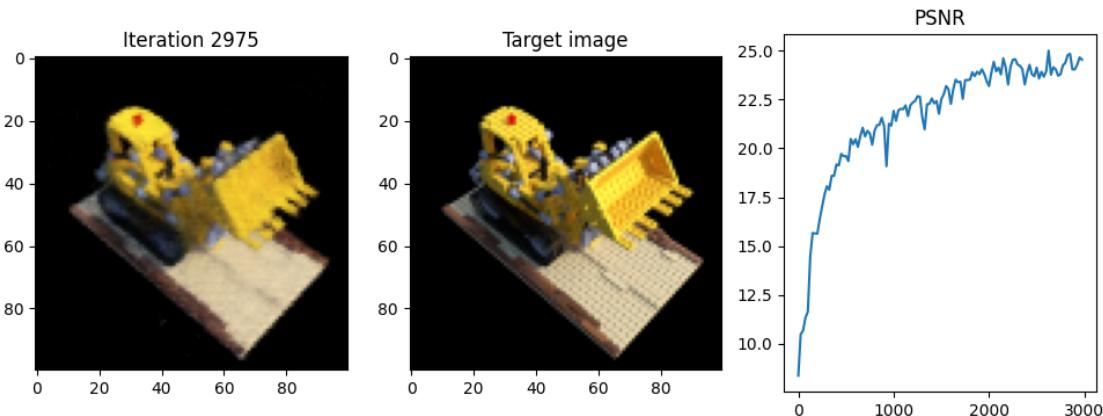
Iteration 2925 Loss: 0.0037 PSNR: 24.27 Time: 0.82 secs per iter, 39.93 mins in total



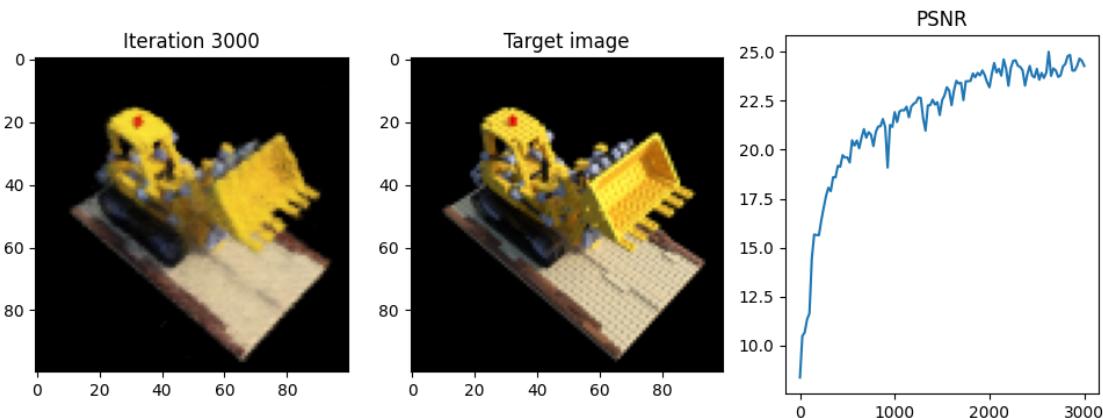
Iteration 2950 Loss: 0.0034 PSNR: 24.65 Time: 0.81 secs per iter, 40.26 mins in total



Iteration 2975 Loss: 0.0035 PSNR: 24.54 Time: 0.81 secs per iter, 40.60 mins in total



Iteration 3000 Loss: 0.0037 PSNR: 24.27 Time: 0.82 secs per iter, 40.94 mins in total



Done!

```
[14]: torch.save(model.state_dict(), 'nerf.pth')
```

```
[15]: !sudo apt-get update &> /dev/null
!sudo apt-get install texlive-xetex texlive-fonts-recommended
↳texlive-plain-generic
!sudo apt-get install texlive-xetex texlive-fonts-recommended
↳texlive-generic-recommended &> /dev/null
!jupyter nbconvert --to pdf cis580_hw5.ipynb
```

Reading package lists... Done
Building dependency tree
Reading state information... Done

The following additional packages will be installed:

```
dvisvgm fonts-droid-fallback fonts-lato fonts-lmodern fonts-noto-mono
fonts-texgyre fonts-urw-base35 javascript-common libapache-pom-java
libcommons-logging-java libcommons-parent-java libfontbox-java libfontenc1
libgs9 libgs9-common libharfbuzz-icu0 libidn11 libijs-0.35 libjbig2dec0
libjs-jquery libkpathsea6 libpdfbox-java libptexenc1 libruby2.7 libsynctex2
libteckit0 libtexlua53 libtexluajit2 libwoff1 libzip-0-13 lmodern
poppler-data preview-latex-style rake ruby ruby-minitest ruby-net-telnet
ruby-power-assert ruby-test-unit ruby-xmlrpc ruby2.7 rubygems-integration
t1utils teckit tex-common tex-gyre texlive-base texlive-binaries
texlive-latex-base texlive-latex-extra texlive-latex-recommended
texlive-pictures tipa xfonts-encodings xfonts-utils
```

Suggested packages:

```
fonts-noto fonts-freefont-otf | fonts-freefont-ttf apache2 | lighttpd
| httpd libavalon-framework-java libcommons-logging-java-doc
libexcalibur-logkit-java liblog4j1.2-java poppler-utils ghostscript
fonts-japanese-mincho | fonts-ipafont-mincho fonts-japanese-gothic
| fonts-ipafont-gothic fonts-aphic-ukai fonts-aphic-uming fonts-nanum ri
ruby-dev bundler debhelper gv | postscript-viewer perl-tk xpdf | pdf-viewer
xzdec texlive-fonts-recommended-doc texlive-latex-base-doc python3-pygments
icc-profiles libfile-which-perl libspreadsheet-parseexcel-perl
texlive-latex-extra-doc texlive-latex-recommended-doc texlive-luatex
texlive-pstricks dot2tex prerex ruby-tcltk | libtcltk-ruby
texlive-pictures-doc vprerex default-jre-headless
```

The following NEW packages will be installed:

```
dvisvgm fonts-droid-fallback fonts-lato fonts-lmodern fonts-noto-mono
fonts-texgyre fonts-urw-base35 javascript-common libapache-pom-java
libcommons-logging-java libcommons-parent-java libfontbox-java libfontenc1
libgs9 libgs9-common libharfbuzz-icu0 libidn11 libijs-0.35 libjbig2dec0
libjs-jquery libkpathsea6 libpdfbox-java libptexenc1 libruby2.7 libsynctex2
libteckit0 libtexlua53 libtexluajit2 libwoff1 libzip-0-13 lmodern
poppler-data preview-latex-style rake ruby ruby-minitest ruby-net-telnet
ruby-power-assert ruby-test-unit ruby-xmlrpc ruby2.7 rubygems-integration
t1utils teckit tex-common tex-gyre texlive-base texlive-binaries
texlive-fonts-recommended texlive-latex-base texlive-latex-extra
texlive-latex-recommended texlive-pictures texlive-plain-generic
texlive-xetex tipa xfonts-encodings xfonts-utils
```

0 upgraded, 58 newly installed, 0 to remove and 25 not upgraded.

Need to get 169 MB of archives.

After this operation, 537 MB of additional disk space will be used.

```
Get:1 http://archive.ubuntu.com/ubuntu focal/main amd64 fonts-droid-fallback all
1:6.0.1r16-1.1 [1,805 kB]
Get:2 http://archive.ubuntu.com/ubuntu focal/main amd64 fonts-lato all 2.0-2
[2,698 kB]
Get:3 http://archive.ubuntu.com/ubuntu focal/main amd64 poppler-data all 0.4.9-2
[1,475 kB]
Get:4 http://archive.ubuntu.com/ubuntu focal/universe amd64 tex-common all 6.13
[32.7 kB]
```

Get:5 http://archive.ubuntu.com/ubuntu focal/main amd64 fonts-urw-base35 all
20170801.1-3 [6,333 kB]
Get:6 http://archive.ubuntu.com/ubuntu focal-updates/main amd64 libgs9-common
all 9.50~dfsg-5ubuntu4.7 [681 kB]
Get:7 http://archive.ubuntu.com/ubuntu focal/main amd64 libidn11 amd64
1.33-2.2ubuntu2 [46.2 kB]
Get:8 http://archive.ubuntu.com/ubuntu focal/main amd64 libijs-0.35 amd64
0.35-15 [15.7 kB]
Get:9 http://archive.ubuntu.com/ubuntu focal/main amd64 libjbig2dec0 amd64
0.18-1ubuntu1 [60.0 kB]
Get:10 http://archive.ubuntu.com/ubuntu focal-updates/main amd64 libgs9 amd64
9.50~dfsg-5ubuntu4.7 [2,173 kB]
Get:11 http://archive.ubuntu.com/ubuntu focal/main amd64 libkpathsea6 amd64
2019.20190605.51237-3build2 [57.0 kB]
Get:12 http://archive.ubuntu.com/ubuntu focal/main amd64 libwoff1 amd64
1.0.2-1build2 [42.0 kB]
Get:13 http://archive.ubuntu.com/ubuntu focal/universe amd64 dvisvgm amd64
2.8.1-1build1 [1,048 kB]
Get:14 http://archive.ubuntu.com/ubuntu focal/universe amd64 fonts-lmodern all
2.004.5-6 [4,532 kB]
Get:15 http://archive.ubuntu.com/ubuntu focal-updates/main amd64 fonts-noto-mono
all 20200323-1build1~ubuntu20.04.1 [80.6 kB]
Get:16 http://archive.ubuntu.com/ubuntu focal/universe amd64 fonts-texgyre all
20180621-3 [10.2 MB]
Get:17 http://archive.ubuntu.com/ubuntu focal/main amd64 javascript-common all
11 [6,066 B]
Get:18 http://archive.ubuntu.com/ubuntu focal/universe amd64 libapache-pom-java
all 18-1 [4,720 B]
Get:19 http://archive.ubuntu.com/ubuntu focal/universe amd64 libcommons-parent-
java all 43-1 [10.8 kB]
Get:20 http://archive.ubuntu.com/ubuntu focal/universe amd64 libcommons-logging-
java all 1.2-2 [60.3 kB]
Get:21 http://archive.ubuntu.com/ubuntu focal/main amd64 libfontenc1 amd64
1:1.1.4-0ubuntu1 [14.0 kB]
Get:22 http://archive.ubuntu.com/ubuntu focal-updates/main amd64 libharfbuzz-
icu0 amd64 2.6.4-1ubuntu4.2 [5,580 B]
Get:23 http://archive.ubuntu.com/ubuntu focal/main amd64 libjs-jquery all
3.3.1~dfsg-3 [329 kB]
Get:24 http://archive.ubuntu.com/ubuntu focal/main amd64 libptexenc1 amd64
2019.20190605.51237-3build2 [35.5 kB]
Get:25 http://archive.ubuntu.com/ubuntu focal/main amd64 rubygems-integration
all 1.16 [5,092 B]
Get:26 http://archive.ubuntu.com/ubuntu focal-updates/main amd64 ruby2.7 amd64
2.7.0-5ubuntu1.10 [95.6 kB]
Get:27 http://archive.ubuntu.com/ubuntu focal/main amd64 ruby amd64 1:2.7+1
[5,412 B]
Get:28 http://archive.ubuntu.com/ubuntu focal/main amd64 rake all 13.0.1-4 [61.6
kB]

Get:29 http://archive.ubuntu.com/ubuntu focal/main amd64 ruby-minitest all
5.13.0-1 [40.9 kB]
Get:30 http://archive.ubuntu.com/ubuntu focal/main amd64 ruby-net-telnet all
0.1.1-2 [12.6 kB]
Get:31 http://archive.ubuntu.com/ubuntu focal/main amd64 ruby-power-assert all
1.1.7-1 [11.4 kB]
Get:32 http://archive.ubuntu.com/ubuntu focal/main amd64 ruby-test-unit all
3.3.5-1 [73.2 kB]
Get:33 http://archive.ubuntu.com/ubuntu focal/main amd64 ruby-xmlrpc all 0.3.0-2
[23.8 kB]
Get:34 http://archive.ubuntu.com/ubuntu focal-updates/main amd64 libruby2.7
amd64 2.7.0-5ubuntu1.10 [3,532 kB]
Get:35 http://archive.ubuntu.com/ubuntu focal/main amd64 libsynctex2 amd64
2019.20190605.51237-3build2 [55.0 kB]
Get:36 http://archive.ubuntu.com/ubuntu focal/universe amd64 libteckit0 amd64
2.5.8+ds2-5ubuntu2 [320 kB]
Get:37 http://archive.ubuntu.com/ubuntu focal/main amd64 libtexlua53 amd64
2019.20190605.51237-3build2 [105 kB]
Get:38 http://archive.ubuntu.com/ubuntu focal/main amd64 libtexluajit2 amd64
2019.20190605.51237-3build2 [235 kB]
Get:39 http://archive.ubuntu.com/ubuntu focal/universe amd64 libzzip-0-13 amd64
0.13.62-3.2ubuntu1 [26.2 kB]
Get:40 http://archive.ubuntu.com/ubuntu focal/main amd64 xfonts-encodings all
1:1.0.5-0ubuntu1 [573 kB]
Get:41 http://archive.ubuntu.com/ubuntu focal/main amd64 xfonts-utils amd64
1:7.7+6 [91.5 kB]
Get:42 http://archive.ubuntu.com/ubuntu focal/universe amd64 lmodern all
2.004.5-6 [9,474 kB]
Get:43 http://archive.ubuntu.com/ubuntu focal/universe amd64 preview-latex-style
all 11.91-2ubuntu2 [184 kB]
Get:44 http://archive.ubuntu.com/ubuntu focal/main amd64 t1utils amd64 1.41-3
[56.1 kB]
Get:45 http://archive.ubuntu.com/ubuntu focal/universe amd64 teckit amd64
2.5.8+ds2-5ubuntu2 [687 kB]
Get:46 http://archive.ubuntu.com/ubuntu focal/universe amd64 tex-gyre all
20180621-3 [6,209 kB]
Get:47 http://archive.ubuntu.com/ubuntu focal/universe amd64 texlive-binaries
amd64 2019.20190605.51237-3build2 [8,041 kB]
Get:48 http://archive.ubuntu.com/ubuntu focal/universe amd64 texlive-base all
2019.20200218-1 [20.8 MB]
Get:49 http://archive.ubuntu.com/ubuntu focal/universe amd64 texlive-fonts-
recommended all 2019.20200218-1 [4,972 kB]
Get:50 http://archive.ubuntu.com/ubuntu focal/universe amd64 texlive-latex-base
all 2019.20200218-1 [990 kB]
Get:51 http://archive.ubuntu.com/ubuntu focal/universe amd64 libfontbox-java all
1:1.8.16-2 [207 kB]
Get:52 http://archive.ubuntu.com/ubuntu focal/universe amd64 libpdfbox-java all
1:1.8.16-2 [5,199 kB]

```
Get:53 http://archive.ubuntu.com/ubuntu focal/universe amd64 texlive-latex-
recommended all 2019.20200218-1 [15.7 MB]
Get:54 http://archive.ubuntu.com/ubuntu focal/universe amd64 texlive-pictures
all 2019.20200218-1 [4,492 kB]
Get:55 http://archive.ubuntu.com/ubuntu focal/universe amd64 texlive-latex-extra
all 2019.20200218-1 [12.5 MB]
Get:56 http://archive.ubuntu.com/ubuntu focal/universe amd64 texlive-plain-
generic all 2019.20200218-1 [24.6 MB]
Get:57 http://archive.ubuntu.com/ubuntu focal/universe amd64 tipa all 2:1.3-20
[2,978 kB]
Get:58 http://archive.ubuntu.com/ubuntu focal/universe amd64 texlive-xetex all
2019.20200218-1 [14.6 MB]
Fetched 169 MB in 26s (6,576 kB/s)
debconf: unable to initialize frontend: Dialog
debconf: (No usable dialog-like program is installed, so the dialog based
frontend cannot be used. at /usr/share/perl5/Debconf/FrontEnd/Dialog.pm line 76,
<> line 58.)
debconf: falling back to frontend: Readline
debconf: unable to initialize frontend: Readline
debconf: (This frontend requires a controlling tty.)
debconf: falling back to frontend: Teletype
dpkg-preconfigure: unable to re-open stdin:
Selecting previously unselected package fonts-droid-fallback.
(Reading database ... 122518 files and directories currently installed.)
Preparing to unpack .../00-fonts-droid-fallback_1%3a6.0.1r16-1.1_all.deb ...
Unpacking fonts-droid-fallback (1:6.0.1r16-1.1) ...
Selecting previously unselected package fonts-lato.
Preparing to unpack .../01-fonts-lato_2.0-2_all.deb ...
Unpacking fonts-lato (2.0-2) ...
Selecting previously unselected package poppler-data.
Preparing to unpack .../02-poppler-data_0.4.9-2_all.deb ...
Unpacking poppler-data (0.4.9-2) ...
Selecting previously unselected package tex-common.
Preparing to unpack .../03-tex-common_6.13_all.deb ...
Unpacking tex-common (6.13) ...
Selecting previously unselected package fonts-urw-base35.
Preparing to unpack .../04-fonts-urw-base35_20170801.1-3_all.deb ...
Unpacking fonts-urw-base35 (20170801.1-3) ...
Selecting previously unselected package libgs9-common.
Preparing to unpack .../05-libgs9-common_9.50~dfsg-5ubuntu4.7_all.deb ...
Unpacking libgs9-common (9.50~dfsg-5ubuntu4.7) ...
Selecting previously unselected package libidn11:amd64.
Preparing to unpack .../06-libidn11_1.33-2.2ubuntu2_amd64.deb ...
Unpacking libidn11:amd64 (1.33-2.2ubuntu2) ...
Selecting previously unselected package libijs-0.35:amd64.
Preparing to unpack .../07-libijs-0.35_0.35-15_amd64.deb ...
Unpacking libijs-0.35:amd64 (0.35-15) ...
Selecting previously unselected package libjbig2dec0:amd64.
```

```
Preparing to unpack .../08-libjbig2dec0_0.18-1ubuntu1_amd64.deb ...
Unpacking libjbig2dec0:amd64 (0.18-1ubuntu1) ...
Selecting previously unselected package libgs9:amd64.
Preparing to unpack .../09-libgs9_9.50~dfsg-5ubuntu4.7_amd64.deb ...
Unpacking libgs9:amd64 (9.50~dfsg-5ubuntu4.7) ...
Selecting previously unselected package libkpathsea6:amd64.
Preparing to unpack .../10-libkpathsea6_2019.20190605.51237-3build2_amd64.deb
...
Unpacking libkpathsea6:amd64 (2019.20190605.51237-3build2) ...
Selecting previously unselected package libwoff1:amd64.
Preparing to unpack .../11-libwoff1_1.0.2-1build2_amd64.deb ...
Unpacking libwoff1:amd64 (1.0.2-1build2) ...
Selecting previously unselected package dvisvgm.
Preparing to unpack .../12-dvisvgm_2.8.1-1build1_amd64.deb ...
Unpacking dvisvgm (2.8.1-1build1) ...
Selecting previously unselected package fonts-lmodern.
Preparing to unpack .../13-fonts-lmodern_2.004.5-6_all.deb ...
Unpacking fonts-lmodern (2.004.5-6) ...
Selecting previously unselected package fonts-noto-mono.
Preparing to unpack .../14-fonts-noto-
mono_20200323-1build1~ubuntu20.04.1_all.deb ...
Unpacking fonts-noto-mono (20200323-1build1~ubuntu20.04.1) ...
Selecting previously unselected package fonts-texgyre.
Preparing to unpack .../15-fonts-texgyre_20180621-3_all.deb ...
Unpacking fonts-texgyre (20180621-3) ...
Selecting previously unselected package javascript-common.
Preparing to unpack .../16-javascript-common_11_all.deb ...
Unpacking javascript-common (11) ...
Selecting previously unselected package libapache-pom-java.
Preparing to unpack .../17-libapache-pom-java_18-1_all.deb ...
Unpacking libapache-pom-java (18-1) ...
Selecting previously unselected package libcommons-parent-java.
Preparing to unpack .../18-libcommons-parent-java_43-1_all.deb ...
Unpacking libcommons-parent-java (43-1) ...
Selecting previously unselected package libcommons-logging-java.
Preparing to unpack .../19-libcommons-logging-java_1.2-2_all.deb ...
Unpacking libcommons-logging-java (1.2-2) ...
Selecting previously unselected package libfontenc1:amd64.
Preparing to unpack .../20-libfontenc1_1%3a1.1.4-0ubuntu1_amd64.deb ...
Unpacking libfontenc1:amd64 (1:1.1.4-0ubuntu1) ...
Selecting previously unselected package libharfbuzz-icu0:amd64.
Preparing to unpack .../21-libharfbuzz-icu0_2.6.4-1ubuntu4.2_amd64.deb ...
Unpacking libharfbuzz-icu0:amd64 (2.6.4-1ubuntu4.2) ...
Selecting previously unselected package libjs-jquery.
Preparing to unpack .../22-libjs-jquery_3.3.1~dfsg-3_all.deb ...
Unpacking libjs-jquery (3.3.1~dfsg-3) ...
Selecting previously unselected package libptexenc1:amd64.
Preparing to unpack .../23-libptexenc1_2019.20190605.51237-3build2_amd64.deb ...
```

```
Unpacking libptexenc1:amd64 (2019.20190605.51237-3build2) ...
Selecting previously unselected package rubygems-integration.
Preparing to unpack .../24-rubygems-integration_1.16_all.deb ...
Unpacking rubygems-integration (1.16) ...
Selecting previously unselected package ruby2.7.
Preparing to unpack .../25-ruby2.7_2.7.0-5ubuntu1.10_amd64.deb ...
Unpacking ruby2.7 (2.7.0-5ubuntu1.10) ...
Selecting previously unselected package ruby.
Preparing to unpack .../26-ruby_1%3a2.7+1_amd64.deb ...
Unpacking ruby (1:2.7+1) ...
Selecting previously unselected package rake.
Preparing to unpack .../27-rake_13.0.1-4_all.deb ...
Unpacking rake (13.0.1-4) ...
Selecting previously unselected package ruby-minitest.
Preparing to unpack .../28-ruby-minitest_5.13.0-1_all.deb ...
Unpacking ruby-minitest (5.13.0-1) ...
Selecting previously unselected package ruby-net-telnet.
Preparing to unpack .../29-ruby-net-telnet_0.1.1-2_all.deb ...
Unpacking ruby-net-telnet (0.1.1-2) ...
Selecting previously unselected package ruby-power-assert.
Preparing to unpack .../30-ruby-power-assert_1.1.7-1_all.deb ...
Unpacking ruby-power-assert (1.1.7-1) ...
Selecting previously unselected package ruby-test-unit.
Preparing to unpack .../31-ruby-test-unit_3.3.5-1_all.deb ...
Unpacking ruby-test-unit (3.3.5-1) ...
Selecting previously unselected package ruby-xmlrpc.
Preparing to unpack .../32-ruby-xmlrpc_0.3.0-2_all.deb ...
Unpacking ruby-xmlrpc (0.3.0-2) ...
Selecting previously unselected package libruby2.7:amd64.
Preparing to unpack .../33-libruby2.7_2.7.0-5ubuntu1.10_amd64.deb ...
Unpacking libruby2.7:amd64 (2.7.0-5ubuntu1.10) ...
Selecting previously unselected package libsynctex2:amd64.
Preparing to unpack .../34-libsynctex2_2019.20190605.51237-3build2_amd64.deb ...
Unpacking libsynctex2:amd64 (2019.20190605.51237-3build2) ...
Selecting previously unselected package libteckit0:amd64.
Preparing to unpack .../35-libteckit0_2.5.8+ds2-5ubuntu2_amd64.deb ...
Unpacking libteckit0:amd64 (2.5.8+ds2-5ubuntu2) ...
Selecting previously unselected package libtexlua53:amd64.
Preparing to unpack .../36-libtexlua53_2019.20190605.51237-3build2_amd64.deb ...
Unpacking libtexlua53:amd64 (2019.20190605.51237-3build2) ...
Selecting previously unselected package libtexluajit2:amd64.
Preparing to unpack .../37-libtexluajit2_2019.20190605.51237-3build2_amd64.deb
...
Unpacking libtexluajit2:amd64 (2019.20190605.51237-3build2) ...
Selecting previously unselected package libzip-0-13:amd64.
Preparing to unpack .../38-libzip-0-13_0.13.62-3.2ubuntu1_amd64.deb ...
Unpacking libzip-0-13:amd64 (0.13.62-3.2ubuntu1) ...
Selecting previously unselected package xfonts-encodings.
```

```
Preparing to unpack .../39-xfonts-encodings_1%3a1.0.5-0ubuntu1_all.deb ...
Unpacking xfonts-encodings (1:1.0.5-0ubuntu1) ...
Selecting previously unselected package xfonts-utils.
Preparing to unpack .../40-xfonts-utils_1%3a7.7+6_amd64.deb ...
Unpacking xfonts-utils (1:7.7+6) ...
Selecting previously unselected package lmodern.
Preparing to unpack .../41-lmodern_2.004.5-6_all.deb ...
Unpacking lmodern (2.004.5-6) ...
Selecting previously unselected package preview-latex-style.
Preparing to unpack .../42-preview-latex-style_11.91-2ubuntu2_all.deb ...
Unpacking preview-latex-style (11.91-2ubuntu2) ...
Selecting previously unselected package t1utils.
Preparing to unpack .../43-t1utils_1.41-3_amd64.deb ...
Unpacking t1utils (1.41-3) ...
Selecting previously unselected package teckit.
Preparing to unpack .../44-teckit_2.5.8+ds2-5ubuntu2_amd64.deb ...
Unpacking teckit (2.5.8+ds2-5ubuntu2) ...
Selecting previously unselected package tex-gyre.
Preparing to unpack .../45-tex-gyre_20180621-3_all.deb ...
Unpacking tex-gyre (20180621-3) ...
Selecting previously unselected package texlive-binaries.
Preparing to unpack .../46-texlive-
binaries_2019.20190605.51237-3build2_amd64.deb ...
Unpacking texlive-binaries (2019.20190605.51237-3build2) ...
Selecting previously unselected package texlive-base.
Preparing to unpack .../47-texlive-base_2019.20200218-1_all.deb ...
Unpacking texlive-base (2019.20200218-1) ...
Selecting previously unselected package texlive-fonts-recommended.
Preparing to unpack .../48-texlive-fonts-recommended_2019.20200218-1_all.deb ...
Unpacking texlive-fonts-recommended (2019.20200218-1) ...
Selecting previously unselected package texlive-latex-base.
Preparing to unpack .../49-texlive-latex-base_2019.20200218-1_all.deb ...
Unpacking texlive-latex-base (2019.20200218-1) ...
Selecting previously unselected package libfontbox-java.
Preparing to unpack .../50-libfontbox-java_1%3a1.8.16-2_all.deb ...
Unpacking libfontbox-java (1:1.8.16-2) ...
Selecting previously unselected package libpdfbox-java.
Preparing to unpack .../51-libpdfbox-java_1%3a1.8.16-2_all.deb ...
Unpacking libpdfbox-java (1:1.8.16-2) ...
Selecting previously unselected package texlive-latex-recommended.
Preparing to unpack .../52-texlive-latex-recommended_2019.20200218-1_all.deb ...
Unpacking texlive-latex-recommended (2019.20200218-1) ...
Selecting previously unselected package texlive-pictures.
Preparing to unpack .../53-texlive-pictures_2019.20200218-1_all.deb ...
Unpacking texlive-pictures (2019.20200218-1) ...
Selecting previously unselected package texlive-latex-extra.
Preparing to unpack .../54-texlive-latex-extra_2019.202000218-1_all.deb ...
Unpacking texlive-latex-extra (2019.202000218-1) ...
```

```

-----
KeyboardInterrupt                                     Traceback (most recent call last)
<ipython-input-15-99dd16edb3c2> in <cell line: 2>()
      1 get_ipython().system('sudo apt-get update &> /dev/null')
--> 2 get_ipython().system('sudo apt-get install texlive-xetex'
   ↪texlive-fonts-recommended texlive-plain-generic')
      3 get_ipython().system('sudo apt-get install texlive-xetex'
   ↪texlive-fonts-recommended texlive-generic-recommended &> /dev/null')
      4 get_ipython().system('jupyter nbconvert --to pdf cis580_hw5.ipynb')

/usr/local/lib/python3.10/dist-packages/google/colab/_shell.py in system(self,*
   ↪*args, **kwargs)
    97         kwargs.update({'also_return_output': True})
    98
--> 99     output = _system_commands._system_compat(self, *args, **kwargs)  #U
   ↪ pylint:disable=protected-access
   100
   101     if pip_warn:

/usr/local/lib/python3.10/dist-packages/google/colab/_system_commands.py in
   ↪_system_compat(shell, cmd, also_return_output)
   451     # is expected to call this function, thus adding one level of nesting
   ↪ to the
   452     # stack.
--> 453     result = _run_command(
   454         shell.var_expand(cmd, depth=2), clear_stamed_output=False
   455     )

/usr/local/lib/python3.10/dist-packages/google/colab/_system_commands.py in
   ↪_run_command(cmd, clear_stamed_output)
   201     os.close(child_pty)
   202
--> 203     return _monitor_process(parent_pty, epoll, p, cmd,*
   ↪update_stdin_widget)
   204     finally:
   205         epoll.close()

/usr/local/lib/python3.10/dist-packages/google/colab/_system_commands.py in
   ↪_monitor_process(parent_pty, epoll, p, cmd, update_stdin_widget)
   231     while True:
   232         try:
--> 233             result = _poll_process(parent_pty, epoll, p, cmd, decoder, state)
   234             if result is not None:
   235                 return result

```

```
/usr/local/lib/python3.10/dist-packages/google/colab/_system_commands.py in __poll_process(parent_pty, epoll, p, cmd, decoder, state)
 279     output_available = False
 280
--> 281     events = epoll.poll()
 282     input_events = []
 283     for _, event in events:
```

```
KeyboardInterrupt:
```